

# Pyg documentation

version

**Yoav Git**

October 25, 2021



# Contents

<b>Welcome to pyg!</b>	<b>1</b>
README	1
pyg.base	1
extensions to dict	1
dictattr	1
ulist	3
Dict	4
dictable	4
perdictable	12
join	13
named_dict	14
decorators	16
wrapper	16
timer	17
try_value	18
try_back	18
loops	18
loop	19
kwargs_support	19
graphs & cells	19
cell	19
cell_go	21
cell_item	22
cell_func	22
cell_clear	23
encode and decode/save and load	23
encode	23
decode	24
pd_to_parquet	24
pd_read_parquet	25
parquet_encode	25
csv_encode	26
convertors to bytes	26
dates and calendar	26
dt	26
ymd	27
dt_bump	27
drange	28
date_range	28
Calendar	28

calendar	31
as_time	31
clock	31
text manipulation	32
lower	32
upper	32
proper	32
capitalize	33
strip	33
split	33
replace	34
common_prefix	34
files & directory	34
mkdir	34
read_csv	34
tree manipulation	34
tree_keys	35
tree_values	35
tree_items	35
tree_update	36
tree_setitem	37
tree_repr	37
items_to_tree	38
tree_to_table	39
list functions	40
as_list	40
as_tuple	40
first	41
last	41
unique	41
Comparing and Sorting	41
cmp	41
Cmp	42
sort	42
eq	42
in	43
bits and pieces	43
type functions	43
zipper	44
reducer	45
reducing	45
logger and get_logger	46

access functions	46
inspection	47
pyg.mongo	50
Query generator	50
q and Q	50
Tables in Mongo	51
mongo_cursor	51
mongo_reader	55
mongo_pk_reader	58
mongo_pk_cursor	58
encoding docs before saving to mongo	58
parquet_write	58
csv_write	59
cells in Mongo	59
db_cell	59
periodic_cell	61
get_cell	62
db_save	62
db_load	62
pyg.timeseries	62
simple functions	63
diff	63
shift	63
ratio	64
ts_count	65
ts_sum	65
ts_mean	66
ts_rms	66
ts_std	67
ts_skew	67
ts_min	68
ts_max	68
ts_median	68
fna	68
v2na/na2v	68
ffill/bfill	69
nona	70
expanding window functions	70
expanding_mean	70
expanding_rms	71
expanding_std	72
expanding_sum	73

expanding_skew	74
expanding_min	75
expanding_max	76
expanding_median	77
expanding_rank	78
cumsum	78
cumprod	79
rolling window functions	80
rolling_mean	80
rolling_rms	82
rolling_std	83
rolling_sum	84
rolling_skew	85
rolling_min	86
rolling_max	87
rolling_median	88
rolling_quantile	89
rolling_rank	90
exponentially weighted moving functions	90
ewma	90
ewmrms	92
ewmstd	93
ewmskew	94
ewmvar	96
ewmcor	97
ewmcorr	98
ewmLR	99
ewmGLM	101
ewmxo	102
functions exposing their state	102
simple functions	102
expanding window functions	103
rolling window functions	103
exponentially weighted moving functions	104
Index handling	104
df_fillna	104
df_index	105
df_reindex	106
presync	106
add/sub/mul/div/pow operators	108
Tutorials	108
pyg.base.Dict	109

initialization	109
members access	109
adding	109
subtracting	109
modifying the keys: rename	110
modifying the values: do	110
Dict can store a calculation flow	110
pyg.base.dictable	111
Motivation: dictable as an organiser of research flow	111
Same code, in dictable	112
Oh, no, we have a bad symbol, how do we remove it?	112
Now if we want to calculate something per symbol and window...	112
dictable functionality	112
construction	112
row access	113
column access	113
d is a dict so supports the usual keys(), values() and items():	113
column and row access are commutative	113
adding records	113
adding/modifying columns	113
do	114
removing columns	114
removing rows	114
sort	114
listby(keys)	114
unlist	114
groupby(keys) and ungroup	115
inner join	115
inner join (with other columns that match names)	115
cross join	115
xor (versus left and right join)	116
pivot	116
a few observations:	116
pyg.mongo	116
q	117
mongo_cursor	117
general objects insertion into documents	118
document reading	118
document writing to files	119
document access	119
filters	119
iteration	120

sorting	120
getitem of a specific document	120
column access	120
add/remove columns	120
add/remove records	120
mongo_pk_table	121
mongo_reader and mongo_pk_reader	121
pyg.base.cell	122
Cell 101	122
Workflow without saving to the database	122
some more functions to calculate the profits & loss as well as the signal/noise ratio	123
Workflow while saving to MongoDB	123
Table creation	123
Any code differences?	123
Accessing the data in MongoDB	124
Calculating the forecasts & saving them	124
Accessing & running the graph once the graph has been created	124
Point-in-time, cache and persistency	124
Comparison of the two workflows	125
To save or not to save?	126
Behind the scene: cell_func	126
pyg.base.join	127
Join	127
Example: Using join function to transfer money to a bank	127
Simple join: inner join between tables	128
Defaults for fields we want to left-join on...	128
Renaming & calculating fields	128
Perdictable	128
perdictable and caching	129
perdictable with the cell framework	129
perdictable API	129
Conclusions	130
pyg.timeseries	130
Agreement between pyg.timeseries and pandas	130
Quick performance comparison	131
pyg and numpy arrays	131
pandas treatment of nan	132
pyg.timeseries treatment of nans	132
Using pyg.timeseries to manage state	132
Example: creating a function exposing its state	133
pyg.timeseries decorators	134
loop	135



presync: manage indexing and date stamps	135
presync and numpy arrays	136
pyg.timeseries.ewma	137
What happens if the clock does not move at all?	138
What are valid time parameters?	140
Indices and tables	140
<b>Python Module Index</b>	<b>141</b>



# Welcome to pyg!

## README

- If you examine data by multiple dimensions, you need `pyg.base.dictable`.
- If you use MongoDB, you need `pyg.mongo`.
- If you use pandas for timeseries analysis, you should consider using `pyg.timeseries`.

pyg is both succinct and powerful and makes your code almost boilerplate free and easy to maintain. As an example, I estimate that Man AHL, a leading quant hedge fund, relies on about 50 coders to replicate the functionality and maintain boilerplate code that pyg would make redundant.

Below is autodoc created by sphinx followed by tutorials created in jupyter notebooks.

## pyg.base

### *extensions to dict*

#### *dictattr*

`class pyg.base._dictattr.dictattr`

**A simple dict with extended member manipulation**

1. access using `d.key`
2. access multiple elements using `d[key1, key2]`

**Example:** members access

```
>>> from pyg import *
>>> d = dictattr(a = 1, b = 2, c = 3)
>>> assert isinstance(d, dict)
>>> assert d.a == 1
>>> assert d['a', 'b'] == [1, 2]
>>> assert d[['a', 'b']] == dictattr(a = 1, b = 2)
```

In addition, it has extended key selection/subsetting

**Example:** subsetting

```
>>> d = dictattr(a = 1, b = 2, c = 3)
>>> assert d - 'a' == dictattr(b = 2, c = 3)
>>> assert d & ['b', 'c', 'not in keys'] == dictattr(b = 2, c = 3)
```

`dictattr` supports not in-place 'update':

**Example:** updating via adding another dict

```
>>> d = dictattr(a = 1, b = 2) + dict(b = 'replacing old value', c = 'new key')
>>> assert d == dictattr(a = 1, b = 'replacing old value', c = 'new key')
```

`copy ()` → a shallow copy of D

`keys ()`

`dictattr` returns an actual list rather than a generator. Further, this recognises that the keys are necessarily unique so it returns a `ulist` which is also a set

**Returns:**

Welcome to pyg!

## ulist

list of keys of dictattr.

### Example:

```
>>> from pyg import *
>>> d = dictattr(a = 1, b = 2)
>>> assert d.keys() == ulist(['a', 'b'])
>>> assert d.keys() & ['a', 'c', 'd'] == ['a']
```

## relabel(\*args, \*\*relabels)

easy relabel/rename of keys

### Parameters:

**\*args** : *str or callable*

- a string ending/starting with `_` will trigger a prefix/suffix to all keys
- callable function will be applied to the keys to update them

**\*\*relabels** : *strings*

individual relabeling of keys

### Returns:

## dictattr

new dict with renamed keys.

### Example: suffix/prefix

```
>>> from pyg import *
>>> d = dictattr(a = 1, b = 2, c = 3)
>>> assert d.relabel('x_') == dictattr(x_a = 1, x_b = 2, x_c = 3) # prefixing
>>> assert d.relabel('_x') == dictattr(a_x = 1, b_x = 2, c_x = 3) # suffixing
```

### Example: callable

```
>>> assert d.rename(upper) == dictattr(A = 1, B = 2, C = 3)
```

### Example: individual relabelling

```
>>> assert d.rename(a = 'A') == dictattr(A = 1, b = 2, c = 3)
>>> assert d.rename(['A', 'B', 'C']) == d.relabel(upper)
```

## rename(\*args, \*\*relabels)

Identical to relabel. See relabel for full docs

**values()** → an object providing a view on D's values

`pyg.base._dictattr.dictattr.__add__` (self, other)

dictattr uses add as a copy + update. Similar to the latest python `|=`

### Example:

```
>>> from pyg import *
>>> d = dictattr(a = 1, b = 2)
>>> assert d + dict(b = 3, c = 5) == dictattr(a = 1, b = 3, c = 5)
```

### Parameters:

**other: dict**

a dict used to update current dict.

`pyg.base._dictattr.dictattr.__sub__` (self, key, copy=True)

Welcome to pyg!

deletes an item but does not throw an exception if not there dictattr uses subtraction to remove key(s)

**Returns:**

updated dictattr

**Example:**

```
>>> from pyg import *
>>> d = dictattr(a = 1, b = 2, c = 3)
>>> assert d - ['b', 'c'] == dictattr(a = 1)
>>> assert d - 'c' == dictattr(a = 1, b = 2)
>>> assert d - 'key not there' == d
>>> #commutative
>>> assert (d - 'c').keys() == d.keys() - 'c'
```

pyg.base.\_dictattr.dictattr.\_\_and\_\_(self, other)  
dictattr uses & as a set operator for key filtering

**Returns:**

updated dictattr

**Example:**

```
>>> from pyg import *
>>> d = dictattr(a = 1, b = 2, c = 3)
>>> assert d & ['a', 'b', 'not_there'] == dictattr(a = 1, b = 2)
>>> #commutative
>>> assert (d & ['a', 'b', 'x']).keys() == d.keys() & ['a', 'b', 'x']
```

## ulist

The dictattr.keys() method returns a ulist: a list with unique elements:

class pyg.base.\_ulist.ulist(\*args, unique=False)

A list whose members are unique. It has +/- operations overloaded while also supporting set operations &/

**Example:**

```
>>> assert ulist([1,3,2,1]) == list([1,3,2])
```

**Example:** addition adds element(s)

```
>>> assert ulist([1,3,2,1]) + 4 == list([1,3,2,4])
>>> assert ulist([1,3,2,1]) + [4,1] == list([1,3,2,4])
>>> assert ulist([1,3,2,1]) + [4,1,5] == list([1,3,2,4,5])
```

**Example:** subtraction removes element(s)

```
>>> assert ulist([1,3,2,1]) - 1 == [3,2]
>>> assert ulist([1,3,2,1]) - [1,3,4] == [2]
```

**Example:** set operations

```
>>> assert ulist([1,3,2,1]) & 1 == [1]
>>> assert ulist([1,3,2,1]) & [1,3,4] == [1,3]
```

```
>>> assert ulist([1,3,2,1]) | 1 == [1,3,2]
>>> assert ulist([1,3,2,1]) | 4 == [1,3,2,4]
>>> assert ulist([1,3,2,1]) | [1,3,4] == [1,3,2,4]
```

**copy()**

Return a shallow copy of the list.

## Dict

`class pyg.base._dict.Dict`

Dict extends dictattr to allow access to *functions* of members

### Example:

```
>>> from pyg import *
>>> d = Dict(a = 1, b=2)
>>> assert d[lambda a, b: a+b] == 3
>>> assert d['a', 'b', lambda a,b: a+b] == [1,2,3]
```

Dict is also callable where the key-value is used to add/update current members

### Example:

```
>>> from pyg import *
>>> d = Dict(a = 1, b=2)
>>> assert d(c = 3) == Dict(a = 1, b = 2, c = 3)
>>> assert d(c = lambda a,b: a+b) == Dict(a = 1, b = 2, c = 3)
```

```
>>> assert d(c = 3) == Dict(a = 1, b = 2) + Dict(c = 3)
>>> assert Dict(a = 1)(b = lambda a: a+1)(c = lambda a,b: a+b) == Dict(a = 1,b = 2,c = 3)
```

**do** (function, \*keys)

applies a function(s) on multiple keys at the same time

### Parameters:

**function** : *callable or list of callables*

function to be applied per column

**\*keys** : *string/list of strings*

list of columns to be applied. If missing, applied to all columns

### Returns:

res : Dict

### Example:

```
>>> from pyg import *
>>> d = Dict(name = 'adam', surname = 'atkins')
>>> assert d.do(proper) == Dict(name = 'Adam', surname = 'Atkins')
```

**Example:** using another key in the calculation

```
>>> from pyg import *
>>> d = Dict(a = 1, b = 5, denominator = 10)
>>> d = d.do(lambda value, denominator: value/denominator, 'a', 'b')
>>> assert d == Dict(a = 0.1, b = 0.5, denominator = 10)
```

`pyg.base._dict.Dict.__call__` (self, \*\*kwargs)

Call self as a function.

## dictable

`class pyg.base._dictable.dictable` (data=None, columns=None, \*\*kwargs)

### What is dictable?:

dictable is a table, a collection of iterable records. It is also a dict with each key being a column. Why not use a pandas.DataFrame? pd.DataFrame leads a dual life:

- by day an index-based optimized numpy array supporting e.g. timeseries analytics etc.
- by night, a table with keys supporting filtering, aggregating, pivoting on keys as well as inner/outer joining on keys.

dictable only tries to do the latter. dictable should be thought of as a 'container for complicated objects' rather than just an array of primitive floats. In general, each cell may contain timeseries, yield\_curves, machine-learning experiments etc. The interface is very succinct and allows the user to concentrate on logic of the calculations rather than boilerplate.

dictable supports quite a flexible construction:

**Example:** construction using records

```
>>> from pyg import *; import pandas as pd
>>> d = dictable([dict(name = 'alan', surname = 'atkins', age = 39, country = 'UK'),
>>>                 dict(name = 'barbara', surname = 'brown', age = 29, country = 'UK')])
```

**Example:** construction using columns and constants

```
>>> d = dictable(name = ['alan', 'barbara'], surname = ['atkins', 'brown'], age = [39, 29])
```

**Example:** construction using pandas.DataFrame

```
>>> original = dictable(name = ['alan', 'barbara'], surname = ['atkins', 'brown'], age = [39, 29])
>>> df_from_dictable = pd.DataFrame(original)
>>> dictable_from_df = dictable(df_from_dictable)
>>> assert original == dictable_from_df
```

**Example:** construction rows and columns

```
>>> d = dictable([['alan', 'atkins', 39, 'UK'], ['barbara', 'brown', 29, 'UK']], columns = ['name', 'surname', 'age', 'country'])
```

**Access:** column access

```
>>> assert d.keys() == ['name', 'surname', 'age', 'country']
>>> assert d.name == ['alan', 'barbara']
>>> assert d['name'] == ['alan', 'barbara']
>>> assert d['name', 'surname'] == [('alan', 'atkins'), ('barbara', 'brown')]
>>> assert d[lambda name, surname: '%s %s'%(name, surname)] == ['alan atkins', 'barbara brown']
```

**Access:** row access & iteration

```
>>> assert d[0] == {'name': 'alan', 'surname': 'atkins', 'age': 39, 'country': 'UK'}
>>> assert [row for row in d] == [{'name': 'alan', 'surname': 'atkins', 'age': 39, 'country': 'UK'},
>>>                                 {'name': 'barbara', 'surname': 'brown', 'age': 29, 'country': 'UK'}]
```

Note that members access is commutative:

```
>>> assert d.name[0] == d[0].name == 'alan'
>>> d[lambda name, surname: name + surname][0] == d[0][lambda name, surname: name + surname]
>>> assert sum([row for row in d], dictable()) == d
```

**Example:** adding records

```
>>> d = dictable(name = ['alan', 'barbara'], surname = ['atkins', 'brown'], age = [39, 29])
>>> d = d + {'name': 'charlie', 'surname': 'chocolate', 'age': 49} # can add a record directly
>>> assert d[-1] == {'name': 'charlie', 'surname': 'chocolate', 'age': 49, 'country': None}
>>> d += dictable(name = ['dana', 'ender'], surname = ['deutch', 'esterhase'], age = [10, 15])
>>> assert d.name == ['alan', 'barbara', 'charlie', 'dana', 'ender']
>>> assert len(dictable.concat([d,d])) == len(d) * 2
```

**Example:** adding columns

```
>>> d = dictable(name = ['alan', 'barbara'], surname = ['atkins', 'brown'], age = [39, 29])

>>> ### all of the below are ways of adding columns ###
>>> d.gender == ['m', 'f']
>>> d = d(gender = ['m', 'f'])
>>> d['gender'] == ['m', 'f']
>>> d2 = dictable(gender = ['m', 'f'], profession = ['astronaut', 'barber'])
>>> d = d(**d2)
```

**Example:** adding derived columns

```
>>> d = dictable(name = ['alan', 'barbara'], surname = ['atkins', 'brown'], age = [39, 29])
>>> d = d(full_name = lambda name, surname: proper('%s %s'%(name, surname)))
>>> d['full_name'] = d[lambda name, surname: proper('%s %s'%(name, surname))]
>>> assert d.full_name == ['Alan Atkins', 'Barbara Brown']
```

**Example:** dropping columns

```
>>> d = dictable(name = ['alan', 'barbara'], surname = ['atkins', 'brown'], age = [39, 29])
>>> del d.country # in place
>>> del d['age'] # in place
>>> assert (d - 'name')[0] == {'surname': 'atkins'} and d[0] == {'name': 'alan', 'surname': 'atkins'}
```

**Example:** row selection, inc/exc

```
>>> d = dictable(name = ['alan', 'barbara'], surname = ['atkins', 'brown'], age = [39, 29])
>>> assert len(d.exc(name = 'alan')) == 1
>>> assert len(d.exc(lambda age: age<30)) == 1 # can filter on *functions* of members, not
>>> assert d.inc(name = 'alan').surname == ['atkins']
>>> assert d.inc(lambda age: age<30).name == ['barbara']
>>> assert d.exc(lambda age: age<30).name == ['alan']
```

**dictable supports:**

- sort
- group-by/ungroup
- list-by/ unlist
- pivot/unpivot
- inner join, outer join and xor

Full details are below.

**classmethod concat** (\*others)

adds together multiple dictables. equivalent to sum(others, self) but a little faster

**Parameters:**

**\*others** : *dictables*

records to be added to current table

**Returns:**

**merged** : *dictable*

sum of all records

**Example:**

```
>>> from pyg import *
>>> d1 = dictable(a = [1,2,3])
>>> d2 = dictable(a = [4,5,6])
>>> d3 = dictable(a = [7,8,9])
```



```
>>> assert dictable.concat(d1,d2,d3) == dictable(a = range(1,10))
>>> assert dictable.concat([d1,d2,d3]) == dictable(a = range(1,10))
```

**do** (function, \*keys)

applies a function(s) on multiple keys at the same time

**Parameters:**

**function** : callable or list of callables

function to be applied per column

**\*keys** : string/list of strings

list of columns to be applied. If missing, applied to all columns

**Returns:**

res : dictable

**Example:**

```
>>> from pyg import *
>>> d = dictable(name = ['adam', 'barbara', 'chris'], surname = ['atkins', 'brown', 'cohen'])
>>> assert d.do(proper) == dictable(name = ['Adam', 'Barbara', 'Chris'], surname = ['Atkins', 'Brown', 'Cohen'])
```

**Example:** using another column in the calculation

```
>>> from pyg import *
>>> d = dictable(a = [1,2,3,4], b = [5,6,9,8], denominator = [10,20,30,40])
>>> d = d.do(lambda value, denominator: value/denominator, 'a', 'b')
>>> assert d == dictable(a = 0.1, b = [0.5,0.3,0.3,0.2], denominator = [10,20,30,40])
```

**exc** (\*functions, \*\*filters)

performs a filter on what rows to exclude

**Parameters:**

**\*functions** : callables or a dict

filters based on functions of each row

**\*\*filters** : value or list of values

filters per each column

**Returns:**

**dictable**

table with rows that satisfy all conditions excluded.

**Example:** filtering on keys

```
>>> from pyg import *; import numpy as np
>>> d = dictable(x = [1,2,3,np.nan], y = [0,4,3,5])
>>> assert d.exc(x = np.nan) == dictable(x = [1,2,3], y = [0,4,3])
>>> assert d.exc(x = 1) == dictable(x = [2,3,np.nan], y = [4,3,5])
>>> assert d.exc(x = [1,2]) == dictable(x = [1,2], y = [0,4])
```

**Example:** filtering on callables

```
>>> from pyg import *; import numpy as np
>>> d = dictable(x = [1,2,3,np.nan], y = [0,4,3,5])
>>> assert d.exc(lambda x,y: x>y) == dictable(x = 1, y = 0)
```

**get** (key, default=None)

Return the value for key if key is in the dictionary, else default.

Welcome to pyg!

**groupby** (\*by, grp='grp')

Similar to pandas groupby but returns a dictable of dictables with a new column 'grp'

**Example:**

```
>>> x = dictable(a = [1,2,3,4], b= [1,0,1,0])
>>> res = x.groupby('b')
>>> assert res.keys() == ['b', 'grp']
>>> assert is_dictable(res[0].grp) and res[0].grp.keys() == ['a']
```

**Parameters:**

**\*by** : str or list of strings

gr.

**grp** : str, optional

The name of the column for the dictables per each key. The default is 'grp'.

**Returns:**

**dictable**

A dictable containing the original keys and a dictable per unique key.

**inc** (\*functions, \*\*filters)

performs a filter on what rows to include

**Parameters:**

**\*functions** : callables or a dict

filters based on functions of each row

**\*\*filters** : value or list of values

filters per each column

**Returns:**

**dictable**

table with rows that satisfy all conditions.

**Example:** filtering on keys

```
>>> from pyg import *; import numpy as np
>>> d = dictable(x = [1,2,3,np.nan], y = [0,4,3,5])
>>> assert d.inc(x = np.nan) == dictable(x = np.nan, y = 5)
>>> assert d.inc(x = 1) == dictable(x = 1, y = 0)
>>> assert d.inc(x = [1,2]) == dictable(x = [1,2], y = [0,4])
```

**Example:** filtering on regex

```
>>> import re
>>> d = dictable(text = ['once', 'upon', 'a', 'time', 'in', 'the', 'west', 1, 2, 3])
>>> assert d.inc(text = re.compile('o')) == dictable(text = ['once', 'upon'])
>>> assert d.exc(text = re.compile('e')) == dictable(text = ['upon', 'a', 'in', 1, 2, 3])
```

**Example:** filtering on callables

```
>>> from pyg import *; import numpy as np
>>> d = dictable(x = [1,2,3,np.nan], y = [0,4,3,5])
>>> assert d.inc(lambda x,y: x>y) == dictable(x = 1, y = 0)
```

**join** (other, lcols=None, rcols=None, mode=None)

Performs either an inner join or a cross join between two dictables

**Example:** inner join

```
>>> from pyg import *
>>> x = dictable(a = ['a','b','c','a'])
>>> y = dictable(a = ['a','y','z'])
>>> assert x.join(y) == dictable(a = ['a', 'a'])
```

**Example:** outer join

```
>>> from pyg import *
>>> x = dictable(a = ['a','b'])
>>> y = dictable(b = ['x','y'])
>>> assert x.join(y) == dictable(a = ['a', 'a', 'b', 'b'], b = ['x', 'y', 'x', 'y'])
```

**pivot (x, y, z, agg=None)**  
pivot table functionality.

**Parameters:**

**x** : *str/list of str*

unique keys per each row

**y** : *str*

unique key per each column

**z** : *str/callable*

A column in the table or an evaluated quantity per each row

**agg** : *None/callable or list of callables, optional*

Each (x,y) cell can potentially contain multiple z values. so if agg = None, a list is returned If you want the data aggregated in any way, then supply an aggregating function(s)

**Returns:**

A dictable which is a pivot table of the original data

**Example:**

```
>>> from pyg import *
>>> timetable_as_list = dictable(x = [1,2,3]) * dictable(y = [1,2,3])
>>> timetable = timetable_as_list.xyz('x','y',lambda x, y: x * y)
>>> assert timetable = dictable(x = [1,2,3], )
```

**Example:**

```
>>> self = dictable(x = [1,2,3]) * dictable(y = [1,2,3])
>>> x = 'x'; y = 'y'; z = lambda x, y: x * y
>>> self.exc(lambda x, y: x+y==5).xyz(x,y,z, len)
```

**sort (\*by)**

Sorts the table either using a key, list of keys or functions of members

**Example:**

```
>>> import numpy as np
>>> self = dictable(a = [_ for _ in 'abracadabra'], b=range(11), c = range(0,33,3))
>>> self.d = list(np.array(self.c) % 11)
>>> res = self.sort('a', 'd')
>>> assert list(res.c) == list(range(11))
```

```
>>> d = dictable(a = ['a', 1, 'c', 0, 'b', 2]).sort('a')
>>> res = d.sort('a','c')
>>> print(res)
>>> assert ''.join(res.a) == 'aaaaabbcdrr' and list(res.c) == [0,4,8,9,10] + [2,3] + [1]
```

Welcome to pyg!

```
>>> d = d.sort(lambda b: b*3 % 11) ## sorting again by c but using a function
>>> assert list(d.c) == list(range(11))
```

**ungroup (grp='grp')**  
Undoes groupby

**Example:**

```
>>> x = dictable(a = [1,2,3,4], b= [1,0,1,0])
>>> self = x.groupby('b')
```

**Parameters:**

**grp** : *str, optional*  
column name where dictables are. The default is 'grp'.

**Returns:**

dictable.

**unlist ()**  
undoes listby...

**Example:**

```
>>> x = dictable(a = [1,2,3,4], b= [1,0,1,0])
>>> x.listby('b')
```

```
dictable[2 x 2] b|a 0|[2, 4] 1|[1, 3]
```

```
>>> assert x.listby('b').unlist().sort('a') == x
```

**Returns:**

**dictable**

a dictable where all rows with list in them have been 'expanded'.

**unpivot (x, y, z)**  
undoes self.xyz / self.pivot

**Example:**

```
>>> from pyg import *
>>> orig = (dictable(x = [1,2,3,4]) * dict(y = [1,2,3,4,5]))(z = lambda x, y: x*y)
>>> pivot = orig.xyz('x', 'y', 'z', last)
>>> unpivot = pivot.unpivot('x','y','z').do(int, 'y') # the conversion to column names me
>>> assert orig == unpivot
```

**Parameters:**

**x** : *str/list of strings*  
list of keys in the pivot table.

**y** : *str*  
name of the columns that wil be used for the values that are currently column headers.

**z** : *str*  
name of the column that describes the data currently within the pivot table.

**Returns:**

dictable

**update ([, E], \*\*F) → None.** Update D from dict/iterable E and F.

If E is present and has a .keys() method, then does: for k in E: D[k] = E[k] If E is present and lacks a .keys() method, then does: for k, v in E: D[k] = v In either case, this is followed by: for k in F: D[k] = F[k]

**xor** (other, lcols=None, rcols=None, mode='l')

returns what is in lhs but NOT in rhs (or vice versa if mode = 'r'). Together with inner joining, can be used as left/right join

**Examples:**

```
>>> from pyg import *
>>> self = dictable(a = [1,2,3,4])
>>> other = dictable(a = [1,2,3,5])
>>> assert self.xor(other) == dictable(a = 4) # this is in lhs but not in rhs
>>> assert self.xor(other, lcols = lambda a: a * 2, rcols = 'a') == dictable(a = [2,3,4])
```

The XOR functionality can be performed using quotient (divide): >>> assert lhs/rhs == dictable(a = 4) >>> assert rhs/lhs == dictable(a = 5)

```
>>> rhs = dictable(a = [1,2], b = [3,4])
>>> left_join_can_be_done_simply_as = lhs * rhs + lhs/rhs
```

**Parameters:**

**other** : dictable (or something that can be turned to one)

what we exclude with.

**lcols** : str/list of strs, optional

the left columns/formulae on which we match. The default is None.

**rcols** : str/list of strs, optional

the right columns/formulae on which we match. The default is None.

**mode** : string, optional

When set to 'r', performs xor the other way. The default is 'l'.

**Returns:**

**dictable**

a dictable containing what is in self but not in ther other dictable.

**xyz** (x, y, z, agg=None)

pivot table functionality.

**Parameters:**

**x** : str/list of str

unique keys per each row

**y** : str

unique key per each column

**z** : str/callable

A column in the table or an evaluated quantity per each row

**agg** : None/callable or list of callables, optional

Each (x,y) cell can potentially contain multiple z values. so if agg = None, a list is returned If you want the data aggregated in any way, then supply an aggregating function(s)

**Returns:**

A dictable which is a pivot table of the original data

**Example:**

```
>>> from pyg import *
>>> timetable_as_list = dictable(x = [1,2,3]) * dictable(y = [1,2,3])
>>> timetable = timetable_as_list.xyz('x','y',lambda x, y: x * y)
>>> assert timetable = dictable(x = [1,2,3], )
```

### Example:

```
>>> self = dictable(x = [1,2,3]) * dictable(y = [1,2,3])
>>> x = 'x'; y = 'y'; z = lambda x, y: x * y
>>> self.exc(lambda x, y: x+y==5).xyz(x,y,z, len)
```

pyg.base.\_dictable.dictable.\_\_call\_\_(self, \*\*kwargs)  
Call self as a function.

## perdictable

pyg.base.\_perdictable.perdictable()  
A decorator that makes a function works per dictable and not just on original value

### Example:

```
>>> f = lambda a, b: a+b
>>> p = perdictable(f, on = ['key'])
```

The new modified function p now works the same on old values:

### Paramaters:

#### function : callable

A function

#### on: str/list of str

perform join based on these keys

#### renames: dict

This tells us which column to grab from which table

#### defaults: dict

If a default is provided for a parameter, we will perform a left join, substituting missing values with the defaults

#### if\_none: bool, list of keys

If historic data is None while the row has expired, should we force a recalculation? if True, will be done.

#### output\_is\_input: bool, list of keys

Some functions want their own output to be presented to them. If you see to True, if cached values exist for these columns, these are provided to the function

#### include\_inputs:

When we return the outputs, do you want the inputs to be included as well in the dictable.

#### col: str

the name of the variable output.

### Example:

```
>>> f = lambda a, b: a+b
>>> p = perdictable(f, include_inputs = True)
>>> assert p(a = 1, b = 2) == 3
>>> assert p(a = dictable(a = [1,2,3]), b = 3) == dictable(a = [1,2,3], b = 3, expiry = No
```

# some parameters are constant, some are tables...

```
>>> assert p(a = 1, b = dictable(key = ['a','b','c'], b = [1,2,3])) == dictable(key = ['a
```

# multiple tables... some unkeyed

```
>>> assert p(a = dictable(a = [1,2]), b = dictable(key = ['a','b','c'], b = [1,2,3])) == d
```

# multiple tables... all keyed

Welcome to pyg!

```
>>> a = dictable(key = ['x', 'y'], data = [1,2])
>>> b = dictable(key = ['y', 'z'], data = [3,4])
>>> assert p(a = a, b = b) == dictable(key = ['y'], data = [5])
```

**Example:** existing data provided using data and expiry

```
>>> a = dictable(key = ['x', 'y', 'z'], data = [1,2,3])
>>> b = dictable(key = ['x', 'y', 'z'], data = [1,3,4])
>>> data = dictable(key = ['x', 'y'], data = ['we calculated this before', 'we calculated this before'])
>>> expiry = dictable(key = ['x', 'y'], data = [dt(2000,1,1), dt(3000,1,1)])
>>> inputs = dict(a = a, b = b)

>>> res = p(a = a, b = b, data = data, expiry = expiry)
>>> assert res.find_data(key = 'x').data == 'we calculated this before'
>>> assert res.find_data(key = 'y').data == 5 # although calculated before, we recalculate
```

## join

`pyg.base._perdictable.join` (inputs, on=None, renames=None, defaults=None)

Suppose we have a function which is defined on simple numbers

**Example:**

```
>>> from pyg import *
>>> profit = lambda amount, price: amount * price
```

The amounts sold are available in one table and prices in another

**Example:**

```
>>> amounts = dictable(product = ['apple', 'orange', 'pear'], amount = [1,2,3])
>>> prices = dictable(product = ['apple', 'orange', 'pear', 'banana'], price = [4,5,6,8])
>>> join(dict(amount = amounts, price = prices), on = 'product')(profit = profit)
```

```
>>> dictable[3 x 4]
>>> product|amount|price|profit
>>> apple  |1      |4      |4
>>> orange |2      |5      |10
>>> pear   |3      |6      |18
```

**Parameters:**

**inputs :** *dict*

a dict of input parameters, some of them may be dictables.

**on :** *str/list of str*

when we have dictables

**renames :** *dict, optional*

remapping. if the datasets contain multiple columns, you can say `renames = dict(price = 'price_in_dollar')` to tell the algo, this is the column to use The default is None.

**defaults :** *dict, optional*

Normally, an inner join is performed. However, if there is a default value/formula for when e.g. a price is missing, use this. The default is None.

**Returns:**

**dictable**

a dictable of an inner join.

**Example:** how column mapping is done

```
>>> on = 'a'
>>> ## if there is only one column apart from keys, then it is selected:
```

```
>>> assert join(dict(x = dictable(a = [1,2], data = [2,3])), on = on) == dictable(a = [1,2])
>>> assert join(dict(x = dictable(a = [1,2], random_name = [2,3])), on = on) == dictable(a = [1,2])
```

```
>>> ## if there are multiple columns, if variable name is there, we use it:
>>> assert join(dict(x = dictable(a = [1,2], z = [2,3], x = [4,5])), on = on) == dictable(a = [1,2], z = [2,3], x = [4,5])
```

```
>>> ## if there are multiple columns, and 'data' is one of the columns, we use it:
>>> assert join(dict(x = dictable(a = [1,2], z = [2,3], data = [4,5])), on = on) == dictable(a = [1,2], z = [2,3], data = [4,5])
```

**Example:** how column mapping is done with rename

```
>>> with pytest.raises(KeyError):
>>>     join(dict(x = dictable(a = [1,2], b = [2,3], c = [4,5])), on = 'a') ## pick b or c
>>> assert join(dict(x = dictable(a = [1,2], b = [2,3], c = [4,5])), on = 'a', renames = {'b': 'x', 'c': 'y'}) == dictable(a = [1,2], b = [2,3], c = [4,5], x = [2,3], y = [4,5])
```

**Example:** joins with partial columns in some tables

```
>>> on = ['a', 'b', 'c']
>>> a = dictable(a = [1,2,3,4], x = [1,2,3,4]) ## only column a here
>>> b = dictable(b = [1,2,3,4], y = [1,2,3,4]) ## only column b here
>>> c = dictable(a = [1,2,3,4], b = [1,2,3,4], c = [1,2,3,4], z = [1,2,3,4])
>>> j = join(dict(x = a, y = b, z = c), on = ['a', 'b', 'c'])
>>> assert len(j) == 4 and sorted(j.keys()) == ['a', 'b', 'c', 'x', 'y', 'z']
```

**Example:** join with defaults

If no defaults are provided, we need all variables to be present. However, if we specify defaults, we left-join on that variable and insert the default value

```
>>> x = dictable(a = [1,2,4], x = [1,2,4])
>>> y = dictable(a = [1,2,3], x = [5,6,7])
>>> on = 'a'
>>> assert join(dict(x = x, y = y), on = on) == dictable(a = [1,2,], x = [1,2], y = [5,6])
>>> assert join(dict(x = x, y = y), on = 'a', defaults = dict(x = None)) == dictable(a = [1,2,3], x = [5,6,7])
>>> assert join(dict(x = x, y = y), on = 'a', defaults = dict(y = 0)) == dictable(a = [1,2,3], x = [5,6,7], y = [0,0,0])
>>> assert join(dict(x = x, y = y), on = 'a', defaults = dict(x = None, y = 0)) == dictable(a = [1,2,3], x = [5,6,7], y = [0,0,0])
```

## named\_dict

`pyg.base._named_dict.named_dict` (name, keys, defaults={}, types={}, casts={}, basedict='pyg.base.dictattr', debug=False)

This forms a base for all classes. It is similar to `named_tuple` but:

- supports additional features such as casting/type checking.
- support default values

The resulting class is a dict so can be stored in MongoDB, sent to json or be used to construct a `pd.Series` automatically.

**Example:** Simple construction

```
>>> Customer = named_dict('Customer', ['name', 'date', 'balance'])
>>> james = Customer('james', 'today', 10)
>>> assert james['balance'] == 10
>>> assert james.date == 'today'
```

**Example:** How `named_dict` works with json/pandas/other `named_dicts`



```
>>> class Customer(named_dict('Customer', ['name', 'date', 'balance'])):
>>>     def add_to_balance(self, value):
>>>         res = self.copy()
>>>         res.balance += value
>>>         return res
```

```
>>> james = Customer('james', 'date', 10)
>>> assert james.add_to_balance(10).balance == 20
>>> import json
>>> assert pd.Series(james).date == 'date'
>>> assert dict(james) == {'name': 'james', 'date': 'date', 'balance': 10}
>>> assert json.dumps(james) == '{"name": "james", "date": "date", "balance": 10}'
```

```
>>> class VIP(named_dict('VIP', ['name', 'date'])):
>>>     def some_method(self):
>>>         return 'inheritance between classes works as long as members can share'
```

```
>>> vip = VIP(james)
>>> assert vip.name == 'james' ## members moved seamlessly
>>> assert vip.some_method() == 'inheritance between classes works as long as members can share'
```

#### Example: Adding defaults

```
>>> Customer = named_dict('Customer', ['name', 'date', 'balance'], defaults = dict(balance=0))
>>> james = Customer('james', 'today')
>>> assert james['balance'] == 0
```

#### Example: types checking

```
>>> import datetime
>>> Customer = named_dict('Customer', ['name', 'date', 'balance'], defaults = dict(balance=0))
>>> james = Customer('james', datetime.datetime.now())
>>> assert james['balance'] == 0
```

#### Example: casting

```
>>> Customer = named_dict('Customer', ['name', 'date', 'balance'], defaults = dict(balance=0))
>>> james = Customer('james', datetime.datetime.now(), balance = '10.3')
>>> assert james['balance'] == 10.3
```

#### Parameters:

**name** : *str*

name of new class.

**keys** : *list*

list of keys that the class must have as members.

**defaults** : *dict, optional*

default values for the keys. The default is {}.

**types** : *type or callable, optional*

A test to be applied for keys either as a callable or as a type. The default is {}.

**casts** : *dict, optional*

function. The default is {}.

**basedict** : *str, optional*

name of the dict class to inherit from. The default is 'dict'.

**debug** : *bool, optional*

output the construction text if set to True. The default is False.

## ValueError

DESCRIPTION.

### Returns:

result : new class that inherits from a dict

## decorators

### wrapper

`class pyg.base._decorators.wrapper (function=None, *args, **kwargs)`

A base class for all decorators. It is similar to `functools.wraps` but better. See below why `wrapt` cannot be used... You basically need to define the wrapped method and everything else is handled for you. - You can then use it either directly to decorate functions - Or use it to create parameterized decorators - the `__name__`, `__wrapped__`, `__doc__` and the `getargspec` will all be taken care of.

### Example:

```
>>> class and_add(wrapper):
>>>     def wrapped(self, *args, **kwargs):
>>>         return self.function(*args, **kwargs) + self.add ## note that we are assuming
```

```
>>> @and_add(add = 3) ## create a decorator and decorate the function
>>> def f(a,b):
>>>     return a+b
```

```
>>> assert f.add == 3
>>> assert f(1,2) == 6
```

Alternatively you can also use it this directly:

```
>>> def f(a,b):
>>>     return a+b
>>>
>>> assert and_add(f, add = 3)(1,2) == 6
```

### Example: Explicit parameter construction

You can make the init more explicit, also adding defaults for the parameters:

```
>>> class and_add_version_2(wrapper):
>>>     def __init__(self, function = None, add = 3):
>>>         super(and_add, self).__init__(function = function, add = add)
>>>     def wrapped(self, *args, **kwargs):
>>>         return self.function(*args, **kwargs) + self.add
```

```
>>> @and_add_version_2
>>> def f(a,b):
>>>     return a+b
>>> assert f(1,2) == 6
```

### Example: No recursion

The decorator is designed to have a single instance of a specific wrapper

```
>>> f = lambda a, b: a+b
>>> assert and_add(and_add(f)) == and_add(f)
```

This holds even for multiple levels of wrapping:

```
>>> x = try_none(and_add(f))
>>> y = try_none(and_add(x))
```

```
>>> assert x == y
>>> assert x(1, 'no can add') is None
```

**Example:** wrapper vs wrapt

wrapt ([wrapt.readthedocs.io](http://wrapt.readthedocs.io)) is an awesome wrapping tool. If you have static library functions, none is better. The problem we face is that wrapt is too good in pretending the wrapped up object is the same as original function:

```
>>> import wrapt
>>> def add_value(value):
>>>     @wrapt.decorator
>>>     def wrapper(wrapped, instance, args, kwargs):
>>>         return wrapped(*args, **kwargs) + value
>>>     return wrapper
```

```
>>> def f(x,y):
>>>     return x*y
```

```
>>> add_three = add_value(value = 3)(f)
>>> add_four = add_value(value = 4)(f)
>>> assert add_four(3,4) == 16 and add_three(3,4) == 15
```

```
>>> ## but here is the problem:
>>> assert encode(add_three) == encode(add_four) == encode(f)
```

So if we ever encode the function and send it across json/Mongo, the wrapping is lost and the user when she receives it cannot use it

```
>>> class add_value(wrapper):
>>>     def wrapped(self, *args, **kwargs):
>>>         return self.function(*args, **kwargs) + self.value
```

```
>>> add_three = add_value(value = 3)(f)
>>> add_four = add_value(value = 4)(f)
>>> encode(add_three)
>>> {'value': 3, 'function': '{"py/function": "__main__.f"}', '_obj': '{"py/type": "__main__"}'}
>>> encode(add_four)
>>> {'value': 4, 'function': '{"py/function": "__main__.f"}', '_obj': '{"py/type": "__main__"}'}
```

## timer

`class pyg.base._decorators.timer` (function, n=1, time=False)

timer is similar to timeit but rather than execution of a Python statement, timer wraps a function to make it log its evaluation time before returning output

**Parameters:**

**function: callable**

The function to be wrapped

**n: int, optional**

Number of times the function is to be evaluated. Default is 1

**time: bool, optional**

If set to True, function will return the TIME it took to evaluate rather than the original function output.

**Example:**

```
>>> from pyg import *; import datetime
>>> f = lambda a, b: a+b
>>> evaluate_100 = timer(f, n = 100, time = True)(1,2)
>>> evaluate_10000 = timer(f, n = 10000, time = True)(1,2)
```

Welcome to pyg!

```
>>> assert evaluate_10000> evaluate_100
>>> assert isinstance(evaluation_time, datetime.timedelta)
```

## ***try\_value***

`pyg.base._decorators.try_value()`

wraps a function to try an evaluation. If an exception is thrown, returns a cached argument

### **Parameters:**

#### **function callable**

The function we want to decorate

#### **value:**

If the function fails, it will return value instead. Default is None

#### **verbose: bool**

If set to True, the logger will warn with the error message.

There are various convenience functions with specific values `try_zero`, `try_false`, `try_true`, `try_nan` and `try_none` will all return specific values if function fails.

### **Example:**

```
>>> from pyg import *
>>> f = lambda a: a[0]
>>> assert try_none(f)(4) is None
>>> assert try_none(f, 'failed')(4) == 'failed'
```

## ***try\_back***

`pyg.base._decorators.try_back()`

wraps a function to try an evaluation. If an exception is thrown, returns first argument

### **Example:**

```
>>> f = lambda a: a[0]
>>> assert try_back(f)('hello') == 'h' and try_back(f)(5) == 5
```

## ***loops***

`class pyg.base._loop.loops (function=None, types=None)`

converts a function to loop over the arguments, depending on the type of the first argument

### **Examples:**

```
>>> @loop(dict, list, pd.DataFrame, pd.Series)
>>> def f(a,b):
>>>     return a+b
```

```
>>> assert f(1,2) == 3
>>> assert f([1,2,3],2) == [3,4,5]
>>> assert f([1,2,3], [4,5,6]) == [5,7,9]
```

```
>>> assert f(dict(x=1,y=2), 3) == dict(x = 4, y = 5)
>>> assert f(dict(x=1,y=2), dict(x = 3, y = 4)) == dict(x = 4, y = 6)
```

```
>>> a = pd.Series(dict(x=1,y=2))
>>> b = dict(x=3,y=4)
>>> assert np.all(f(a,b) == pd.Series(dict(x=4,y=6)))
```

Welcome to pyg!

```
>>> a = pd.DataFrame(dict(x=[1,1],y=[2,2])); a.index = [5,10]
>>> b = dict(x=3,y=4)
>>> res = f(a,b)
>>> assert np.all(res == pd.DataFrame(dict(x=[4,4],y=[6,6]), index = [5,10]))
```

```
>>> a = pd.DataFrame(dict(x=[1,1],y=[2,2])); a.index = [5,10]
>>> res = f(a,[3,4])
>>> assert np.all(res == pd.DataFrame(dict(x=[4,4],y=[6,6]), index = [5,10]))
```

## loop

`pyg.base._loop.loop(*types)`  
returns an instance of `loops(types = types)`

`loop_all` is an instance of `loops` that loops over `dict`, `list`, `tuple`, `np.ndarray` and `pandas.DataFrame/Series`

## kwargs\_support

`pyg.base._decorators.kwargs_support()`  
Extends a function to support **kwargs** inputs

**Example:**

```
>>> from pyg import *
>>> @kwargs_support
>>> def f(a,b):
>>>     return a+b
```

```
>>> assert f(1,2, what_is_this = 3, not_used = 4, ignore_this_too = 5) == 3
```

## graphs & cells

### cell

`class pyg.base._cell.cell(function=None, output=None, **kwargs)`

`cell` is a `Dict` that can be thought of as a node in a calculation graph. The nearest parallel is actually an Excel cell:

- `cell` contains both its function and its output. `cell.output` defines the keys where the output is supposed to be
- `cell` contains reference to all the function outputs
- `cell` contains its locations and the means to manage its own persistency

**Parameters:**

- `function` is the function to be called
- **kwargs** are the function named key value args. NOTE: NO SUPPORT for **args** nor **kwargs** in function
- `output`: where should the function output go?

**Example:** simple construction

```
>>> from pyg import *
>>> c = cell(lambda a, b: a+b, a = 1, b = 2)
>>> assert c.a == 1
>>> c = c.go()
>>> assert c.output == ['data'] and c.data == 3
```

**Example:** make output go to 'value' key

```
>>> c = cell(lambda a, b: a+b, a = 1, b = 2, output = 'value')
>>> assert c.go().value == 3
```

**Example:** multiple outputs by function

```
>>> f = lambda a, b: dict(sum = a+b, prod = a*b)
>>> c = cell(f, a = 1, b = 2, output = ['sum', 'prod'])
>>> c = c.go()
>>> assert c.sum == 3 and c.prod == 2
```

### Methods:

- `cell.run()` returns bool if cell needs to be run
- `cell.go()` calculates the cell and returns the function with `cell.output` keys now populated.
- `cell.load()/cell.save()` interface for self load/save persistence

`copy ()` → a shallow copy of D

`go (go=1, mode=0, **kwargs)`

calculates the cell (if needed). By default, will then run `cell.save()` to save the cell. If you don't want to save the output (perhaps you want to check it first), use `cell._go()`

### Parameters:

**go** : *int, optional*

a parameter that forces calculation. The default is 0. `go = 0`: calculate cell only if `cell.run()` is True `go = 1`: calculate THIS cell regardless. calculate the parents only if their `cell.run()` is True `go = 2`: calculate THIS cell and PARENTS cell regardless, calculate grandparents if `cell.run()` is True etc. `go = -1`: calculate the entire tree again.

**\*\*kwargs** : *parameters*

You can actually allocate the variables to the function at runtime

Note that by default, `cell.go()` will default to `go = 1` and force a calculation on cell while `cell()` is lazy and will default to assuming `go = 0`

### Returns:

**cell**

the cell, calculated

**Example:** different values of go

```
>>> from pyg import *
>>> f = lambda x=None,y=None: max([dt(x), dt(y)])
>>> a = cell(f)()
>>> b = cell(f, x = a)()
>>> c = cell(f, x = b)()
>>> d = cell(f, x = c)()
```

```
>>> e = d.go()
>>> e0 = d.go(0)
>>> e1 = d.go(1)
>>> e2 = d.go(2)
>>> e_1 = d.go(-1)
```

```
>>> assert not d.run() and e.data == d.data
>>> assert e0.data == d.data
>>> assert e1.data > d.data and e1.x.data == d.x.data
>>> assert e2.data > d.data and e2.x.data > d.x.data and e2.x.x.data == d.x.x.data
>>> assert e_1.data > d.data and e_1.x.data > d.x.data and e_1.x.x.data > d.x.x.data
```

**Example:** adding parameters on the run

```
>>> c = cell(lambda a, b: a+b)
>>> d = c(a = 1, b =2)
>>> assert d.data == 3
```

**load (mode=0)**

Loads the cell from the database based on primary keys of cell perhaps. Not implemented for simple cell. see db\_cell

**Returns:**

**cell**

self, updated with values from database.

**register (inputs=None)**

registers the cell so that it calculates when the inputs of the cells are calculated

**\*inputs : str**

list of inputs

**cell**

self.

**run ()**

checks if the cell needs calculation. This depends on the nature of the cell. By default (for cell and db\_cell), if the cell is already calculated so that cell.\_output exists, then returns False. otherwise True

**bool**

run cell?

**Example:**

```
>>> c = cell(lambda x: x+1, x = 1)
>>> assert c.run()
>>> c = c()
>>> assert c.data == 2 and not c.run()
```

**save ()**

Saves the cell for persistency. Not implemented for simple cell. see db\_cell

**Returns:**

**cell**

self, saved.

## **cell\_go**

pyg.base.\_cell.**cell\_go** (value, go=0, mode=0)

cell\_go makes a cell run (using cell.go(go)) and returns the calculated cell. If value is not a cell, value is returned.

**Parameters:**

**value : cell**

The cell (or anything else).

**go : int**

same inputs as per cell.go(go). 0: run if cell.run() is True 1: run this cell regardless, run parent cells only if they need to calculate too n: run this cell & its nth parents regardless.

**Returns:**

The calculated cell

**Example:** calling non-cells

Welcome to pyg!

```
>>> assert cell_go(1) == 1
>>> assert cell_go(dict(a=1,b=2)) == dict(a=1,b=2)
```

**Example:** calling cells

```
>>> c = cell(lambda a, b: a+b, a = 1, b = 2)
>>> assert cell_go(c) == c(data = 3)
```

## cell\_item

pyg.base.\_cell.**cell\_item**(value, key=None)

returns an item from a cell (if not cell, returns back the value). If no key is provided, will return the output of the cell

**Parameters:**

**value :** *cell or object or list of cells/objects*

cell

**key :** *str, optional*

The key within cell we are interested in. Note that key is treated as GUIDANCE only. Our strong preference is to return valid output from cell\_output(cell)

**Example:** non cells

```
>>> assert cell_item(1) == 1
>>> assert cell_item(dict(a=1,b=2)) == dict(a=1,b=2)
```

**Example:** cells, simple

```
>>> c = cell(lambda a, b: a+b, a = 1, b = 2)
>>> assert cell_item(c) is None
>>> assert cell_item(c.go()) == 3
```

## cell\_func

pyg.base.\_cell.**cell\_func** ()

cell\_func is a wrapper and wraps a function to act on cells rather than just on values When called, it will returns not just the function, but also args, kwargs used to call it.

**In order to present the itemized value in the cell, inputs for the function that are cells will:**

1. be loaded (from the persistency layer)
2. called and calculated
3. itemized: i.e. cell\_item(input)

**Example:**

```
>>> from pyg import *
>>> a = cell(lambda x: x**2, x = 3)
>>> b = cell(lambda y: y**3, y = 2)
>>> function = lambda a, b: a+b
>>> self = cell_func(function)
>>> result, args, kwargs = self(a,b)
```

```
>>> assert result == 8 + 9
>>> assert kwargs['a'].data == 3 ** 2
>>> assert kwargs['b'].data == 2 ** 3
```

**Parameters:**

**function :** *callable*



Welcome to pyg!

The function to be wrapped

**relabels** : *dict or None*

Allows a redirect of variable names. For example:

```
>>> from pyg import *
>>> a = cell(a = 1, b = 2, c = 3)
>>> f = cell_func(lambda x: x+1, relabels = dict(x = 'c')) ## please use key 'c' to grab x value
>>> assert f(a)[0] == 4
```

**unloaded**: list or str

By defaults, if a cell is in the inputs, it will be loaded (`cell.load()`) prior to being presented to the function. If an arg is in unloaded, it will not be loaded.

**uncalled**: list or str

By defaults, if a cell is in the inputs, it will be called (`cell.call()`) prior to being presented to the function. If an arg is in uncalled, it will not be called.

**unitemized**: list or str

If a cell is in the inputs, once run, we itemize and grab its data. If an arg is in unitemized, it will be presented 'as is'. Example:

```
from pyg import *
cell(passthru, data = 'this is the value presented')
assert cell_func(lambda x: len(x))(x)[0] == len('this is the value presented') ## function run on item
assert cell_func(lambda x: len(x), unitemized = 'x')(x)[0] == len(x) ## function run on x as-is
```

## ***cell\_clear***

`pyg.base._cell.cell_clear` (value)

`cell_clear` clears a cell of its output so that it contains only the essential stuff to do its calculations. This will be used when we save the cell or we want to recalculate it.

**Example:**

```
>>> from pyg import *
>>> a = cell(add_, a = 1, b = 2)
>>> b = cell(add_, a = 2, b = 3)
>>> c = cell(add_, a = a, b = b)()
>>> assert c.data == 8
>>> assert c.a.data == 3
```

```
>>> bare = cell_clear(c)
>>> assert 'data' not in bare and 'data' not in bare.a
>>> assert bare() == c
```

**Parameters:**

**value**: obj

cell (or list/dict of) to be cleared of output

## ***encode and decode/save and load***

### ***encode***

`pyg.base._encode.encode` (value)

encode/decode are performed prior to sending to mongodb or after retrieval from db. The idea is to make object embedding in Mongo transparent to the user.

- We use `jsonpickle` package to embed general objects. These are encoded as strings and can be decoded as long as the original library exists when decoding.

- pandas.DataFrame are encoded to bytes using pickle while numpy arrays are encoded using the faster `array.tobytes()` with arrays' shape & type exposed and searchable.

**Example:**

```
>>> from pyg import *; import numpy as np
>>> value = Dict(a=1,b=2)
>>> assert encode(value) == {'a': 1, 'b': 2, '_obj': '{"py/type": "pyg.base._dict.Dict"}'}
>>> assert decode({'a': 1, 'b': 2, '_obj': '{"py/type": "pyg.base._dict.Dict"}'}) == Dict(a=1,b=2)
>>> value = dictable(a=[1,2,3], b = 4)
>>> assert encode(value) == {'a': [1, 2, 3], 'b': [4, 4, 4], '_obj': '{"py/type": "pyg.base._dict.Dict"}'}
>>> assert decode(encode(value)) == value
>>> assert encode(np.array([1,2])) == {'data': bytes([1, 2]), 'shape': (2,), 'dtype': 'numpy.dtype(int64)', '_obj': '{"py/function": "pyg.base._encode.numpy2py"}'}
```

**Example:** functions and objects

```
>>> from pyg import *; import numpy as np
>>> assert encode(ewma) == '{"py/function": "pyg.timeseries._ewm.ewma"}'
>>> assert encode(Calendar) == '{"py/type": "pyg.base._drange.Calendar"}'
```

**Parameters:**

**value :** *obj*

An object to be encoded

**Returns:**

A pre-json object

## decode

`pyg.base._encode.decode` (value, date=None)  
decodes a string or an object dict

**Parameters:**

**value :** *str or dict*

usually a json

**date :** *None, bool or a regex expression, optional*

date format to be decoded

**Returns:**

**obj**

the json decoded.

**Examples:**

```
>>> from pyg import *
>>> class temp(dict):
>>>     pass
```

```
>>> orig = temp(a = 1, b = dt(0))
>>> encoded = encode(orig)
>>> assert eq(decode(encoded), orig) # type matching too...
```

## pd\_to\_parquet

`pyg.base._parquet.pd_to_parquet` (value, path, compression='GZIP')  
a small utility to save df to parquet, extending both pd.Series and non-string columns

### Example:

```
>>> from pyg import *
>>> import pandas as pd
>>> import pytest
```

```
>>> df = pd.DataFrame([[1,2],[3,4]], drange(-1), columns = [0, dt(0)])
>>> s = pd.Series([1,2,3], drange(-2))
```

```
>>> with pytest.raises(ValueError): ## must have string column names
    df.to_parquet('c:/temp/test.parquet')
```

```
>>> with pytest.raises(AttributeError): ## pd.Series has no to_parquet
    s.to_parquet('c:/temp/test.parquet')
```

```
>>> df_path = pd_to_parquet(df, 'c:/temp/df.parquet')
>>> series_path = pd_to_parquet(s, 'c:/temp/series.parquet')
```

```
>>> df2 = pd_read_parquet(df_path)
>>> s2 = pd_read_parquet(series_path)
```

```
>>> assert eq(df, df2)
>>> assert eq(s, s2)
```

## ***pd\_read\_parquet***

`pyg.base._parquet.pd_read_parquet` (path)

a small utility to read df/series from parquet, extending both `pd.Series` and non-string columns

### Example:

```
>>> from pyg import *
>>> import pandas as pd
>>> import pytest
```

```
>>> df = pd.DataFrame([[1,2],[3,4]], drange(-1), columns = [0, dt(0)])
>>> s = pd.Series([1,2,3], drange(-2))
```

```
>>> with pytest.raises(ValueError): ## must have string column names
    df.to_parquet('c:/temp/test.parquet')
```

```
>>> with pytest.raises(AttributeError): ## pd.Series has no to_parquet
    s.to_parquet('c:/temp/test.parquet')
```

```
>>> df_path = pd_to_parquet(df, 'c:/temp/df.parquet')
>>> series_path = pd_to_parquet(s, 'c:/temp/series.parquet')
```

```
>>> df2 = pd_read_parquet(df_path)
>>> s2 = pd_read_parquet(series_path)
```

```
>>> assert eq(df, df2)
>>> assert eq(s, s2)
```

## ***parquet\_encode***

`pyg.mongo._encoders.parquet_encode` (value, path, compression='GZIP')

encodes a single DataFrame or a document containing dataframes into a an object that can be decoded

```
>>> from pyg import *
>>> path = 'c:/temp'
>>> value = dict(key = 'a', n = np.random.normal(0,1, 10), data = dictable(a = [pd.Series(
>>> encoded = parquet_encode(value, path)
>>> assert encoded['n']['file'] == 'c:/temp/n.npy'
>>> assert encoded['data'].a[0]['path'] == 'c:/temp/data/a/0.parquet'
>>> assert encoded['other']['df']['path'] == 'c:/temp/other/df.parquet'

>>> decoded = decode(encoded)
>>> assert eq(decoded, value)
```

## csv\_encode

`pyg.mongo._encoders.csv_encode` (value, path)

encodes a single DataFrame or a document containing dataframes into a an object that can be decoded while saving dataframes into csv

```
>>> path = 'c:/temp'
>>> value = dict(key = 'a', data = dictable(a = [pd.Series([1,2,3]), pd.Series([4,5,6])],
>>> encoded = csv_encode(value, path)
>>> assert encoded['data'].a[0]['path'] == 'c:/temp/data/a/0.csv'
>>> assert encoded['other']['df']['path'] == 'c:/temp/other/df.csv'

>>> decoded = decode(encoded)
>>> assert eq(decoded, value)
```

## convertors to bytes

`pyg.base._encode.pd2bson` (value)

converts a value (usually a pandas.DataFrame/Series) to bytes using pickle

`pyg.base._encode.np2bson` (value)

converts a numpy array to bytes using `value.tobytes()`. This is much faster than pickle but does not save shape/type info which we save separately.

`pyg.base._encode.bson2np` (data, dtype, shape)

converts a byte with dtype and shape information into a numpy array.

`pyg.base._encode.bson2pd` (data)

converts a pickled object back to an object. We insist that new object has `.shape` to ensure we did not unpickle gibberish.

## dates and calendar

### dt

`pyg.base._dates.dt` (\*args, dialect='uk', none=<built-in method now of type object>)

A more generic constructor for `datetime.datetime`.

**Example:** Simple construction

```
>>> assert dt(2000,1,1) == datetime.datetime(2000, 1, 1, 0, 0) # name of month
>>> assert dt(2000,'jan',1) == datetime.datetime(2000, 1, 1, 0, 0) # name of month
>>> assert dt(2000,'f',1) == datetime.datetime(2000, 1, 1, 0, 0) # future month code
>>> assert dt('01-02-2002') == datetime.datetime(2002, 2, 1)
>>> assert dt('01-02-2002', dialect = 'US') == datetime.datetime(2002, 1, 2)
>>> assert dt('01 March 2002') == datetime.datetime(2002, 3, 1)
>>> assert dt('01 March 2002', dialect = 'US') == datetime.datetime(2002, 3, 1)
>>> assert dt('01 March 2002 10:20:30') == datetime.datetime(2002, 3, 1, 10, 20, 30)
```

```
>>> assert dt(20020301) == datetime.datetime(2002, 3, 1)
>>> assert dt(37316) == datetime.datetime(2002, 3, 1) # excel date
>>> assert dt(730180) == datetime.datetime(2000,3,1) # ordinal for 1/3/2000
>>> assert dt(2000,3,1).timestamp() == 951868800.0
>>> assert dt(951868800.0) == datetime.datetime(2000,3,1) # utc timestamp
>>> assert dt(np.datetime64(dt(2000,3,1))) == dt(2000,3,1) ## numpy.datetime64 object
```

```
>>> assert dt(2000) == datetime.datetime(2000,1,1)
>>> assert dt(2000,3) == datetime.datetime(2000,3,1)
>>> assert dt(2000,3, 1) == datetime.datetime(2000,3,1)
>>> assert dt(2000,3, 1, 10,20,30) == datetime.datetime(2000,3,1,10,20,30)
>>> assert dt(2000,'march', 1) == datetime.datetime(2000,3,1)
>>> assert dt(2000,'h', 1) == datetime.datetime(2000,3,1) # future codes
```

**Example:** date as offset from today

```
>>> today = dt(0);
>>> import datetime
>>> day = datetime.timedelta(1)
>>> assert dt(-3) == today - 3 * day
>>> assert dt('-10b') == today - 14 * day
```

**Example:** datetime arithmetic:

dt has an interesting logic in implementing datetime arithmetic:

- day and month parameters can be negative or bigger than the days of month
- dt() will roll back/forward from the date which is valid

```
>>> assert dt(2000,4,1) == datetime.datetime(2000, 4, 1, 0, 0)
>>> assert dt(2000,4,0) == datetime.datetime(2000, 3, 31, 0, 0) # a day before dt(2000,4,1)
```

and rolling back months:

```
>>> assert dt(2000,0,1) == datetime.datetime(1999, 12, 1, 0, 0) # a month before dt(2000,1,1)
>>> assert dt(2000,13,1) == datetime.datetime(2001, 1, 1, 0, 0) # a month after dt(2000,12,1)
```

This may feel unnatural at first, but does allow for much nicer code, e.g.: [dt(2000,i,1) for i in range(-10,10)]

**Parameters:**

**\*args :** *str, int or dates*

argument to be converted into dates

**dialect :** *str, optional*

parsing of 01/02/2020 is it 1st Feb or 2nd Jan? The default is 'uk', i.e. dd/mm/yyyy

**none :** *callable, optional*

What is dt()? The default is datetime.datetime.now()

## ymd

pyg.base.\_dates.**ymd** (\*args, dialect='uk', none=<built-in method now of type object>)  
just like dt() but always returns date only (year/month/date) without fractions. see dt() for full documentation  
datetime.datetime

## dt\_bump

pyg.base.\_dates.**dt\_bump** (t, \*bumps)

**Example:**

Welcome to pyg!

```
>>> from pyg import *
>>> t = pd.Series([1,2,3], drange(dt(2000,1,1),2))
>>> assert eq(dt_bump(t, 1), pd.Series([1,2,3], drange(dt(2000,1,2),2)))
```

## drange

pyg.base.\_drange.**drange** (t0=None, t1=None, bump=None)

A quick and happy wrapper for dateutil.rrule

### Examples:

```
>>> drange(2000, 10, 1) # 10 days starting from dt(2000,1,1)
>>> drange(2000, '10b', '1b') # weekdays between dt(2000,1,1) and dt(2000,1,17)
>>> drange('-10b', 0, '1b') # business days since 10 bdays ago
>>> drange('-10b', '10b', '1w') # starting 10b days ago, to 10b from now, counting in week
```

### Parameters:

**t0** : *date, optional*

start date. The default is None.

**t1** : *date, optional*

end date. The default is None.

**bump** : *timedelta, int, string, optional*

bump period. The default is None.

### Returns:

list of dates

### Example:

```
>>> t0 = 2000; t1 = 1999
>>> bump = '-1b'
```

### Example:

```
>>> t0 = dt(2020); t1 = dt(2021); bump = datetime.timedelta(hours = 4)
```

## date\_range

pyg.base.\_drange.**date\_range** (t0=None, t1=None)

## Calendar

`class pyg.base._drange.Calendar` (key=None, holidays=None, weekend=None, t0=None, t1=None, adj='m')

### Calendar is

- a dict
- containing holiday dates
- implementing business day arithmetic

Calendar is restricted to operate between cal.t0 and cal.t1 which default to TMIN = 1900 and TMAX = 2300

### Calendar does this by having two key members:

- dt2int: a mapping from all business dates to their integer 'clock'
- int2dt: a mapping from integer value to the date

Since Calendar is an 'expensive' memory wise, we assign a key to the calendar and the Calendar is stored in the singleton calendars under this key

### Example:

```
>>> from pyg import *
>>> holidays = dictable([[1,'2012-01-02','New Year Day'],,
                        [2,'2012-01-16','Martin Luther King Jr. Day'],,
                        [3,'2012-02-20','Presidents Day (Washingtons Birthday)'],,
                        [4,'2012-05-28','Memorial Day'],,
                        [5,'2012-07-04','Independence Day'],,
                        [6,'2012-09-03','Labor Day'],,
                        [7,'2012-10-08','Columbus Day'],,
                        [8,'2012-11-12','Veterans Day'],,
                        [9,'2012-11-22','Thanksgiving Day'],,
                        [10,'2012-12-25','Christmas Day'],,
                        [11,'2013-01-01','New Year Day'],,
                        [12,'2013-01-21','Martin Luther King Jr. Day'],,
                        [13,'2013-02-18','Presidents Day (Washingtons Birthday)'],,
                        [14,'2013-05-27','Memorial Day'],,
                        [15,'2013-07-04','Independence Day'],,
                        [16,'2013-09-02','Labor Day'],,
                        [17,'2013-10-14','Columbus Day'],,
                        [18,'2013-11-11','Veterans Day'],,
                        [19,'2013-11-28','Thanksgiving Day'],,
                        [20,'2013-12-25','Christmas Day'],,
                        [21,'2014-01-01','New Year Day'],,
                        [22,'2014-01-20','Martin Luther King Jr. Day'],,
                        [23,'2014-02-17','Presidents Day (Washingtons Birthday)'],,
                        [24,'2014-05-26','Memorial Day'],,
                        [25,'2014-07-04','Independence Day'],,
                        [26,'2014-09-01','Labor Day'],,
                        [27,'2014-10-13','Columbus Day'],,
                        [28,'2014-11-11','Veterans Day'],,
                        [29,'2014-11-27','Thanksgiving Day'],,], ['i', 'date', 'name']).do
```

```
>>> cal = calendar('US', holidays.date, t0 = 2012, t1 = 2015)
>>> assert not cal.is_bday(dt(2013,9,2))          # Labor day
```

```
>>> cached_calendar = calendar('US')
>>> assert not cached_calendar.is_bday(dt(2013,9,2))  # Labor day
```

```
>>> assert cal.adjust(dt(2013,9,2)) == dt(2013,9,3)
>>> assert cal.drangle(dt(2013,9,0), dt(2013,9,7), '1b') == [dt(2013,8,30), dt(2013,9,3), d
```

```
>>> assert cal.bdays(dt(2013,9,0), dt(2013,9,7)) == 5
```

### **add** (date, days, adj=None)

adjusts the start date to a business day. Then add business days on top. add will initiate self.\_populate unless we are bumping date by exactly one day

**date** : *datetime*

start date.

**days** : *int*

days to bump.

**adj** : *str, optional*

adjusting method for date if it isn't a business day to start with.

**datetime**

bumped date.

### **adjust** (date, adj=None)

Welcome to pyg!

adjust a non-business day to prev/following bussiness date

**Parameters:**

**date** : *datetime*. **adj** : None or p/f/m

adjustment convention: 'prev/following/modified following'

**Returns:**

**datetime**

nearby business day

**dt\_bump** (t, bump, adj=None)

adds a bump to a date

**Parameters:**

**t** : *datetime*

date to bump.

**bump** : *int, str*

bump e.g. '-1y' or '1b' or 3

**adj** : *adjustment type*

The default is None.

**Returns:**

**datetime**

bumped date.

**is\_trading** (date=None, day\_start=0, day\_end=235959)

calculates if we are within a trading session

**Parameters:**

**date** : *datetime, optional*

the time & date we want to check. The default is None (i.e. now)

**Returns:**

**bool:**

are we within a trading session

**trade\_date** (date=None, adj=None, day\_start=0, day\_end=235959)

This is very similar for adjust, but it also takes into account the time of the day. if day\_start = 0 and day\_end = 23:59:59 then this is exactly adjust.

**Parameters:**

**date** : *datetime, optional*

date (with time). The default is None.

**adj** : *f/p, optional*

If date isn't within trading day, which direction to adjust to? The default is None.

**Example:**

```
>>> from pyg import *; import datetime
```

```
>>> uk = calendar('UK', day_start = 8, day_end = 17)
>>> assert uk.trade_date(dt(2021,2,9,5), 'f') == dt(2021, 2, 9) # Tuesday morning rolls
>>> assert uk.trade_date(dt(2021,2,9,5), 'p') == dt(2021, 2, 8) # Tuesday morning back
>>> assert uk.trade_date(dt(2021,2,7,5), 'f') == dt(2021, 2, 8) # Sunday rolls into Monday
>>> assert uk.trade_date(dt(2021,2,7,5), 'p') == dt(2021, 2, 5) # Sunday rolls back to Friday
```



```
>>> assert uk.trade_date(date = dt(2021,2,9,23), adj = 'f') == dt(2021, 2, 10) # Tuesday
>>> assert uk.trade_date(date = dt(2021,2,9,23), adj = 'p') == dt(2021, 2, 9) # Tuesday
>>> assert uk.trade_date(date = dt(2021,2,7,23), adj = 'f') == dt(2021, 2, 8) # Sunday
>>> assert uk.trade_date(date = dt(2021,2,7,23), adj = 'p') == dt(2021, 2, 5) # Sunday
```

```
>>> assert uk.trade_date(date = dt(2021,2,9,12), adj = 'f') == dt(2021, 2, 9) # Tuesday
>>> assert uk.trade_date(date = dt(2021,2,9,12), adj = 'p') == dt(2021, 2, 9) # Tuesday
```

```
>>> au = calendar('AU', day_start = 2230, day_end = 1300)
>>> assert au.trade_date(dt(2021,2,9,5), 'f') == dt(2021, 2, 9) # Tuesday morning in ses
>>> assert au.trade_date(dt(2021,2,9,5), 'p') == dt(2021, 2, 9) # Tuesday morning in ses
>>> assert au.trade_date(dt(2021,2,7,5), 'f') == dt(2021, 2, 8) # Sunday rolls into Mond
>>> assert au.trade_date(dt(2021,2,7,5), 'p') == dt(2021, 2, 5) # Sunday rolls back to F
```

```
>>> assert au.trade_date(date = dt(2021,2,9,23), adj = 'f') == dt(2021, 2, 10) # Tuesday
>>> assert au.trade_date(date = dt(2021,2,9,23), adj = 'p') == dt(2021, 2, 10) # Already
>>> assert au.trade_date(date = dt(2021,2,7,23), adj = 'f') == dt(2021, 2, 8) # Sunday
>>> assert au.trade_date(date = dt(2021,2,7,23), adj = 'p') == dt(2021, 2, 8) # Already
>>> assert au.trade_date(date = dt(2021,2,5,23), adj = 'f') == dt(2021, 2, 8) # Friday a
```

```
>>> assert au.trade_date(date = dt(2021,2,9,14), adj = 'f') == dt(2021, 2, 10) # Tuesday
>>> assert au.trade_date(date = dt(2021,2,9,14), adj = 'p') == dt(2021, 2, 9) # roll bac
```

## calendar

pyg.base.\_drange.**calendar** (key=None, holidays=None, weekend=None, t0=None, t1=None)

A function to returns either an existing calendar or construct a new one. - calendar('US') will return a US calendar if that is already cached - calendar('US', us\_holiday\_dates) will construct a calendar with holiday dates and then cache it

## as\_time

pyg.base.\_drange.**as\_time** (t=None)

parses t into a datetime.time object

### Example:

```
>>> assert as_time('10:30:40') == datetime.time(10, 30, 40)
>>> assert as_time('103040') == datetime.time(10, 30, 40)
>>> assert as_time('10:30') == datetime.time(10, 30)
>>> assert as_time('1030') == datetime.time(10, 30)
>>> assert as_time('05') == datetime.time(5)
>>> assert as_time(103040) == datetime.time(10, 30, 40)
>>> assert as_time(13040) == datetime.time(1, 30, 40)
>>> assert as_time(130) == datetime.time(1, 30)
>>> assert as_time(datetime.time(1, 30)) == datetime.time(1, 30)
>>> assert as_time(datetime.datetime(2000, 1, 1, 1, 30)) == datetime.time(1, 30)
```

### Parameters:

**t** : str/int/datetime.time/datetime.datetime  
time of day

### Returns:

datetime.time

## clock

Welcome to pyg!

`pyg.base._drange.clock` (ts, time=None, t=None)  
returns a vector marking the passage of time.

**Parameters:**

ts : timeseries time : None, a string or a Calendar, or already a timeseries of times

None: Will increment by 1 every non-nan observation 'i' : increment by 1 every date in index (nan or not) 'b' : weekdays distance 'd' : day-distance (ignore intraday stamp) 'f' : fraction-of-day-distance (do not ignore intraday stamp) 'm' : month-distance 'q' : quarter-distance 'y' : year-distance calendar: uses the business-days distance between any two dates

t: starting time in the past.

**Returns:**

an array

an increasing array of time such that distance between points match the above.

**Example:**

```
>>> from pyg import *
>>> assert eq(clock(pd.Series(np.arange(10), drange(2000, 9))), np.arange(1,11))
>>> assert eq(clock(pd.Series(np.arange(10), drange(2000, 9)), t = 5), np.arange(6,16))
>>> assert eq(clock(pd.Series(np.arange(10), drange(2000, 9)), 'i'), np.arange(1,11))
>>> assert eq(clock(pd.Series(np.arange(10), drange(2000, 9)), 'b'), np.array([26090, 2609
>>> assert eq(clock(pd.Series(np.arange(10), drange(2000, '9b', '1b')), 'b'), np.arange(26
```

```
>>> assert eq(clock(np.arange(10)), np.arange(1,11))
>>> assert eq(clock(pd.Series(np.arange(10)), t = 5), np.arange(6,16))
>>> assert eq(clock(np.arange(10), 'i'), np.arange(1,11))
```

## text manipulation

### lower

`pyg.base._txt.lower` (value)

equivalent to `txt.lower()` but:

- does not throw on non-string
- supports lists/dicts

**Example:**

```
>>> assert lower(['The Brown Fox',1]) == ['the brown fox',1]
>>> assert lower(dict(a = 'The Brown Fox', b = 3.0)) == {'a': 'the brown fox', 'b': 3.0}
```

### upper

`pyg.base._txt.upper` (value)

equivalent to `txt.upper()` but:

- does not throw on non-string
- supports lists/dicts

**Example:**

```
>>> assert upper(['The Brown Fox',1]) == ['THE BROWN FOX',1]
>>> assert upper(dict(a = 'The Brown Fox', b = 3.0)) == {'a': 'THE BROWN FOX', 'b': 3.0}
```

### proper

Welcome to pyg!

`pyg.base._txt.proper (value)`

**equivalent to Excel's PROPER(txt) but:**

- does not throw on non-string
- supports lists/dicts

**Example:**

```
>>> assert proper(['THE BROWN FOX',1]) == ['The Brown Fox',1]
>>> assert proper(dict(a = 'THE BROWN FOX', b = 3.0)) == {'a': 'The Brown Fox', 'b': 3.0}
```

## **capitalize**

`pyg.base._txt.capitalize (value)`

**equivalent to `text.capitalize()` but:**

- does not throw on non-string
- supports lists/dicts

**Example:**

```
>>> assert capitalize('alan howard') == 'Alan howard' # use proper to get Alan Howard
>>> assert capitalize(['alan howard', 'donald trump']) == ['Alan howard', 'Donald trump']
```

## **strip**

`pyg.base._txt.strip (value)`

**equivalent to `txt.strip()` but:**

- does not throw on non-string
- supports lists/dicts

**Example:**

```
>>> assert strip([' whatever you say ', ' whatever you do.. ']) == ['whatever you say', 'whatever you do..']
>>> assert strip(dict(a = ' whatever you say ', b = 3.0)) == {'a': 'whatever you say', 'b': 3.0}
```

## **split**

`pyg.base._txt.split (text, sep=' ', dedup=False)`

**equivalent to `txt.split(sep)` but supports:**

- does not throw on non-string
- removal of multiple seps
- ensuring there is a unique single separator

**Parameters:**

**text** : *str*

text to be stripped.

**sep** : *str, list of str, optional*

text used to strip. The default is ' '.

**dedup** : *bool, optional*

If True, will remove duplicated instances of seps. The default is False.

**Returns:**

**str**

Welcome to pyg!

splitted text

**Example:**

```
>>> text = '    The quick... brown .. fox... '
>>> assert split(text) == ['', '', '', 'The', 'quick...', 'brown', '..', 'fox...', '']
>>> assert split(text, [' ', '.'], True) == ['The', 'quick', 'brown', 'fox']
>>> text = dict(a = 'Can split this', b = '..and split this too')
>>> assert split(text, [' ', '.'], True) == {'a': ['Can', 'split', 'this'], 'b': ['and', '']}
```

## replace

`pyg.base._txt.replace` (text, old, new=None)

A souped up version of `text.replace(old, new)`

**Example:** `replace` continues to replace until no-more is found

```
>>> assert replace('this    has lots of double    spaces', ' '*2, ' ') == 'this has lot
>>> assert replace('this, sentence? has! too, many, punctuations!', list(',?!.')) == 'this
>>> assert replace(dict(a = 1, b = [' text within a list ', 'and within a dict']), ' ') ==
```

## common\_prefix

`pyg.base._txt.common_prefix` (\*values)

**Parameters:**

**\*values** : *list of iterables*

values for which we want to find common prefix

**Returns:**

**iterable**

the common prefix.

**Example:**

```
>>> assert common_prefix(['abra', 'abba', 'abacus']) == 'ab'
>>> assert common_prefix('abra', 'abba', 'abacus') == 'ab'
>>> assert common_prefix() is None
>>> assert common_prefix([1,2,3,4], [1,2,3,5,8]) == [1,2,3]
```

## files & directory

### mkdir

`pyg.base._file.mkdir` (path)

makes a new directory if not exists. It works if path is a filename too.

### read\_csv

`pyg.base._file.read_csv` (path, errors='replace', \*\*fmt)

A light-weight csv reader, no conversion is done, nor do we insist equal number of columns per row. - by default, encoding error (unicode characters) are replaced. - fmt parameters are parameters for the `csv.reader` object, see <https://docs.python.org/3/library/csv.html>

## tree manipulation

Trees are dicts of dicts. just like an item in a dict is (key, value), tree items are just longer tuples: (key1, key2, key3, value) We deliberately avoid creating a tree class so that the functionality is available on ordinary tree-like structures.

## ***tree\_keys***

`pyg.base._dict.tree_keys` (tree, types=None)  
returns the keys (branches) of a tree as a list of tuples

### **Example:**

```
>>> tree = dict(a = 1, b = dict(c = 2, d = 3, e = dict(f = 4)))
>>> assert tree_keys(tree) == [('a',), ('b', 'c'), ('b', 'd'), ('b', 'e', 'f')]
```

### **Parameters:**

**tree** : tree (dict of dicts) **types** : types of dicts, optional

## ***tree\_values***

`pyg.base._dict.tree_values` (tree, types=None)  
returns the values (leaf) of a tree (a collection of tuples)

### **Example:**

```
>>> tree = dict(a = 1, b = dict(c = 2, d = 3, e = dict(f = 4)))
>>> assert tree_values(tree) == [1,2,3,4]
```

### **Parameters:**

**tree** : tree (dict of dicts) **types** : types of dicts, optional

## ***tree\_items***

`pyg.base._dict.tree_items` (tree, types=None)  
An extension of `dict.items()`, returning a list of tuples but of varying length, each a branch of a tree

### **Parameters:**

**tree** : *dict of dicts*

a tree of data.

**types** : *dict or a list of dict-types, optional*

The types that we consider as 'branches' of the tree. Default is (dict, Dict, dictattr).

### **Returns:**

#### **a list of tuples**

these are an extension of `dict.items()` and are of varying length

### **Example:**

```
>>> school = dict(pupils = dict(id1 = dict(name = 'james', surname = 'maxwell', gender = 'm'),
                                id2 = dict(name = 'adam', surname = 'smith', gender = 'm'),
                                id3 = dict(name = 'michell', surname = 'obama', gender = 'f'),
                                ),
                  teachers = dict(math = dict(name = 'albert', surname = 'einstein', grade = 3),
                                  english = dict(name = 'william', surname = 'shakespeare', grade = 3),
                                  physics = dict(name = 'richard', surname = 'feyman', grade = 4)
                                  ))
```

```
>>> items = tree_items(school)
>>> items
```

```
>>> [('pupils', 'id1', 'name', 'james'),
>>>  ('pupils', 'id1', 'surname', 'maxwell'),
>>>  ('pupils', 'id1', 'gender', 'm'),
>>>  ('pupils', 'id2', 'name', 'adam'),
```

```
>>> ('pupils', 'id2', 'surname', 'smith'),
>>> ('pupils', 'id2', 'gender', 'm'),
>>> ('pupils', 'id3', 'name', 'michell'),
>>> ('pupils', 'id3', 'surname', 'obama'),
>>> ('pupils', 'id3', 'gender', 'f'),
>>> ('teachers', 'math', 'name', 'albert'),
>>> ('teachers', 'math', 'surname', 'einstein'),
>>> ('teachers', 'math', 'grade', 3),
>>> ('teachers', 'english', 'name', 'william'),
>>> ('teachers', 'english', 'surname', 'shakespeare'),
>>> ('teachers', 'english', 'grade', 3),
>>> ('teachers', 'physics', 'name', 'richard'),
>>> ('teachers', 'physics', 'surname', 'feyman'),
>>> ('teachers', 'physics', 'grade', 4)]
```

#To reverse this, we call:

```
>>> assert items_to_tree(items) == school
```

## **tree\_update**

`pyg.base._dict.tree_update` (tree, update, types=(`<class 'dict'>`, `<class 'pyg.base._dict.Dict'>`, `<class 'pyg.base._dictattr.dictattr'>`), ignore=None)  
equivalent to `dict.update()` except: not in-place and also updates further down the tree

**Example:**

```
>>> ranking = dict(cambridge = dict(trinity = 1, stjohms = 2, christ = 3),
                   oxford = dict(trinity = 1, jesus = 2, magdalene = 3))
>>> new_ranking = dict(oxford = dict(wolfson = 3, magdalene = 4))
```

```
>>> print(tree_repr(tree_update(ranking, new_ranking)))
```

```
>>> cambridge:
>>>     {'trinity': 1, 'stjohms': 2, 'christ': 3}
>>> oxford:
>>>     {'trinity': 1, 'jesus': 2, 'magdalene': 4, 'wolfson': 3}
```

Note how values for magdalene in Oxford were overwritten even though they are further down the tree

**Example:** using ignore

```
>>> update = dict(a = None, b = np.nan, c = 0)
>>> tree = dict(a = 1, b = 2, c = 3)
>>> assert tree_update(tree, update) == update
>>> assert tree_update(tree, update, ignore = [None]) == dict(a = 1, b = np.nan, c = 0)
>>> assert tree_update(tree, update, ignore = [None, np.nan]) == dict(a = 1, b = 2, c = 0)
>>> assert tree_update(tree, update, ignore = [None, np.nan, 0]) == tree
```

**Parameters:**

**tree :** *tree*

existing tree.

**update :** *tree*

new information.

**types :** *types, optional*

see `tree_items`. The default is (`dict`, `Dict`, `dictattr`).

**Returns:**

**tree**

Welcome to pyg!

updated tree.

## ***tree\_setitem***

`pyg.base._dict.tree_setitem` (tree, key, value, ignore=None, types=None)  
sets an item of a tree

### **Parameters:**

**tree** : tree (dicts of dict) **key** : a dot-separated string or a tuple of values  
the branch to hang value on

**value** : *object*  
the leaf at the end of the branch

**ignore** : *None or list, optional*  
what values of leaf will be ignored and not overwrite existing data. The default is None.

**types** : *types, optional*  
As we go down the tree, when do we stop and say: what we have is a leaf already?

### **Example:**

```
>>> tree = dict()
>>> tree_setitem(tree, 'a', 1)
>>> assert tree == dict(a = 1)
>>> tree_setitem(tree, 'b.c', 2)
>>> assert tree == {'a': 1, 'b': {'c': 2}}
>>> tree_setitem(tree, ('b', 'c', 'd'), 2)
>>> tree_setitem(tree, ('b', 'c', 'e'), 3)
>>> assert tree == {'a': 1, 'b': {'c': {'d': 2, 'e': 3}}}
```

### **Example: types**

```
>>> from pyg import *
>>> tree = dict(mycell = cell(lambda a, b: a * b, b = 2, a = cell(lambda x: x**2, x = cell(
>>> # We are missing y...
>>> tree_setitem(tree, 'mycell.a.x.y', 3, types = (dict, cell)) ## drill into cell
>>> assert tree['mycell'].a.x.y == 3
>>> tree_setitem(tree, 'mycell.a.x.y', 1) ## stop when you hit cell
>>> assert tree['mycell'].a.x == dict(y = 1)
```

None.

## ***tree\_repr***

`pyg.base._tree_repr.tree_repr` (value, offset=0)  
a cleaner representation of a tree

### **Example:**

```
>>> school = dict(pupils = dict(id1 = dict(name = 'james', surname = 'maxwell', gender = 'm'),
>>>                                id2 = dict(name = 'adam', surname = 'smith', gender = 'm'),
>>>                                id3 = dict(name = 'michell', surname = 'obama', gender = 'f'),
>>>                                ),
>>> teachers = dict(math = dict(name = 'albert', surname = 'einstein', grade = 3),
>>>                   english = dict(name = 'william', surname = 'shakespeare', grade = 4),
>>>                   physics = dict(name = 'richard', surname = 'feynman', grade = 4)
>>>                   ))
```

```
>>> print(tree_repr(school, 4))
>>> pupils:
>>> id1:
```

Welcome to pyg!

```
>>>      {'name': 'james', 'surname': 'maxwell', 'gender': 'm'}
>>>   id2:
>>>      {'name': 'adam', 'surname': 'smith', 'gender': 'm'}
>>>   id3:
>>>      {'name': 'michell', 'surname': 'obama', 'gender': 'f'}
>>> teachers:
>>>   math:
>>>      {'name': 'albert', 'surname': 'einstein', 'grade': 3}
>>>   english:
>>>      {'name': 'william', 'surname': 'shakespeare', 'grade': 3}
>>>   physics:
>>>      {'name': 'richard', 'surname': 'feyman', 'grade': 4}
```

#### Parameters:

value : a tree

**offset** : *int, optional*

offset from the left for printing. The default is 0.

#### Returns:

**string**

a tree-like string representation of a dict-of-dicts.

### items\_to\_tree

`pyg.base._dict.items_to_tree` (items, tree=None, raise\_if\_duplicate=True, ignore=None, types=None)  
converts **items** to branches of a tree. If an original **tree** is provided, hang the additional branches on the existing tree. If **ignore** is provided as a list of values, will not overwrite branches with last value (the leaf) in these values

#### Example:

```
>>> items = [('cambridge', 'smith', 'economics',),
             ('cambridge', 'keynes', 'economics'),
             ('cambridge', 'lyons', 'maths'),
             ('cambridge', 'maxwell', 'maths'),
             ('oxford', 'penrose', 'maths'),
             ]
```

```
>>> tree = items_to_tree(items)
>>> print(tree_repr(tree))
```

```
>>> cambridge:
>>>   smith:
>>>       economics
>>>   keynes:
>>>       economics
>>>   lyons:
>>>       maths
>>>   maxwell:
>>>       maths
>>> oxford:
>>>   {'penrose': 'maths'}
```

We can add to tree:

#### Parameters:

**items** : *list of tuples*,

items are just like dict items, only longer,

**tree** : *tree, optional*



Welcome to pyg!

a pre-existing tree of trees. The default is None.

**raise\_if\_duplicate** : *TYPE, optional*

DESCRIPTION. The default is True.

**ignore** : *list, optional*

list of values that when over-writing an existing tree, should ignore. The default is None.

**Example:** using ignore

```
>>> tree = dict(a = 1, b = 'keep_old_value')
>>> update = dict(a = 'valid_new_value', b = None, c = None)
>>> tree_update(tree, update, ignore = [None])
>>> {'a': valid_new_value, 'b': 'keep_old_value', 'c': None}
```

- a is over-riden as the new value is valid
- b is not over-riden since the update b = None is considered invalid
- c is added as it did not exist before, even though c = None is invalid value

**Returns:**

tree : dict of dicts

## ***tree\_to\_table***

pyg.base.\_tree.**tree\_to\_table** (tree, pattern)

The best way to understand is to give an example:

**Examples:**

```
>>> school = dict(pupils = dict(id1 = dict(name = 'james', surname = 'maxwell', gender = 'm'),
                                id2 = dict(name = 'adam', surname = 'smith', gender = 'm'),
                                id3 = dict(name = 'michell', surname = 'obama', gender = 'f'),
                                ),
                  teachers = dict(math = dict(name = 'albert', surname = 'einstein', grade = 3),
                                  english = dict(name = 'william', surname = 'shakespeare', grade = 4),
                                  physics = dict(name = 'richard', surname = 'feynman', grade = 4)
                                  ))
```

Suppose we wanted to identify all male students:

```
>>> res = tree_to_table(school, 'pupils/%id/gender/m')
>>> assert res == [dict(id = 'id1'), dict(id = 'id2')]
```

or grades:

```
>>> res = tree_to_table(school, 'teachers/%subject/grade/%grade')
>>> assert res == [{'grade': 3, 'subject': 'math'},
                   {'grade': 3, 'subject': 'english'},
                   {'grade': 4, 'subject': 'physics'}]
```

**Parameters:**

**tree** : *tree (dict of dicts)*

tree is a yaml-like structure

**pattern** : *string*

The pattern whose instances we wish to find in tree

**Returns:**

list of dicts

## *list functions*

### *as\_list*

`pyg.base._as_list.as_list` (value, none=False)  
returns a list of the original object.

#### Example:

```
>>> assert as_list(None) == []
>>> assert as_list(4) == [4]
>>> assert as_list((1,2,)) == [1,2]
>>> assert as_list([1,2,]) == [1,2]
>>> assert eq(as_list(np.array([1,2,])) , [np.array([1,2,])])
>>> assert as_list(dict(a = 1)) == [dict(a=1)]
```

In practice, this function is has an incredible useful usage:

**Example:** using `as_list` to give flexibility on `*args`

```
>>> def my_sum(*values):
>>>     values = as_list(values)
>>>     return sum(values)
```

```
>>> assert my_sum(1,2,3) == 6
>>> assert my_sum([1,2,3]) == 6 ## This is nice... wasn't possible before
```

#### Parameters:

value : anything none : bool optional

Shall I return None as a value? The default is False and we return [], if True, returns [None]

#### Returns:

**list**

a list of original objects.

### *as\_tuple*

`pyg.base._as_list.as_tuple` (value, none=False)  
returns a tuple of the original object.

#### Example:

```
>>> assert as_tuple(None) == ()
>>> assert as_tuple(4) == (4,)
>>> assert as_tuple((1,2,)) == (1,2)
>>> assert as_tuple([1,2,]) == (1,2)
>>> assert eq(as_tuple(np.array([1,2,])) , (np.array([1,2,]),))
>>> assert as_tuple(dict(a = 1)) == (dict(a=1),)
```

In practice, this function is has an incredible useful usage:

**Example:** using `as_list` to give flexibility on `*args`

```
>>> def my_sum(*values):
>>>     values = as_tuple(values)
>>>     return sum(values)
```

```
>>> assert my_sum(1,2,3) == 6
>>> assert my_sum([1,2,3]) == 6 ## This is nice... wasn't possible before
```

Welcome to pyg!

**Parameters:**

value : anything none : bool optional

Shall I return None as a value? The default is False and we return [], if True, returns [None]

**Returns:**

**tuple**

a tuple of original objects.

## *first*

`pyg.base._as_list.first` (value)

returns the first value in a list (None if empty list) or the original if value not a list

**Example:**

```
>>> assert first(5) == 5
>>> assert first([5,5]) == 5
>>> assert first([]) is None
>>> assert first([1,2]) == 1
```

## *last*

`pyg.base._as_list.last` (value)

returns the last value in a list (None if empty list) or the original if value not a list

**Example:**

```
>>> assert last(5) == 5
>>> assert last([5,5]) == 5
>>> assert last([]) is None
>>> assert last([1,2]) == 2
```

## *unique*

`pyg.base._as_list.unique` (value)

returns the asserted unique value in a list (None if empty list) or the original if value not a list. Throws an exception if list non-unique

**Example:**

```
>>> assert unique(5) == 5
>>> assert unique([5,5]) == 5
>>> assert unique([]) is None
>>> with pytest.raises(ValueError):
>>>     unique([1,2])
```

## *Comparing and Sorting*

### *cmp*

`pyg.base._sort.cmp` (x, y)

Implements lexcompare while allowing for comparison of different types. First compares by type, then by length, then by keys and finally on values

**Parameters:**

**x** : *obj*

1st object to be compared.

**y** : *obj*

Welcome to pyg!

2nd object to be compared.

**Returns:**

**int**

returns -1 if  $x < y$  else 1 if  $x > y$  else 0

**Examples:**

```
>>> assert cmp('2', 2) == 1
>>> assert cmp(np.int64(2), 2) == 0
>>> assert cmp(None, 2.0) == -1 # None is smallest
>>> assert cmp([1,2,3], [4,5]) == 1 # [1,2,3] is longer
>>> assert cmp([1,2,3], [1,2,0]) == 1 # lexical sorting
>>> assert cmp(dict(a = 1, b = 2), dict(a = 1, c = 2)) == -1 # lexical sorting on keys
>>> assert cmp(dict(a = 1, b = 2), dict(b = 2, a = 1)) == 0 # order does not matter
```

## Cmp

`pyg.base._sort.Cmp(x)`

class wrapper of cmp, allowing us to compare objects of different types

**Example:**

```
>>> with pytest.raises(TypeError):
>>>     sorted([1,2,3,None])
```

```
>>> # but this is fine:
>>> assert sorted([1,3,2,None], key = Cmp) == [None, 1, 2, 3]
```

## sort

`pyg.base._sort.sort(iterable)`

implements sorting allowing for comparing of not-same-type objects

**Parameters:**

**iterable** : *iterable*

values to be sorted

**Returns:**

**list**

sorted list.

**Example:**

```
>>> with pytest.raises(TypeError):
>>>     sorted([1,2,3,None])
>>> # but this is fine:
>>> sort([1,3,2,None]) == [None, 1, 2, 3]
```

## eq

`pyg.base._eq.eq(x, y)`

A better nan-handling equality comparison. Here is the problem:

```
>>> import numpy as np
>>> assert not np.nan == np.nan ## What?
```

The nan issue extends to `np.arrays`...

```
>>> assert list(np.array([np.nan,2]) == np.array([np.nan,2])) == [False, True]
```

Welcome to pyg!

but not to lists...

```
>>> assert [np.nan] == [np.nan]
```

But wait, if the lists are derived from np.arrays, then no equality...

```
>>> assert not list(np.array([np.nan])) == list(np.array([np.nan]))
```

The other issue is inheritance:

```
>>> class FunnyDict(dict):
>>>     def __getitem__(self, key):
>>>         return 5
>>> assert dict(a = 1) == FunnyDict(a=1) ## equality seems to ignore any type mismatch
>>> assert not dict(a = 1)['a'] == FunnyDict(a = 1)['a']
```

There are also issues with partial

```
>>> from functools import partial
>>> f = lambda a: a + 1
>>> x = partial(f, a = 1)
>>> y = partial(f, a = 1)
>>> assert not x == y
```

```
>>> import pandas as pd
>>> import pytest
>>> from pyg import eq
```

```
>>> assert eq(np.nan, np.nan) ## That's better
>>> assert eq(x = np.array([np.nan,2]), y = np.array([np.nan,2]))
>>> assert eq(np.array([np.array([1,2]),2], dtype = 'object'), np.array([np.array([1,2]),2]))
>>> assert eq(np.array([np.nan,2]), np.array([np.nan,2]))
>>> assert eq(dict(a = np.array([np.array([1,2]),2], dtype = 'object')) , dict(a = np.array([np.array([1,2]),2], dtype = 'object')))
>>> assert eq(dict(a = np.array([np.array([1,np.nan]),np.nan])), dict(a = np.array([np.array([1,np.nan]),np.nan])))
>>> assert eq(np.array([np.array([1,2]),dict(a = np.array([np.array([1,2]),2]))]), np.array([np.array([1,2]),dict(a = np.array([np.array([1,2]),2]))]))
```

```
>>> assert not eq(dict(a = 1), FunnyDict(a=1))
>>> assert eq(1, 1.0)
>>> assert eq(x = pd.DataFrame([1,2]), y = pd.DataFrame([1,2]))
>>> assert eq(pd.DataFrame([np.nan,2]), pd.DataFrame([np.nan,2]))
>>> assert eq(pd.DataFrame([1,np.nan], columns = ['a']), pd.DataFrame([1,np.nan], columns = ['a']))
>>> assert not eq(pd.DataFrame([1,np.nan], columns = ['a']), pd.DataFrame([1,np.nan], columns = ['a']))
```

## in

pyg.base.\_eq.in\_(x, seq)

Evaluates if x is in seq, avoiding issues such as these:

```
>>> s = pd.Series([1,2,3])
>>> with pytest.raises(ValueError):
>>>     s in [None]
>>> assert not in_(s, [None])
>>> assert in_(s, [None, s])
```

## bits and pieces

### type functions

pyg.base.\_types.is\_arr(value)

is value a numpy array of non-zero-size

pyg.base.\_types.is\_bool(value)

Welcome to pyg!

is value a Bool, or a `np.bool_` type

`pyg.base._types.is_date` (value)  
is value a date type: either `datetime.date`, `datetime.datetime` or `np.datetime64`

`pyg.base._types.is_df` (value)  
is value a `pd.DataFrame`

`pyg.base._types.is_dict` (value)  
is value a dict

`pyg.base._types.is_float` (value)  
is value an float, or any variant of `np.float`

`pyg.base._types.is_int` (value)  
is value an int, or any variant of `np.intN` type

`pyg.base._types.is_iterable` (value)  
is value Iterable excluding a string

`pyg.base._types.is_len` (value)  
is value of zero length (or has no len at all)

`pyg.base._types.is_list` (value)  
is value a list

`pyg.base._types.is_nan` (value)  
is value a nan or an inf. Unlike `np.isnan`, works for non numeric

`pyg.base._types.is_none` (value)  
is value None

`pyg.base._types.is_num` (value)  
is `_int(value)` or `_float(value)`

`pyg.base._types.is_pd` (value)  
is value a `pd.DataFrame`/`pd.Series`

`pyg.base._types.is_series` (value)  
is value a `pd.Series`

`pyg.base._types.is_str` (value)  
is value a str, or a `np.str_` type

`pyg.base._types.is_ts` (value)  
is value a pandas dataframe which is indexed by datetimes

`pyg.base._types.is_tuple` (value)  
is value a tuple

`pyg.base._types.nan2none` (value)  
convert `np.nan`/`np.inf` to None

## zipper

`pyg.base._zip.zipper` (\*values)  
a safer version of zip

**Examples:** zipper works with single values as well as full list:

```
>>> assert list(zipper([1,2,3], 4)) == [(1, 4), (2, 4), (3, 4)]
>>> assert list(zipper([1,2,3], [4,5,6])) == [(1, 4), (2, 5), (3, 6)]
>>> assert list(zipper([1,2,3], [4,5,6], [7])) == [(1, 4, 7), (2, 5, 7), (3, 6, 7)]
>>> assert list(zipper([1,2,3], [4,5,6], None)) == [(1, 4, None), (2, 5, None), (3, 6, None)]
>>> assert list(zipper([1,2,3], np.array([4,5,6]), None)) == [(1, 4, None), (2, 5, None), (3, 6, None)]
```

**Examples:** zipper rejects multi-length lists

```
>>> import pytest
>>> with pytest.raises(ValueError):
>>>     zipper([1,2,3], [4,5])
```

**Parameters:**

**\*values** : *lists*

values to be zipped

**Returns:**

zipped values

## reducer

`pyg.base._reducer.reducer` (function, sequence, default=None)  
reduce adds stuff to zero by defaults. This is not needed.

**Parameters:**

**function** : *callable*

binary function.

**sequence** : *iterable*

list of inputs to be applied iteratively to reduce.

**default** : *TYPE, optional*

A default value to be returned with an empty sequence

**Example:**

```
>>> from operator import add, mul
>>> from functools import reduce
>>> import pytest
```

```
>>> assert reducer(add, [1,2,3,4]) == 10
>>> assert reducer(mul, [1,2,3,4]) == 24
>>> assert reducer(add, [1]) == 1
```

```
>>> assert reducer(add, []) is None
>>> with pytest.raises(TypeError):
>>>     reduce(add, [])
```

## reducing

`class pyg.base._reducer.reducing` (function=None, \*args, \*\*kwargs)  
Makes a bivariate-function being able to act on a sequence of elements using reduction

**Example:**

```
>>> from operator import mul
>>> assert reducing(mul)([1,2,3,4]) == 24
>>> assert reducing(mul)(6,4) == 24
```

Since `a.join(b).join(c).join(d)` is also very common, we provide a simple interface for that:

**Example:** chaining

```
>>> assert reducing('__add__')( [1,2,3,4] ) == 10
>>> assert reducing('__add__')( 6,4 ) == 10
```

```
d = dictable(a = [1,2,3,5,4]) reducing('inc')(d, dict(a=1))
```

## ***logger and get\_logger***

`pyg.base._logger.get_logger` (name='pyg', level='info', fmt='%(%asctime)s - %(name)s - %(levelname)s - %(message)s', file=False, console=True)

quick utility to simplify loggers creation and ensure we cache them and do not add to many handlers

### **Parameters:**

**name** : *str, optional*

name of logger. The default is 'pyg'.

**level** : *str, optional*

DEBUG/INFO/WARN etc. The default is 'info'.

**fmt** : *str, optional*

string formatting for messages. The default is '%(asctime)s - %(name)s - %(levelname)s - %(message)s'.

**file** : *bool/str, optional*

the name of the file to log to. The default is False = do not log to file.

**console** : *bool, optional*

log to console? The default is True.

### **Returns:**

logging.logger

## ***access functions***

These are useful to convert object-oriented code to declarative functions

`pyg.base._getitem.callattr` (value, attr, args=None, kwargs=None)

gets the attribute(s) from a value and calls its

### **Example:**

```
>>> from pyg import *
>>> value = Dict(function = lambda a, b: a + b)
>>> assert callattr(value, 'function', kwargs = dict(a = 1, b = 2)) == 3
>>> assert callattr(value, attr = 'function', args = (1, 2), kwargs = None) == 3
```

```
>>> ts = pd.Series(np.random.normal(0,1,1000))
>>> assert ts.std() == callattr(ts, 'std')
>>> assert eq(ts.ewm(com = 10).mean(), callattr(ts, ['ewm','mean'], kwargs = [{'com':10}],
```

```
>>> d = dictable(a = [1,2,3,4,1,2], b = list('abcdef'))
>>> assert callattr(d, ['inc', 'exc'], kwargs = [dict(a = 2), dict(b = 'f')]) == d.inc(a =
```

**value** : *obj*

object that contains an item.

**attr** : *string(s)*

key within object.

**args** : *tuple, optional*

tuple of values to be fed to function. The default is None.

**kwargs** : *dict, optional*

kwargs to be fed to the method. The default is None.

`pyg.base._getitem.callitem` (value, key, args=None, kwargs=None)

gets an item and calls it

### **Example:**



```
>>> c = dict(function = lambda a, b: a + b)
>>> assert callitem(c, 'function', kwargs = dict(a = 1, b = 2)) == 3
>>> assert callitem(c, 'function', args = (1, 2)) == 3
```

**value** : *obj*

object that contains an item.

**key** : *string*

key within object.

**args** : *tuple, optional*

tuple of values to be fed to function. The default is None.

**kwargs** : *dict, optional*

kwargs to be fed to the method. The default is None.

`pyg.base._getitem.getitem` (value, key, \*default)  
gets an item, like `getattr`

**Example:**

```
>>> a = dict(a = 1)
>>> assert getitem(a, 'a') == 1
>>> assert getitem(a, 'b', 2) == 2
```

```
>>> import pytest
>>> with pytest.raises(KeyError):
>>>     getitem(a, 'b')
```

## inspection

There are a few functions extending the inspect module.

`pyg.base._inspect.argspec_add` (fullargspec, \*\*update)  
adds new args with default values at the end of the existing args

**Parameters:**

**fullargspec** : *FullArgSpec*

DESCRIPTION.

**\*\*update** : *dict*

parameter names with their default values.

**Returns:**

*FullArgSpec*

**Example:**

```
>>> f = lambda b : b
>>> argspec = getargspec(f)
>>> updated = argspec_add(argspec, axis = 0)
>>> assert updated.args == ['b', 'axis'] and updated.defaults == (0,)
```

```
>>> f = lambda b, axis : None ## axis already exists without a default
>>> argspec = getargspec(f)
>>> updated = argspec_add(argspec, axis = 0)
>>> assert updated == argspec
```

```
>>> f = lambda b, axis = 1 : None ## axis already exists with a different default
>>> argspec = getargspec(f)
>>> updated = argspec_add(argspec, axis = 0)
>>> assert updated == argspec
```

Welcome to pyg!

`pyg.base._inspect.argspec_defaults` (function)

**Returns:** the function defaults as a dict rather than using the inspect structure

**Example:**

```
>>> f = lambda a, b = 1: a+b
>>> assert argspec_defaults(f) == dict(b=1)
```

```
>>> g = partial(f, b = 2)
>>> assert argspec_defaults(g) == dict(b=2)
```

**Parameters:**

function : callable

**Returns:**

defaults as a dict.

`pyg.base._inspect.argspec_required` (function)

**Parameters:**

function : callable

**Returns:**

list

parameters that *must* be provided in order to run the function

`pyg.base._inspect.argspec_update` (argspec, \*\*kwargs)

generic function to create new FullArgSpec (python 3) or normal ArgSpec (python 2)

**Parameters:**

**argspec** : *FullArgSpec*

The argspec of the dunction

**\*\*kwargs** : *TYPE*

updates

**Returns:**

FullArgSpec

**Example:**

```
>>> f = lambda a, b =1 : a + b
>>> argspec = getargspec(f)
>>> assert argspec_update(argspec, args = ['a', 'b', 'c']) == inspect.FullArgSpec(**{'anno
                                                                 'args': ['a',
                                                                 'defaults': (
                                                                 'kwonlyargs':
                                                                 'kwonlydefault
                                                                 'varargs': No
                                                                 'varkw': None
```

`pyg.base._inspect.call_with_callargs` (function, callargs)

replicates inspect.getcallargs with support to functions within decorators

**Example:**

```
>>> function = lambda a, b, *args, **kwargs: 1+b+len(args)+10*len(kwargs)
>>> args = (1,2,3,4,5); kwargs = dict(c = 6, d = 7)
>>> assert function(*args, **kwargs) == 26
>>> callargs = getcallargs(function, *args, **kwargs)
>>> assert call_with_callargs(function, callargs) == 26
```

`pyg.base._inspect.getargs` (function, n=0)

### Parameters:

**function** : *callable*

The function for which we want the args

**n** : *int optional*

get the name of the args after allowing for n args to be set by \*args. The default is 0.

### Returns:

None or a list of args

pyg.base.\_inspect.**getargspec** (function)

Extends inspect.getfullargspec to allow us to decorate functions with a signature.

### Parameters:

**function** : *callable*

function for which we want to know argspec.

### Returns:

inspect.FullArgSpec

pyg.base.\_inspect.**getcallargs** (function, \*args, \*\*kwargs)

replicates inspect.getcallargs with support to functions within decorators

### Example:

```
>>> from pyg import *; import inspect
>>> function = lambda a, b, *myargs, **mykwargs: 1
>>> args = (1,2,3,4,5); kwargs = dict(c = 6, d = 7)
>>> assert getcallargs(function, *args, **kwargs) == inspect.getcallargs(function, *args,
```

```
>>> function = lambda a: a + 1
>>> args = (); kwargs = dict(a=1)
>>> assert getcallargs(function, *args, **kwargs) == inspect.getcallargs(function, *args,
```

```
>>> function = lambda a, b = 1: 1
>>> args = (); kwargs = dict(a=1)
>>> assert getcallargs(function, *args, **kwargs) == inspect.getcallargs(function, *args,
>>> args = (); kwargs = dict(a=1, b = 2)
>>> assert getcallargs(function, *args, **kwargs) == inspect.getcallargs(function, *args,
>>> args = (1,); kwargs = {}
>>> assert getcallargs(function, *args, **kwargs) == inspect.getcallargs(function, *args,
>>> args = (1,2); kwargs = {}
>>> assert getcallargs(function, *args, **kwargs) == inspect.getcallargs(function, *args,
>>> args = (1,); kwargs = {'b' : 2}
>>> assert getcallargs(function, *args, **kwargs) == inspect.getcallargs(function, *args,
```

pyg.base.\_inspect.**kwargs2args** (function, args, kwargs)

converts a list of parameters that were provided as kwargs, into args

### Example:

```
>>> assert kwargs2args(lambda a, b: a+b, (), dict(a = 1, b=2)) == ([1,2], {})
```

### Parameters:

**function** : *callable* **args** : *tuple*

parameters of function.

**kwargs** : *dict*

key-word parameters of function.

### Returns:

**tuple**

a pair of a function args, kwargs.

## pyg.mongo

A few words on MongoDB, a no-SQL database versus SQL:

- Mongo has 'collections' that are the equivalent of tables
- Mongo will refer to 'documents' instead of traditional records. Those records are unstructured and look like trees: dicts of dicts. They contain arbitrary objects as well as just the primary types a SQL database is designed to support.
- Mongo collections do not have the concept of primary keys
- Mongo WHERE SQL clause is replaced by a query in a form of a dict "presented" to the collection object.
- Mongo SELECT SQL clause is replaced by a 'projection' on the cursor, specifying what fields are retrieved.

## Query generator

We start by simplifying the way we generate mongo query dictionaries.

### q and Q

```
class pyg.mongo._q.Q (keys=None)
```

The MongoDB interface for query of a collection (table) is via a creation of a complicated looking dict: <https://docs.mongodb.com/manual/tutorial/query-documents/>

This is rather complicated for the average user so Q simplifies it greatly. Q is based on TinyDB and users of TinyDB will recognise it. <https://tinydb.readthedocs.io/en/latest/usage.html>

q is the singleton of Q.

q supports both *calling* to generate the querying dict

```
>>> q(a = 1, b = 2)
```

or

```
>>> (q.a == 1) & (q.b == 2) # {"$and": [{"a": {"$eq": 1}}, {"b": {"$eq": 2}}]}
>>> (q.a == 1) | (q.b == 2) # {"$or": [{"a": {"$eq": 1}}, {"b": {"$eq": 2}}]}
```

or indeed

```
>>> q(q.a == 1, q.b == 2)
```

### Example:

```
>>> from pyg import q
>>> import re
```

```
>>> assert dict(q.a == 1) == {"a": {"$eq": 1}}
>>> assert dict(q(a = [1,2])) == {'a': {'$in': [1, 2]}}
>>> assert dict(q(q.a == [1,2], q.b > 3)) == {'$and': [{"a": {"$in": [1, 2]}}, {"b": {"$gt": 3}}]}
>>> assert dict(q(a = re.compile('^hello'))) == {'a': {'regex': '^hello'}} # a regex q
```

```
>>> assert dict(q.a.exists + q.b.not_exists) == {"$and": [{"a": {"$exists": true}}, {"b": {"$exists": false}}]}
```

```
>>> assert dict(~(q.a==1)) == {'$not': {"a": {"$eq": 1}}}
```

## Tables in Mongo

### mongo\_cursor

mongo\_cursor has hybrid functionality of a Mongo cursor and Mongo collection objects.

`class pyg.mongo._cursor.mongo_cursor(cursor, writer=None, reader=None, query=None, **_)`  
 mongo\_cursor is a souped-up combination of mongo.Cursor and mongo.Collection with a simple API.

#### Parameters:

**cursor** : MongoDB cursor or MongoDB collection

**writer** : *True/False/string, optional*

The default is None.

writer allows you to transform the data before saving it in Mongo. You can create a function yourself or use built-in options:

- False: do nothing, save the document as is
- True/None: use `pyg.base.encode` to encode objects. This will transform numpy array/dataframes into bytes that can be stored
- '.csv': save dataframes into csv files and then save reference of these files to mongo
- '.parquet' save dataframes into .parquet and np.ndarray into .npy files.

For .csv and .parquet to work, you will need to specify WHERE the document is to be saved. This can be done either:

- the document has a 'root' key, specifying the root.
- you specify root by setting `writer = 'c:/%name%surname.parquet'`

**reader** : *callable or None, optional*

The default is None, using decode. Use `reader = False` to passthru

**query** : *dict, optional*

This is used to specify the Mongo query, e.g. `q.a==1`.

**\*\*\_** :

#### Example:

```
>>> from pyg import *
>>> cursor = mongo_table('test', 'test')
>>> cursor.drop()
```

## insert some data

```
>>> table = dictable(a = range(5)) * dictable(b = range(5))
>>> cursor.insert_many(table)
>>> cursor.set(c = lambda a, b: a * b)
```

#### Filtering:

```
>>> assert len(cursor) == 25
>>> assert len(cursor.find(a = 3)) == 5
>>> assert len(cursor.exc(a = 3)) == 20
>>> assert len(cursor.find(a = [3,2]).find(q.b<3)) == 6 ## can chain queries as well as us
```

#### Row access:

```
>>> cursor[0]
```

```
{'_id': ObjectId('603aec85cd15e2c090c07b87'), 'a': 0, 'b': 0}
```

```
>>> cursor[::] - '_id' == dictable(cursor) - '_id'
```

```
>>> dictable[25 x 3]
>>> a|b|c
>>> 0|0|0
>>> 0|1|0
>>> 0|2|0
>>> ...25 rows...
>>> 4|2|8
>>> 4|3|12
>>> 4|4|16
```

#### Column access:

```
>>> cursor[['a', 'b']] ## just columns 'a' and 'b'
>>> del cursor['c'] ## delete all c
>>> cursor.set(c = lambda a, b: a * b)
>>> assert cursor.find_one(a = 3, b = 2)[0].c == 6
```

#### Example: root specification

```
>>> from pyg import *
>>> t = mongo_table('test', 'test', writer = 'c:/temp/%name/%surname.parquet')
>>> t.drop()
>>> doc = dict(name = 'adam', surname = 'smith', ts = pd.Series(np.arange(10)))
>>> t.insert_one(doc)
>>> assert eq(pd_read_parquet('c:/temp/adam/smith/ts.parquet'), doc['ts'])
>>> assert eq(t[0]['ts'], doc['ts'])
>>> doc = dict(name = 'beth', surname = 'brown', a = np.arange(10))
>>> t.insert_one(doc)
```

Since `mongo_cursor` is too powerful, we also have a `mongo_reader` version which is read-only.

#### `delete_many()`

Equivalent to `drop()`: deletes all documents the cursor currently points to.

##### Note:

If you want to drop a subset of the data, then use `c.find(criteria).delete_many()`

##### Returns:

itself

#### `delete_one(*args, **kwargs)`

drops a specific record after verifying exactly one exists.

##### Parameters:

*\*args* : query *\*\*kwargs* : query

##### Returns:

itself

#### `drop()`

Equivalent to `drop()`: deletes all documents the cursor currently points to.

##### Note:

If you want to drop a subset of the data, then use `c.find(criteria).delete_many()`

##### Returns:

itself

#### `insert_many(table)`

inserts multiple documents into the collection

**table** : *sequence of documents*

list of dicts or dictable

mongo\_cursor

**Example:** simple insertion

```
>>> from pyg import *
>>> t = mongo_table('test', 'test')
>>> t = t.drop()
>>> values = dictable(a = [1,2,3,4,], b = [5,6,7,8])
>>> t = t.insert_many(values)
>>> t[::]
```

```
>>> dictable[4 x 3]
>>> _id |a|b
>>> 602daee68c336f6429a77bdd|1|5
>>> 602daee68c336f6429a77bde|2|6
>>> 602daee68c336f6429a77bdf|3|7
>>> 602daee68c336f6429a77be0|4|8
```

**Example:** update

```
>>> table = t[::]
>>> modified = table(b = lambda b: b**2)
>>> t = t.insert_many(modified)
```

Since each of the documents we uploaded already has an `_id`...

```
>>> assert len(t) == 4
>>> t[::]
>>> dictable[4 x 3]
>>> _id |a|b
>>> 602daee68c336f6429a77bdd|1|25
>>> 602daee68c336f6429a77bde|2|36
>>> 602daee68c336f6429a77bdf|3|49
>>> 602daee68c336f6429a77be0|4|64
```

**insert\_one (doc)**

inserts/updates a single document.

If the document ALREADY has `_id` in it, it updates that document If the document has no `_id` in it, it inserts it as a new document

**Parameters:**

**doc** : *dict*

document.

**Example:**

```
>>> from pyg import *
>>> t = mongo_table('test', 'test')
>>> t = t.drop()
>>> values = dictable(a = [1,2,3,4,], b = [5,6,7,8])
>>> t = t.insert_many(values)
```

**Example:** used to update an existing document

```
>>> doc = t[0]
>>> doc['c'] = 8
>>> str(doc)
>>> '{"_id": ObjectId('602d36150a5cd32717323197'), 'a': 1, 'b': 5, 'c': 8}"
```

```
>>> t = t.insert_one(doc)
>>> assert len(t) == 4
>>> assert t[0] == doc
```

**Example:** used to insert

```
>>> doc = Dict(a = 1, b = 8, c = 10)
>>> t = t.insert_one(doc)
>>> assert len(t) == 5
>>> t.drop()
```

*property raw*

returns an unfiltered mongo\_reader

**set (\*\*kwargs)**

updates all documents in current cursor based on the kwargs. It is similar to update\_many but supports also functions

**Parameters:**

kwargs: dict of values to be updated

**Example:**

```
>>> from pyg import *
>>> t = mongo_table('test', 'test')
>>> t = t.drop()
>>> values = dictable(a = [1,2,3,4,], b = [5,6,7,8])
>>> t = t.insert_many(values)
>>> assert t[::]-'_id' == values
```

```
>>> t.set(c = lambda a, b: a+b)
>>> assert t[::]-'_id' == values(c = [6,8,10,12])
>>> t.set(d = 1)
>>> assert t[::]-'_id' == values(c =lambda a, b: a+b)(d = 1)
```

**Returns:**

itself

**update\_many (doc, upsert=False)**

updates all documents in current cursor based on the doc. The two are equivalent:

```
>>> cursot.update_many(doc)
>>> collection.update_many(cursor.spec, { 'set' : update})
```

**Parameters:**

doc : dict of values to be updated

**Returns:**

itself

**update\_one (doc, upsert=True)**

- updates a document if an \_id is present in doc.
- insert a document if an \_id is not present and upsert is true

**Parameters:**

**doc :** *document*

doc to be upserted.

**upsert :** *bool, optional*



insert if no document present? The default is True.

**Returns:**

**doc**

document updated.

## ***mongo\_reader***

mongo\_reader is a read-only version of the cursor to avoid any unintentional damage to database.

`class pyg.mongo._reader.mongo_reader (cursor, writer=None, reader=None, query=None, **_)`

mongo\_reader is a read-only version of the mongo\_cursor. You can instantiate it with a mongo\_reader(cursor) call where cursor can be a mongo\_cursor, a pymongo.Cursor or a pymongo.Collection

*property* address

**Returns:**

**tuple**

A unique combination of the client address, database name and collection name, identifying the collection uniquely.

**clone (\*\*params)**

**Returns:**

**mongo\_reader**

Returns a cloned version of current mongo\_reader but allows additional parameters to be set (see spec and project)

*property* collection

**Returns:**

pymongo.Collection object

**count ()**

cursor.count() and len(cursor) are the same and return the number of documents matching current specification.

**distinct (key)**

returns the distinct values of the key

**key : str**

a key in the documents.

**list of strings**

distinct values

**Example:**

```
>>> from pyg import *; import pymongo
>>> table = pymongo.MongoClient()['test']['test']
>>> table.insert_one(dict(name = 'alan', surname = 'abrahams', age = 39, marriage = dt(2010, 1, 1)))
>>> table.insert_one(dict(name = 'barbara', surname = 'brown', age = 50, marriage = dt(2010, 1, 1)))
>>> table.insert_one(dict(name = 'charlie', surname = 'cohen', age = 20))
```

```
>>> t = mongo_reader(table)
>>> assert t.name == t.distinct('name') == ['alan', 'barbara', 'charlie']
>>> table.drop()
```

**docs (doc='doc', \*keys)**

self[:,] flattens the entire document. At times, we want to see the full documents, indexed by keys and docs does that. returns a dictable with both keys and the document in the 'doc' column

**exc (\*\*kwargs)**

filters 'negatively' removing documents that match the criteria specified.

**cursor**

filtered documents.

**Example:**

```
>>> from pyg import *; import pymongo
>>> table = pymongo.MongoClient()['test']['test']
>>> table.insert_one(dict(name = 'alan', surname = 'abrahams', age = 39, marriage = dt(2010, 1, 1)))
>>> table.insert_one(dict(name = 'barbara', surname = 'brown', age = 50, marriage = dt(2010, 1, 1)))
>>> table.insert_one(dict(name = 'charlie', surname = 'cohen', age = 20))
```

```
>>> t = mongo_reader(table)
>>> assert len(t.exc(name = 'alan')) == 2
>>> assert len(t.exc(name = ['alan', 'barbara'])) == 1
>>> table.drop()
```

**find(\*args, \*\*kwargs)**

Same as self.specify()

The 'spec' is the cursor's filter on documents (can think of it as row-selection) within the collection. We use q (see pyg.mongo.\_q.q) to specify the filter on the cursor.

**Returns:**

A filtered mongo\_reader cursor

**Example:**

```
>>> from pyg import *; import pymongo
>>> table = pymongo.MongoClient()['test']['test']
>>> table.insert_one(dict(name = 'alan', surname = 'abrahams', age = 39, marriage = dt(2010, 1, 1)))
>>> table.insert_one(dict(name = 'barbara', surname = 'brown', age = 50, marriage = dt(2010, 1, 1)))
>>> table.insert_one(dict(name = 'charlie', surname = 'cohen', age = 20))
```

```
>>> t = mongo_reader(table)
>>> assert len(t.find(name = 'alan')) == 1
>>> assert len(t.find(q.age>25)) == 2
>>> assert len(t.find(q.age>25, q.marriage<dt(2010))) == 1
```

```
>>> table.drop()
```

**find\_one(doc=None, \*args, \*\*kwargs)**

searches for records based either on the doc, or the args/kwargs specified. Unlike mongo cursor which finds one of many, here, when we ask for find\_one, we will throw an exception if more than one documents are found.

**Returns:**

A cursor pointing to a single record (document)

**inc(\*args, \*\*kwargs)**

Same as self.specify()

The 'spec' is the cursor's filter on documents (can think of it as row-selection) within the collection. We use q (see pyg.mongo.\_q.q) to specify the filter on the cursor.

**Returns:**

A filtered mongo\_reader cursor

**Example:**

```
>>> from pyg import *; import pymongo
>>> table = pymongo.MongoClient()['test']['test']
```

```
>>> table.insert_one(dict(name = 'alan', surname = 'abrahams', age = 39, marriage = dt(20
>>> table.insert_one(dict(name = 'barbara', surname = 'brown', age = 50, marriage = dt(20
>>> table.insert_one(dict(name = 'charlie', surname = 'cohen', age = 20))
```

```
>>> t = mongo_reader(table)
>>> assert len(t.find(name = 'alan')) == 1
>>> assert len(t.find(q.age>25)) == 2
>>> assert len(t.find(q.age>25, q.marriage<dt(2010))) == 1
```

```
>>> table.drop()
```

**project (projection=None)**

The 'projection' is the cursor's column selection on documents. If in SQL we write SELECT col1, col2 FROM ..., in Mongo, the cursor.projection = ['col1', 'col2']

**Parameters:**

projection: a list/str of keys we are interested in reading. Note that nested keys are OK: 'level1.level2.name' is perfectly good

**Returns:**

A mongo\_reader cursor filtered to read just these keys

**property projection****Returns:**

The 'projection' is the cursor's column selection on documents. If in SQL we write SELECT col1, col2 FROM ..., in Mongo, the cursor.projection = ['col1', 'col2']

**property raw**

returns an unfiltered mongo\_reader

**read (item=0, reader=None)**

reads the next document from the collection.

**Parameters:**

**item** : *int, optional*

Please read the ith record. The default is 0.

**reader** : *callable/list of callables, optional*

When we read the document from the collection, we first transform them. The default behaviour is to use pyg.base.\_encode.decode but you may pass reader = False to grab the raw data from mongo

**Returns:****document**

The document from Mongo

**sort (\*by)**

sorting on server side, per key(s)

by : str/list of strs

sorted cursor.

**property spec****Returns:**

The 'spec' is the cursor's filter on documents (can think of it as row-selection) within the collection

**specify (\*args, \*\*kwargs)**

The 'spec' is the cursor's filter on documents (can think of it as row-selection) within the collection. We use q (see pyg.mongo.\_q.q) to specify the filter on the cursor.

**Returns:**

A filtered mongo\_reader cursor

**mongo\_pk\_reader**

mongo\_pk\_reader extends the standard reader to handle tables with primary keys (pk) while being read-only.

`class pyg.mongo._pk_reader.mongo_pk_reader (cursor, pk, writer=None, reader=None, query=None, **_)`  
 we set up a system in Mongo to ensure we can mimin tables with primary keys. The way we do this is two folds:

- At document insertion, we mark as `_deleted` old documents sharing the same keys by adding a key `_deleted` to the old doc
- At reading, we filter for documents where `q._deleted.not_exists`.

`clone (**kwargs)`

**Returns:**

**mongo\_reader**

Returns a cloned version of current mongo\_reader but allows additional parameters to be set (see spec and project)

`create_index (*keys)`

creates a sorted index on the collection

**Parameters:**

**\*keys** : *strings*

if missing, use the primary keys.

`dedup ()`

Although in principle, if a single process reads/writes to Mongo, we should not get duplicates. In practice, when multiple clients access the database, we occasionally get multiple records with the same primary keys. When this happens, we also end up with poor mongo\_ids

**mongo\_pk\_cursor**

Hopefully, a table with unique keys.

`docs (doc='doc', *keys)`

`self[:,]` flattens the entire document. At times, we want to see the full documents, indexed by keys and docs does that. returns a dictable with both keys and the document in the 'doc' column

**mongo\_pk\_cursor**

mongo\_pk\_cursor is our go-to object and it manages all our primary-keyed tables. .. autoclass::  
 pyg.mongo.\_pk\_cursor.mongo\_pk\_cursor

**members:****encoding docs before saving to mongo**

Before we save data to Mongo, we may need to transform it, especially if we are to save `pd.DataFrame`. By default, we encode them into bytes and push to mongo. You can choose to save pandas dataframes/series as .parquet files and numpy arrays into .npy files.

**parquet\_write**

`pyg.mongo._encoders.parquet_write (doc, root=None)`

MongoDB is great for manipulating/searching dict keys/values. However, the actual dataframes in each doc, we may want to save in a file system. - The DataFrames are stored as bytes in MongoDB anyway, so they are not

searchable - Storing in files allows other non-python/non-MongoDB users easier access, allowing data to be detached from app - MongoDB free version has limitations on size of document - file based system may be faster, especially if saved locally not over network - for data licensing issues, data must not sit on servers but stored on local computer

Therefore, the doc encode will cycle through the elements in the doc. Each time it sees a pd.DataFrame/pd.Series, it will - determine where to write it (with the help of the doc) - save it to a .parquet file

## csv\_write

`pyg.mongo._encoders.csv_write(doc, root=None)`

MongoDB is great for manipulating/searching dict keys/values. However, the actual dataframes in each doc, we may want to save in a file system. - The DataFrames are stored as bytes in MongoDB anyway, so they are not searchable - Storing in files allows other non-python/non-MongoDB users easier access, allowing data to be detached from original application - MongoDB free version has limitations on size of document - file based system may be faster, especially if saved locally not over network - for data licensing issues, data must not sit on servers but stored on local computer

Therefore, the doc encode will cycle through the elements in the doc. Each time it sees a pd.DataFrame/pd.Series, it will - determine where to write it (with the help of the doc) - save it to a .csv file

## cells in Mongo

Now that we have a database, we construct cells that can load/save data to collections.

## db\_cell

`class pyg.mongo._db_cell.db_cell(function=None, output=None, db=None, **kwargs)`

a db\_cell is a specialized cell with a 'db' member pointing to a database where cell is to be stored. We use this to implement save/load for the cell.

It is important to recognize the duality in the design: - the job of the cell.db is to be able to save/load based on the primary keys. - the job of the cell is to provide the primary keys to the db object.

The cell saves itself by 'presenting' itself to cell.db() and say... go on, load my data based on my keys.

**Example:** saving & loading

```
>>> from pyg import *
>>> people = partial(mongo_table, db = 'test', table = 'test', pk = ['name', 'surname'])
>>> anna = db_cell(db = people, name = 'anna', surname = 'abramzon', age = 46).save()
>>> bob = db_cell(db = people, name = 'bob', surname = 'brown', age = 25).save()
>>> james = db_cell(db = people, name = 'james', surname = 'johnson', age = 39).save()
```

Now we can pull the data directly from the database

```
>>> people()['name', 'surname', 'age'][:, :]
>>> dictable[3 x 4]
>>> _id | age | name | surname
>>> 601e732e0ef13bec9cd8a6cb | 39 | james | johnson
>>> 601e73db0ef13bec9cd8a6d4 | 46 | anna | abramzon
>>> 601e73db0ef13bec9cd8a6d7 | 25 | bob | brown
```

db\_cell can implement a function:

```
>>> def is_young(age):
>>>     return age < 30
>>> bob.function = is_young
>>> bob = bob.go()
>>> assert bob.data is True
```

When run, it saves its new data to Mongo and we can load its own data:

```
>>> new_cell_with_just_db_and_keys = db_cell(db = people, name = 'bob', surname = 'brown')
>>> assert 'age' not in new_cell_with_just_db_and_keys
```

```
>>> now_with_the_data_from_database = new_cell_with_just_db_and_keys.load()
>>> assert now_with_the_data_from_database.age == 25
```

```
>>> people()['name', 'surname', 'age', 'data'][:, :]
>>> dictable[3 x 4]
>>> _id | age | name | surname | data
>>> 601e732e0ef13bec9cd8a6cb | 39 | james | johnson | None
>>> 601e73db0ef13bec9cd8a6d4 | 46 | anna | abramzon | None
>>> 601e73db0ef13bec9cd8a6d7 | 25 | bob | brown | True
>>> people().raw.drop()
```

**go (go=1, mode=0, \*\*kwargs)**

calculates the cell (if needed). By default, will then run cell.save() to save the cell. If you don't want to save the output (perhaps you want to check it first), use cell.\_go()

#### Parameters:

**go** : *int, optional*

a parameter that forces calculation. The default is 0. go = 0: calculate cell only if cell.run() is True go = 1: calculate THIS cell regardless. calculate the parents only if their cell.run() is True go = 2: calculate THIS cell and PARENTS cell regardless, calculate grandparents if cell.run() is True etc. go = -1: calculate the entire tree again.

**\*\*kwargs** : *parameters*

You can actually allocate the variables to the function at runtime

Note that by default, cell.go() will default to go = 1 and force a calculation on cell while cell() is lazy and will default to assuming go = 0

#### Returns:

**cell**

the cell, calculated

**Example:** different values of go

```
>>> from pyg import *
>>> f = lambda x=None,y=None: max([dt(x), dt(y)])
>>> a = cell(f)()
>>> b = cell(f, x = a)()
>>> c = cell(f, x = b)()
>>> d = cell(f, x = c)()
```

```
>>> e = d.go()
>>> e0 = d.go(0)
>>> e1 = d.go(1)
>>> e2 = d.go(2)
>>> e_1 = d.go(-1)
```

```
>>> assert not d.run() and e.data == d.data
>>> assert e0.data == d.data
>>> assert e1.data > d.data and e1.x.data == d.x.data
>>> assert e2.data > d.data and e2.x.data > d.x.data and e2.x.x.data == d.x.x.data
>>> assert e_1.data > d.data and e_1.x.data > d.x.data and e_1.x.x.data > d.x.x.data
```

**Example:** adding parameters on the run

```
>>> c = cell(lambda a, b: a+b)
>>> d = c(a = 1, b =2)
>>> assert d.data == 3
```

**load (mode=0)**

loads a document from the database and updates various keys.

**Persistency:**

Since we want to avoid hitting the database, there is a singleton GRAPH, a dict, storing the cells by their address. Every time we load/save from/to Mongo, we also update GRAPH.

We use the GRAPH often so if you want to FORCE the cell to go to the database when loading, use this:

```
>>> cell.load(-1)
>>> cell.load(-1).load(0) # clear GRAPH and load from db
>>> cell.load([0]) # same thing: clear GRAPH and then load if available
```

**Merge of  
cached cell  
and calling  
cell:**

Once we load from memory (either MongoDB or GRAPH), we `tree_update` the cached cell with the new values in the current cell. This means that it is next to impossible to actually *delete* keys. If you want to delete keys in a cell/cells in the database, you need to:

```
>>> del db.inc(filters)['key.subkey']
```

**Parameters:**

**mode** : *int* , *datetime*, *optional*

if -1, then does not load and clears the GRAPH if 0, then will load from database if found. If not found, will return original document if 1, then will throw an exception if no document is found in the database if mode is a date, will return the version alive at that date The default is 0.

IF you enclose any of these in a list, then GRAPH is cleared prior to running and the database is called.

**Returns:**

document

**save ()**

Saves the cell for persistency. Not implemented for simple cell. see `db_cell`

**Returns:**

**cell**

self, saved.

***periodic\_cell***

`class pyg.mongo._periodic_cell.periodic_cell (function=None, output=None, db=None, period='1b', updated=None, **kwargs)`

`periodic_cell` inherits from `db_cell` its ability to save itself in MongoDB using its `db` members Its calculation schedule depends on when it was last updated.

**Example:**

```
>>> from pyg import *
>>> c = periodic_cell(lambda a: a + 1, a = 0)
```

We now assert it needs to be calculated and calculate it...

```
>>> assert c.run()
>>> c = c.go()
>>> assert c.data == 1
>>> assert not c.run()
```

Now let us cheat and tell it, it was last run 3 days ago...

```
>>> c.updated = dt(-3)
>>> assert c.run()
```

**run()**

checks if the cell needs calculation. This depends on the nature of the cell. By default (for cell and db\_cell), if the cell is already calculated so that cell.\_output exists, then returns False. otherwise True

**bool**

run cell?

**Example:**

```
>>> c = cell(lambda x: x+1, x = 1)
>>> assert c.run()
>>> c = c()
>>> assert c.data == 2 and not c.run()
```

**get\_cell****db\_save**

pyg.mongo.\_db\_cell.**db\_save** (value)

saves a db\_cell from the database. Will iterates through lists and dicts

**Parameters:****value: obj**

db\_cell (or list/dict of) to be loaded

**Example:**

```
>>> from pyg import *
>>> db = partial(mongo_table, table = 'test', db = 'test', pk = ['a','b'])
>>> c = db_cell(add_, a = 2, b = 3, key = 'test', db = db)
>>> c = db_save(c)
>>> assert get_cell('test', 'test', a = 2, b = 3).key == 'test'
```

**db\_load**

pyg.mongo.\_db\_cell.**db\_load** (value, mode=0)

loads a db\_cell from the database. Iterates through lists and dicts

**Parameters:****value: obj**

db\_cell (or list/dict of) to be loaded

**mode: int**

loading mode -1: dont load, +1: load and throw an exception if not found, 0: load if found

**pyg.timeseries**

Given pandas, why do we need this timeseries library? pandas is amazing but there are a few features in pyg.timeseries designed to enhance it. There are three issues with pandas that pyg.timeseries tries to address:

- pandas works on pandas objects (obviously) but not on numpy arrays.
- **pandas handles TimeSeries with nan inconsistently across its functions. This makes your results sensitive to reindexing/resampling. E.g.:**
  - a.expanding() & a.ewm() **ignore** nan's for calculation and then ffill the result.
  - a.diff(), a.rolling() **include** any nans in the calculation, leading to nan propagation.



- pandas is great if you have the full timeseries. However, if you now want to run the same calculations in a live environment, on recent data, pandas cannot help you: you have to stick the new data at the end of the DataFrame and rerun.

pyg.timeseries tries to address this:

- pyg.timeseries agrees with pandas 100% on DataFrames (with no nan) while being of comparable (if not faster) speed
- pyg.timeseries works seamlessly on pandas objects and on numpy arrays, with no code change.
- pyg.timeseries handles nan consistently across all its functions, 'ignoring' all nan, making your results consistent regardless of resampling.
- **pyg.timeseries exposes the state of the internal function calculation. The exposure of internal states allows us to calculate the output of additional data without re-running history. This speeds up of two very common problems in finance:**
  - risk calculations, Monte Carlo scenarios: We can run a trading strategy up to today and then generate multiple scenarios and see what-if, without having to rerun the full history.
  - live versus history: pandas is designed to run a full historical simulation. However, once we reach "today", speed is of the essence and running a full historical simulation every time we ingest a new price, is just too slow. That is why most fast trading is built around fast state-machines. Of course, making sure research & live versions do the same thing is tricky. pyg gives you the ability to run two systems in parallel with almost the same code base: run full history overnight and then run today's code base instantly, instantiated with the output of the historical simulation.

## simple functions

### diff

pyg.timeseries.\_rolling.diff(a, n=1, axis=0, data=None, state=None)

equivalent to a.diff(n) in pandas if there are no nans. If there are, we SKIP nans rather than propagate them.

#### Parameters:

**a** : array/timeseries

array/timeseries

**n**: int, optional, default = 1

window size

**data**: None.

unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** : matching pandas no nan's

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> assert eq(timer(diff, 1000)(a), timer(lambda a, n=1: a.diff(n), 1000)(a))
```

**Example:** : nan skipping

```
>>> a = np.array([1., np.nan, 3., 9.])
>>> assert eq(diff(a), np.array([np.nan, np.nan, 2.0, 6.0]))
>>> assert eq(pd.Series(a).diff().values, np.array([np.nan, np.nan, np.nan, 6.0]))
```

### shift

pyg.timeseries.\_rolling.shift(a, n=1, axis=0, data=None, state=None)

Equivalent to a.shift() with support to arra

**Parameters:**

**a** : array, *pd.Series*, *pd.DataFrame* or list/dict of these  
timeseries

**n**: int

size of rolling window

**data**: None.

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:**

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series([1.,2,3,4,5], range(-4))
>>> assert eq(shift(a), pd.Series([np.nan,1,2,3,4], range(-4)))
>>> assert eq(shift(a,2), pd.Series([np.nan,np.nan,1,2,3], range(-4)))
>>> assert eq(shift(a,-1), pd.Series([2,3,4,5,np.nan], range(-4)))
```

**Example:** np.ndarrays

```
>>> assert eq(shift(a.values), shift(a).values)
```

**Example:** nan skipping

```
>>> a = pd.Series([1.,2,np.nan,3,4], range(-4))
>>> assert eq(shift(a), pd.Series([np.nan,1,np.nan, 2,3], range(-4)))
>>> assert eq(a.shift(), pd.Series([np.nan,1,2,np.nan,3], range(-4))) # the location of t
```

**Example:** state management

```
>>> old = a.iloc[:3]
>>> new = a.iloc[3:]
>>> old_ts = shift_(old)
>>> new_ts = shift(new, **old_ts)
>>> assert eq(new_ts, shift(a).iloc[3:])
```

**ratio**

`pyg.timeseries._rolling.ratio` (a, n=1, data=None, state=None)

Equivalent to `a.diff()` but in log-space..

**Parameters:**

**a** : array, *pd.Series*, *pd.DataFrame* or list/dict of these  
timeseries

**n**: int

size of rolling window

**data**: None.

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:**

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series([1.,2,3,4,5], range(-4))
>>> assert eq(ratio(a), pd.Series([np.nan, 2, 1.5, 4/3,1.25], range(-4)))
>>> assert eq(ratio(a,2), pd.Series([np.nan, np.nan, 3, 2, 5/3], range(-4)))
```

***ts\_count***

`pyg.timeseries._ts.ts_count(a)` is equivalent to `a.count()` (though slightly slower)

- supports numpy arrays
- skips nan
- supports state management

**Example:** pandas matching

```
>>> # create sample data:
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999)); a[a>0] = np.nan
>>> assert ts_count(a) == a.count()
```

**Example:** numpy

```
>>> assert ts_count(a.values) == ts_count(a)
```

**Example:** state management

```
>>> old = ts_count_(a.iloc[:2000])
>>> new = ts_count(a.iloc[2000:], state = old.state)
>>> assert new == ts_count(a)
```

***ts\_sum***

`pyg.timeseries._ts.ts_sum(a)` is equivalent to `a.sum()`

- supports numpy arrays
- handles nan
- supports state management

**Parameters:**

**a** : *array, pd.Series, pd.DataFrame or list/dict of these timeseries*

**axis** : *int, optional*  
0/1/-1. The default is 0.

**data**: **None**

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: **dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** pandas matching

```
>>> # create sample data:
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999)); a[a>0] = np.nan
>>> assert ts_sum(a) == a.sum()
```

**Example:** numpy

```
>>> assert ts_sum(a.values) == ts_sum(a)
```

**Example:** state management

```
>>> old = ts_sum_(a.iloc[:2000])
>>> new = ts_sum(a.iloc[2000:], vec = old.vec)
>>> assert new == ts_sum(a)
```

## ts\_mean

pyg.timeseries.\_ts.**ts\_mean**(a) is equivalent to a.mean()

- supports numpy arrays
- handles nan
- supports state management
- pandas is actually faster on count

### Parameters:

**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**axis** : int, optional

0/1/-1. The default is 0.

**data**: None

unused at the moment. Allow code such as func(live, **\*\*func\_**(history)) to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** pandas matching

```
>>> # create sample data:
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999)); a[a>0] = np.nan
>>> assert ts_mean(a) == a.mean()
```

**Example:** numpy

```
>>> assert ts_mean(a.values) == ts_mean(a)
```

**Example:** state management

```
>>> old = ts_mean_(a.iloc[:2000])
>>> new = ts_mean(a.iloc[2000:], vec = old.vec)
>>> assert new == ts_mean(a)
```

## ts\_rms

pyg.timeseries.\_ts.**ts\_rms**(a, axis=0, data=None, state=None)

ts\_rms(a) is equivalent to (a\*\*2).mean()<sup>0.5</sup>

### Parameters:

**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**axis** : int, optional

0/1/-1. The default is 0.

**data**: None

unused at the moment. Allow code such as func(live, **\*\*func\_**(history)) to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

- supports numpy arrays
- handles nan
- supports state management

```
>>> # create sample data:
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999)); a[a>0] = np.nan
```

**Example:** pandas matching

```
>>> assert abs(ts_std(a) - a.std())<1e-13
```

**Example:** numpy

```
>>> assert ts_std(a.values) == ts_std(a)
```

**Example:** state management

```
>>> old = ts_rms_(a.iloc[:2000])
>>> new = ts_rms(a.iloc[2000:], vec = old.vec)
>>> assert new == ts_rms(a)
```

## ts\_std

pyg.timeseries.\_ts.**ts\_std**(a) is equivalent to a.std()

- supports numpy arrays
- handles nan
- supports state management

```
>>> # create sample data:
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999)); a[a>0] = np.nan
```

**Example:** pandas matching

```
>>> assert abs(ts_std(a) - a.std())<1e-13
```

**Example:** numpy

```
>>> assert ts_std(a.values) == ts_std(a)
```

**Example:** state management

```
>>> old = ts_std_(a.iloc[:2000])
>>> new = ts_std(a.iloc[2000:], vec = old.vec)
>>> assert new == ts_std(a)
```

## ts\_skew

pyg.timeseries.\_ts.**ts\_skew**(a, 0) is equivalent to a.skew()

- supports numpy arrays
- handles nan
- faster than pandas
- supports state management

**Parameters:**

**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**axis** : int, optional  
0/1/-1. The default is 0.

**min\_sample**: float, optional

This refers to the denominator when we calculate the skew. Over time, the denominator converges to 1 but initially, it is small. Also, if there is a gap in the data, older datapoints weight may have decayed while there are not enough “new point”. min\_sample ensures that in both cases, if denominator < 0.25 (default value) we return nan.

**data**: None

unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to ‘a’ provided.

**Example:** pandas matching

```
>>> # create sample data:
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999)); a[a>0] = np.nan
>>> assert abs(ts_skew(a, 0) - a.skew())<1e-13
```

**Example:** numpy

```
>>> assert ts_skew(a.values) == ts_skew(a)
```

**Example:** state management

```
>>> old = ts_skew_(a.iloc[:2000])
>>> new = ts_skew(a.iloc[2000:], vec = old.vec)
>>> assert new == ts_skew(a)
```

## ts\_min

pyg.timeseries.\_ts.**ts\_max**(a) is equivalent to pandas a.min()

## ts\_max

pyg.timeseries.\_ts.**ts\_max**(a) is equivalent to pandas a.min()

## ts\_median

pyg.timeseries.\_ts.**ts\_median**(a, axis=0)

## fnna

pyg.timeseries.\_rolling.**fnna**(a, n=1, axis=0)  
returns the index in a of the nth first non-nan.

**Parameters:**

a : array/timeseries n: int, optional, default = 1

**Example:**

```
>>> a = np.array([np.nan, np.nan, 1, np.nan, np.nan, 2, np.nan, np.nan, np.nan])
>>> fnna(a, n=-2)
```

## v2na/na2v

`pyg.timeseries._rolling.v2na` (*a*, *old*=0.0, *new*=nan)  
 replaces an old value with a new value (default is nan)

#### Examples:

```
>>> from pyg import *
>>> a = np.array([1., np.nan, 1., 0.])
>>> assert eq(v2na(a), np.array([1., np.nan, 1., np.nan]))
>>> assert eq(v2na(a,1), np.array([np.nan, np.nan, np.nan, 0]))
>>> assert eq(v2na(a,1,0), np.array([0., np.nan, 0., 0.]))
```

#### Parameters:

*a* : array/timeseries *old*: float

value to be replaced

**new** : float, optional

new value to be used, The default is np.nan.

#### Returns:

array/timeseries

`pyg.timeseries._rolling.na2v` (*a*, *new*=0.0)  
 replaces a nan with a new value

#### Example:

```
>>> from pyg import *
>>> a = np.array([1., np.nan, 1.])
>>> assert eq(na2v(a), np.array([1., 0.0, 1.]))
>>> assert eq(na2v(a,1), np.array([1., 1., 1.]))
```

#### Parameters:

*a* : array/timeseries *new* : float, optional

DESCRIPTION. The default is 0.0.

#### Returns:

array/timeseries

## ffill/bfill

`pyg.timeseries._rolling.ffill` (*a*, *n*=0, *axis*=0, *data*=None, *state*=None)  
 returns a forward filled array, up to *n* values forward. supports state management which is needed if we want only *nth*

#### Parameters:

**a** : array/timeseries

array/timeseries

**n**: int, optional, default = 1

window size

**data**: None.

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

#### Example:

```
>>> a = np.array([np.nan, np.nan, 1, np.nan, np.nan, 2, np.nan, np.nan, np.nan])
>>> fnna(a, n=-2)
```

`pyg.timeseries._rolling.bfill` (*a*, *n*=-1, *axis*=0)  
 equivalent to `a.fillna('bfill')`. There is no state-aware as this function is forward looking

**Example:**

```
>>> from pyg import *
>>> a = np.array([np.nan, 1., np.nan])
>>> b = np.array([1., 1., np.nan])
>>> assert eq(bfill(a), b)
```

**Example:** `pd.Series`

```
>>> ts = pd.Series(a, range(-2))
>>> assert eq(bfill(ts).values, b)
```

***nona***

`pyg.timeseries._ts.nona` (*a*, *value*=nan)  
 removes rows that are entirely nan (or a specific other value)

**Parameters:**

*a* : dataframe/ndarray

**value** : *float, optional*

value to be removed. The default is `np.nan`.

**Example:**

```
>>> from pyg import *
>>> a = np.array([1,np.nan,2,3])
>>> assert eq(nona(a), np.array([1,2,3]))
```

**Example:** multiple columns

```
>>> a = np.array([[1,np.nan,2,np.nan], [np.nan, np.nan, np.nan, 3]]).T
>>> b = np.array([[1,2,np.nan], [np.nan, np.nan, 3]]).T ## 2nd row has nans across
>>> assert eq(nona(a), b)
```

***expanding window functions******expanding\_mean***

`pyg.timeseries._expanding.expanding_mean` (*a*, *axis*=0, *data*=None, *state*=None)  
 equivalent to `pandas a.expanding().mean()`.

- works with `np.arrays`
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a** : *array, pd.Series, pd.DataFrame or list/dict of these*  
 timeseries

**axis** : *int, optional*

0/1/-1. The default is 0.

**data**: **None.**

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: **dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.



**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.expanding().mean(); ts = expanding_mean(a)
>>> assert eq(ts,panda)
```

**Example:** nan handling

Unlike pandas, timeseries does not forward fill the nans.

```
>>> a[a<0.1] = np.nan
>>> panda = a.expanding().mean(); ts = expanding_mean(a)
```

```
>>> pd.concat([panda,ts], axis=1)
>>>
>>> 1993-09-23  1.562960  1.562960
>>> 1993-09-24  0.908910  0.908910
>>> 1993-09-25  0.846817  0.846817
>>> 1993-09-26  0.821423  0.821423
>>> 1993-09-27  0.821423      NaN
>>>
>>> 2021-02-03  0.870358  0.870358
>>> 2021-02-04  0.870358      NaN
>>> 2021-02-05  0.870358      NaN
>>> 2021-02-06  0.870358      NaN
>>> 2021-02-07  0.870353  0.870353
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = expanding_mean(a)
>>> old_ts = expanding_mean_(old)
>>> new_ts = expanding_mean(new, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(expanding_mean(dict(x = a, y = a**2)), dict(x = expanding_mean(a), y = expanding_mean(a**2)))
>>> assert eq(expanding_mean([a,a**2]), [expanding_mean(a), expanding_mean(a**2)])
```

**expanding\_rms**

pyg.timeseries.\_expanding.**expanding\_rms** (a, axis=0, data=None, state=None)  
equivalent to pandas (a\*\*2).expanding().mean()\*\*0.5). - works with np.arrays - handles nan without forward filling. - supports state parameters

**Parameters:**

**a** : array, pd.Series, pd.DataFrame, list/dict of these timeseries

**axis** : int, optional

0/1/-1. The default is 0.

**data**: None.

unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = (a**2).expanding().mean()**0.5; ts = expanding_rms(a)
>>> assert eq(ts,panda)
```

**Example:** nan handling

Unlike pandas, timeseries does not forward fill the nans.

```
>>> a[a<0.1] = np.nan
>>> panda = (a**2).expanding().mean()**0.5; ts = expanding_rms(a)
```

```
>>> pd.concat([panda,ts], axis=1)
>>>
      0      1
>>> 1993-09-23  0.160462  0.160462
>>> 1993-09-24  0.160462      NaN
>>> 1993-09-25  0.160462      NaN
>>> 1993-09-26  0.160462      NaN
>>> 1993-09-27  0.160462      NaN
>>>
      ...      ...
>>> 2021-02-03  1.040346  1.040346
>>> 2021-02-04  1.040346      NaN
>>> 2021-02-05  1.040338  1.040338
>>> 2021-02-06  1.040337  1.040337
>>> 2021-02-07  1.040473  1.040473
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = expanding_rms(a)
>>> old_ts = expanding_rms_(old)
>>> new_ts = expanding_rms(new, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(expanding_rms(dict(x = a, y = a**2)), dict(x = expanding_rms(a), y = expanding_rms(a**2)))
>>> assert eq(expanding_rms([a,a**2]), [expanding_rms(a), expanding_rms(a**2)])
```

***expanding\_std***

pyg.timeseries.\_expanding.**expanding\_std** (a, axis=0, data=None, state=None)  
equivalent to pandas a.expanding().std().

- works with np.arrays
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a** : array, pd.Series, pd.DataFrame or list/dict of these  
timeseries

**axis** : int, optional

0/1/-1. The default is 0.

**data**: None.

unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state: dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.expanding().std(); ts = expanding_std(a)
>>> assert abs(ts-panda).max()<1e-10
```

**Example:** nan handling

Unlike pandas, timeseries does not forward fill the nans.

```
>>> a[a<0.1] = np.nan
>>> panda = a.expanding().std(); ts = expanding_std(a)
```

```
>>> pd.concat([panda,ts], axis=1)
>>>
      0      1
>>> 1993-09-23    NaN    NaN
>>> 1993-09-24    NaN    NaN
>>> 1993-09-25    NaN    NaN
>>> 1993-09-26    NaN    NaN
>>> 1993-09-27    NaN    NaN
>>>
      ...    ...
>>> 2021-02-03  0.590448  0.590448
>>> 2021-02-04  0.590448    NaN
>>> 2021-02-05  0.590475  0.590475
>>> 2021-02-06  0.590475    NaN
>>> 2021-02-07  0.590411  0.590411
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = expanding_std(a)
>>> old_ts = expanding_std(old)
>>> new_ts = expanding_std(new, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(expanding_std(dict(x = a, y = a**2)), dict(x = expanding_std(a), y = expanding_std(a**2)))
>>> assert eq(expanding_std([a,a**2]), [expanding_std(a), expanding_std(a**2)])
```

**expanding\_sum**

pyg.timeseries.\_expanding.**expanding\_sum** (a, axis=0, data=None, state=None)  
equivalent to pandas a.expanding().sum().

- works with np.arrays
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**axis** : int, optional

0/1/-1. The default is 0.

**data: None**

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state: dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.expanding().sum(); ts = expanding_sum(a)
>>> assert eq(ts,panda)
```

**Example:** nan handling

Unlike pandas, timeseries does not forward fill the nans.

```
>>> a[a<0.1] = np.nan
>>> panda = a.expanding().sum(); ts = expanding_sum(a)
```

```
>>> pd.concat([panda,ts], axis=1)
>>>
      0      1
1993-09-23    NaN    NaN
1993-09-24    NaN    NaN
1993-09-25    0.645944    0.645944
1993-09-26    2.816321    2.816321
1993-09-27    2.816321    NaN
>>>
      ...      ...
2021-02-03  3976.911348  3976.911348
2021-02-04  3976.911348    NaN
2021-02-05  3976.911348    NaN
2021-02-06  3976.911348    NaN
2021-02-07  3976.911348    NaN
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = expanding_sum(a)
>>> old_ts = expanding_sum_(old)
>>> new_ts = expanding_sum(new, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(expanding_sum(dict(x = a, y = a**2)), dict(x = expanding_sum(a), y = expanding_sum(a**2)))
>>> assert eq(expanding_sum([a,a**2]), [expanding_sum(a), expanding_sum(a**2)])
```

**expanding\_skew**

`pyg.timeseries._expanding.expanding_skew` (a, bias=False, axis=0, data=None, state=None)  
equivalent to `pandas a.expanding().skew()` which doesn't exist

- works with `np.arrays`
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a** : *array, pd.Series, pd.DataFrame or list/dict of these timeseries*

**axis** : *int, optional*

0/1/-1. The default is 0.

**data**: **None**.

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: **dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
```

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = expanding_skew(a)
>>> old_ts = expanding_skew_(old)
>>> new_ts = expanding_skew(new, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(expanding_skew(dict(x = a, y = a**2)), dict(x = expanding_skew(a), y = expanding_skew(a**2)))
>>> assert eq(expanding_skew([a,a**2]), [expanding_skew(a), expanding_skew(a**2)])
```

## expanding\_min

`pyg.timeseries._min.expanding_min(a, axis=0, data=None, state=None)`

equivalent to `pandas a.expanding().min()`.

- works with `np.arrays`
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a** : *array, pd.Series, pd.DataFrame or list/dict of these timeseries*

**axis** : *int, optional*

0/1/-1. The default is 0.

**data**: **None**.

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: **dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.expanding().min(); ts = expanding_min(a)
>>> assert eq(ts,panda)
```

**Example:** nan handling

Unlike pandas, timeseries does not forward fill the nans.

```
>>> a[a<0.1] = np.nan
>>> panda = a.expanding().min(); ts = expanding_min(a)
```

```
>>> pd.concat([panda,ts], axis=1)
>>>
>>> 1993-09-24      NaN      NaN
>>> 1993-09-25      NaN      NaN
>>> 1993-09-26  0.775176  0.775176
>>> 1993-09-27  0.691942  0.691942
>>> 1993-09-28  0.691942      NaN
>>>
>>> 2021-02-04  0.100099  0.100099
>>> 2021-02-05  0.100099      NaN
>>> 2021-02-06  0.100099      NaN
>>> 2021-02-07  0.100099  0.100099
>>> 2021-02-08  0.100099  0.100099
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = expanding_min(a)
>>> old_ts = expanding_min_(old)
>>> new_ts = expanding_min(new, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(expanding_min(dict(x = a, y = a**2)), dict(x = expanding_min(a), y = expanding_min(a**2)))
>>> assert eq(expanding_min([a,a**2]), [expanding_min(a), expanding_min(a**2)])
```

**expanding\_max**

`pyg.timeseries._max.expanding_max` (`a`, `axis=0`, `data=None`, `state=None`)  
equivalent to `pandas a.expanding().max()`.

- works with `np.arrays`
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a** : *array, pd.Series, pd.DataFrame or list/dict of these timeseries*

**axis** : *int, optional*

0/1/-1. The default is 0.

**data**: **None**.

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: **dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.expanding().max(); ts = expanding_max(a)
>>> assert eq(ts,panda)
```

**Example:** nan handling

Unlike pandas, timeseries does not forward fill the nans.

```
>>> a[a<0.1] = np.nan
>>> panda = a.expanding().max(); ts = expanding_max(a)
```

```
>>> pd.concat([panda,ts], axis=1)
>>>
      0      1
>>> 1993-09-24    NaN    NaN
>>> 1993-09-25    NaN    NaN
>>> 1993-09-26  0.875409  0.875409
>>> 1993-09-27  0.875409    NaN
>>> 1993-09-28  0.875409    NaN
>>>
      ...      ...
>>> 2021-02-04  3.625858  3.625858
>>> 2021-02-05  3.625858    NaN
>>> 2021-02-06  3.625858  3.625858
>>> 2021-02-07  3.625858    NaN
>>> 2021-02-08  3.625858    NaN
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = expanding_max(a)
>>> old_ts = expanding_max_(old)
>>> new_ts = expanding_max(new, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(expanding_max(dict(x = a, y = a**2)), dict(x = expanding_max(a), y = expanding_max(a**2)))
>>> assert eq(expanding_max([a,a**2]), [expanding_max(a), expanding_max(a**2)])
```

**expanding\_median**

pyg.timeseries.\_median.**expanding\_median** (a, axis=0)  
equivalent to pandas a.expanding().median().

- works with np.arrays
- handles nan without forward filling.
- There is no state-aware version since this requires essentially the whole history to be stored.

**Parameters:**

**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**axis** : int, optional

0/1/-1. The default is 0.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.expanding().median(); ts = expanding_median(a)
>>> assert eq(ts,panda)
```

**Example:** nan handling

Unlike pandas, timeseries does not forward fill the nans.

```
>>> a[a<0.1] = np.nan
>>> panda = a.expanding().median(); ts = expanding_median(a)
```

```
>>> pd.concat([panda,ts], axis=1)
>>>
      0      1
>>> 1993-09-23  1.562960  1.562960
>>> 1993-09-24  0.908910  0.908910
>>> 1993-09-25  0.846817  0.846817
>>> 1993-09-26  0.821423  0.821423
>>> 1993-09-27  0.821423      NaN
>>>
      ...      ...
>>> 2021-02-03  0.870358  0.870358
>>> 2021-02-04  0.870358      NaN
>>> 2021-02-05  0.870358      NaN
>>> 2021-02-06  0.870358      NaN
>>> 2021-02-07  0.870353  0.870353
```

**Example:** dict/list inputs

```
>>> assert eq(expanding_median(dict(x = a, y = a**2)), dict(x = expanding_median(a), y = e
>>> assert eq(expanding_median([a,a**2]), [expanding_median(a), expanding_median(a**2)])
```

## expanding\_rank

pyg.timeseries.\_rank.**expanding\_rank** (a, axis=0)

returns a rank of the current value within history, scaled to be -1 if it is the smallest and +1 if it is the largest - works on mumpy arrays too - skips nan, no ffill

### Parameters:

**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**axis** : int, optional

0/1/-1. The default is 0.

### Example:

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series([1.,2., np.nan, 0.,4.,2.], drange(-5))
>>> rank = expanding_rank(a)
>>> assert eq(rank, pd.Series([0, 1, np.nan, -1, 1, 0.25], drange(-5)))
>>> #
>>> # 2 is largest in [1,2] so goes to 1;
>>> # 0 is smallest in [1,2,0] so goes to -1 etc.
```

**Example:** numpy equivalent

```
>>> assert eq(expanding_rank(a.values), expanding_rank(a).values)
```

## cumsum

pyg.timeseries.\_expanding.**cumsum** (a, axis=0, data=None, state=None)

equivalent to pandas a.expanding().sum().

- works with np.arrays
- handles nan without forward filling.
- supports state parameters

### Parameters:



**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**axis** : int, optional  
0/1/-1. The default is 0.

**data**: None  
unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state**: dict, optional  
state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.expanding().sum(); ts = expanding_sum(a)
>>> assert eq(ts,panda)
```

**Example:** nan handling

Unlike pandas, timeseries does not forward fill the nans.

```
>>> a[a<0.1] = np.nan
>>> panda = a.expanding().sum(); ts = expanding_sum(a)
```

```
>>> pd.concat([panda,ts], axis=1)
>>>
>>> 1993-09-23      NaN      NaN
>>> 1993-09-24      NaN      NaN
>>> 1993-09-25    0.645944    0.645944
>>> 1993-09-26    2.816321    2.816321
>>> 1993-09-27    2.816321      NaN
>>>
>>> 2021-02-03  3976.911348  3976.911348
>>> 2021-02-04  3976.911348      NaN
>>> 2021-02-05  3976.911348      NaN
>>> 2021-02-06  3976.911348      NaN
>>> 2021-02-07  3976.911348      NaN
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = expanding_sum(a)
>>> old_ts = expanding_sum_(old)
>>> new_ts = expanding_sum(new, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(expanding_sum(dict(x = a, y = a**2)), dict(x = expanding_sum(a), y = expanding_sum(a**2)))
>>> assert eq(expanding_sum([a,a**2]), [expanding_sum(a), expanding_sum(a**2)])
```

## cumprod

pyg.timeseries.\_expanding.cumprod(a, axis=0, data=None, state=None)  
equivalent to pandas np.exp(np.log(a).expanding().sum()).

- works with np.arrays
- handles nan without forward filling.

- supports state parameters

**Parameters:**

**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**axis** : int, optional

0/1/-1. The default is 0.

**data**: None

unused at the moment. Allow code such as func(live, **\*\*func\_**(history)) to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = 1 + pd.Series(np.random.normal(0.001,0.05,10000), drange(-9999))
>>> panda = np.exp(np.log(a).expanding().sum()); ts = cumprod(a)
>>> assert abs(ts-panda).max() < 1e-10
```

**Example:** nan handling

Unlike pandas, timeseries does not forward fill the nans.

```
>>> a = 1 + pd.Series(np.random.normal(-0.01,0.05,100), drange(-99, 2020))
>>> a[a<0.975] = np.nan
>>> panda = np.exp(np.log(a).expanding().sum()); ts = cumprod(a)
```

```
>>> pd.concat([panda,ts], axis=1)
>>> 2019-09-24  1.037161  1.037161
>>> 2019-09-25  1.050378  1.050378
>>> 2019-09-26  1.158734  1.158734
>>> 2019-09-27  1.158734      NaN
>>> 2019-09-28  1.219402  1.219402
>>> ...
>>> 2019-12-28  4.032919  4.032919
>>> 2019-12-29  4.032919      NaN
>>> 2019-12-30  4.180120  4.180120
>>> 2019-12-31  4.180120      NaN
>>> 2020-01-01  4.244261  4.244261
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:50]
>>> new = a.iloc[50:]
>>> ts = cumprod(a)
>>> old_ts = cumprod_(old)
>>> new_ts = cumprod(new, **old_ts)
>>> assert eq(new_ts, ts.iloc[50:])
```

**Example:** dict/list inputs

```
>>> assert eq(cumprod(dict(x = a, y = a**2)), dict(x = cumprod(a), y = cumprod(a**2)))
>>> assert eq(cumprod([a,a**2]), [cumprod(a), cumprod(a**2)])
```

## rolling window functions

### rolling\_mean

`pyg.timeseries._rolling.rolling_mean(a, n, time=None, axis=0, data=None, state=None)`  
 equivalent to `pandas a.rolling(n).mean()`.

- works with `np.arrays`
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a :** *array, pd.Series, pd.DataFrame or list/dict of these timeseries*

**n: int**  
 size of rolling window

**time: a sequence of rising values of time**  
 passage of time

**axis :** *int, optional*  
 0/1/-1. The default is 0.

**data: None.**  
 unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state: dict, optional**  
 state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.rolling(10).mean(); ts = rolling_mean(a,10)
>>> assert abs(ts-panda).max()<1e-10
```

**Example:** nan handling

Unlike pandas, timeseries does not include the nans in the rolling calculation: it skips them. Since pandas rolling engine does not skip nans, they propagate. In fact, having removed half the data points, `rolling(10)` will return 99% of nans

```
>>> a[a<0.1] = np.nan
>>> panda = a.rolling(10).mean(); ts = rolling_mean(a,10)
>>> print('#original:', len(nona(a)), 'timeseries:', len(nona(ts)), 'panda:', len(nona(panda)))
>>> #original: 4534 timeseries: 4525 panda: 6 data points
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = rolling_mean(a,10)
>>> old_ts = rolling_mean_(old,10)
>>> new_ts = rolling_mean(new, 10, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(rolling_mean(dict(x = a, y = a**2),10), dict(x = rolling_mean(a,10), y = rolling_mean(a**2,10)))
>>> assert eq(rolling_mean([a,a**2],10), [rolling_mean(a,10), rolling_mean(a**2,10)])
```

**Example:** passage of time

```
>>> a = np.array([1.,2.,3.,4.])
>>> time = np.array([0,0,1,1]) ## i.e. the first two observations are from day 0 and the next two from day 1
```

```
>>> rolling_mean(a,2,time)
>>> array([nan, nan, 2.5, 3. ])
```

## The first two observations are from day 0 so cannot have a mean until we get a second point ## The mean is then calculated from last observation in day 0 (i.e. 2) and then 3. and then with 4. since these are again, from same day

## ***rolling\_rms***

pyg.timeseries.\_rolling.rolling\_rms (a, n, time=None, axis=0, data=None, state=None)  
equivalent to pandas (a\*\*2).rolling(n).mean()<sup>2</sup>0.5.

- works with np.arrays
- handles nan without forward filling.
- supports state parameters

### **Parameters:**

**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**n**: int

size of rolling window

**time**: a sequence of rising values of time

passage of time

**axis** : int, optional

0/1/-1. The default is 0.

**data**: None.

unused at the moment. Allow code such as func(live, <sup>2</sup>func\_(history)) to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = (a**2).rolling(10).mean()20.5; ts = rolling_rms(a,10)
>>> assert abs(ts-panda).max()<1e-10
```

**Example:** nan handling

Unlike pandas, timeseries does not include the nans in the rolling calculation: it skips them. Since pandas rolling engine does not skip nans, they propagate. In fact, having removed half the data points, rolling(10) will return 99% of nans

```
>>> a[a<0.1] = np.nan
>>> panda = (a**2).rolling(10).mean()20.5; ts = rolling_rms(a,10)
>>> print('#original:', len(nona(a)), 'timeseries:', len(nona(ts)), 'panda:', len(nona(panda)))
>>> #original: 4534 timeseries: 4525 panda: 6 data points
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = rolling_rms(a,10)
>>> old_ts = rolling_rms(old,10)
>>> new_ts = rolling_rms(new, 10, 2old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(rolling_rms(dict(x = a, y = a**2),10), dict(x = rolling_rms(a,10), y = rolling_rms(a**2,10)))
>>> assert eq(rolling_rms([a,a**2],10), [rolling_rms(a,10), rolling_rms(a**2,10)])
```

**Example:** passage of time

```
>>> a = np.array([1.,2.,3.,4.])
>>> time = np.array([0,0,1,1]) ## i.e. the first two observations are from day 0 and the next two from day 1
>>> (rolling_rms(a,2,time) ** 2) * 2
>>> array([nan, nan, 13., 20.]) == array([nan, nan, 4 + 9, 4 + 16])
```

## The first two observations are from day 0 so cannot have a mean until we get a second point ## The rms is then calculated from last observation in day 0 (i.e. 2) and then 3. and then with 4. since these are again, from same day

**rolling\_std**

pyg.timeseries.\_rolling.rolling\_std (a, n, time=None, axis=0, data=None, state=None)  
equivalent to pandas a.rolling(n).std().

- works with np.arrays
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a** : array, pd.Series, pd.DataFrame or list/dict of these  
timeseries

**n**: int  
size of rolling window

**time**: a sequence of rising values of time  
passage of time

**axis** : int, optional  
0/1/-1. The default is 0.

**data**: None.  
unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state**: dict, optional  
state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.rolling(10).std(); ts = rolling_std(a,10)
>>> assert abs(ts-panda).max()<1e-10
```

**Example:** nan handling

Unlike pandas, timeseries does not include the nans in the rolling calculation: it skips them. Since pandas rolling engine does not skip nans, they propagate. In fact, having removed half the data points, rolling(10) will return 99.9% nans

```
>>> a[a<0.1] = np.nan
>>> panda = a.rolling(10).std(); ts = rolling_std(a,10)
>>> print('#original:', len(nona(a)), 'timeseries:', len(nona(ts)), 'panda:', len(nona(panda)))
>>> #original: 4534 timeseries: 4525 panda: 2 data points
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = rolling_std(a,10)
>>> old_ts = rolling_std_(old,10)
>>> new_ts = rolling_std(new, 10, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(rolling_std(dict(x = a, y = a**2),10), dict(x = rolling_std(a,10), y = rolling_std(a**2,10)))
>>> assert eq(rolling_std([a,a**2],10), [rolling_std(a,10), rolling_std(a**2,10)])
```

**Example:** passage of time

```
>>> a = np.array([1.,2.,3.,4.])
>>> time = np.array([0,0,1,1]) ## i.e. the first two observations are from day 0 and the next two from day 1
>>> 2*(rolling_std(a,2,time) ** 2)
>>> array([nan, nan, 1., 4.])
```

## The first two observations are from day 0 so cannot have a mean until we get a second point ## The sum is then calculated from last observation in day 0 (i.e. 2) and then 3. and then with 4. since these are again, from same day

## rolling\_sum

pyg.timeseries.\_rolling.**rolling\_sum** (a, n, time=None, axis=0, data=None, state=None)  
equivalent to pandas a.rolling(n).sum().

- works with np.arrays
- handles nan without forward filling.
- supports state parameters

### Parameters:

**a :** array, pd.Series, pd.DataFrame or list/dict of these timeseries

**n: int**  
size of rolling window

**time: a sequence of rising values of time**  
passage of time

**axis :** int, optional  
0/1/-1. The default is 0.

**data: None.**  
unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state: dict, optional**  
state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.rolling(10).sum(); ts = rolling_sum(a,10)
>>> assert abs(ts-panda).max()<1e-10
```

**Example:** nan handling

Unlike pandas, timeseries does not include the nans in the rolling calculation: it skips them. Since pandas rolling engine does not skip nans, they propagate. In fact, having removed half the data points, rolling(10) will return 99.9% nans

```
>>> a[a<0.1] = np.nan
>>> panda = a.rolling(10).sum(); ts = rolling_sum(a,10)
>>> print('#original:', len(nona(a)), 'timeseries:', len(nona(ts)), 'panda:', len(nona(panda)))
>>> #original: 4534 timeseries: 4525 panda: 2 data points
```

#### Example: state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = rolling_sum(a,10)
>>> old_ts = rolling_sum_(old,10)
>>> new_ts = rolling_sum(new, 10, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

#### Example: dict/list inputs

```
>>> assert eq(rolling_sum(dict(x = a, y = a**2),10), dict(x = rolling_sum(a,10), y = rolling_sum(a**2,10)))
>>> assert eq(rolling_sum([a,a**2],10), [rolling_sum(a,10), rolling_sum(a**2,10)])
```

#### Example: passage of time

```
>>> a = np.array([1.,2.,3.,4.])
>>> time = np.array([0,0,1,1]) ## i.e. the first two observations are from day 0 and the next two from day 1
>>> rolling_sum(a,2,time)
>>> array([nan, nan, 5., 6.]) == array([nan, nan, 2+3, 2+4])
```

## The first two observations are from day 0 so cannot have a mean until we get a second point ## The sum is then calculated from last observation in day 0 (i.e. 2) and then 3. and then with 4. since these are again, from same day

## rolling\_skew

pyg.timeseries.\_rolling.rolling\_skew(a, n, bias=False, time=None, axis=0, data=None, state=None)  
equivalent to pandas a.rolling(n).skew().

- works with np.arrays
- handles nan without forward filling.
- supports state parameters

#### Parameters:

**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**n**: int  
size of rolling window

**time**: a sequence of rising values of time  
passage of time

**bias**:  
affects the skew calculation definition, see scipy documentation for details.

**axis** : int, optional  
0/1/-1. The default is 0.

**data**: None.  
unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state: dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.rolling(10).skew(); ts = rolling_skew(a,10)
>>> assert abs(ts-panda).max()<1e-10
```

**Example:** nan handling

Unlike pandas, timeseries does not include the nans in the rolling calculation: it skips them. Since pandas rolling engine does not skip nans, they propagate. In fact, having removed half the data points, rolling(10) will return 99.9% nans

```
>>> a[a<0.1] = np.nan
>>> panda = a.rolling(10).skew(); ts = rolling_skew(a,10)
>>> print('#original:', len(nona(a)), 'timeseries:', len(nona(ts)), 'panda:', len(nona(panda)))
>>> #original: 4534 timeseries: 4525 panda: 2 data points
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = rolling_skew(a,10)
>>> old_ts = rolling_skew_(old,10)
>>> new_ts = rolling_skew(new, 10, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(rolling_skew(dict(x = a, y = a**2),10), dict(x = rolling_skew(a,10), y = rolling_skew(a**2,10)))
>>> assert eq(rolling_skew([a,a**2],10), [rolling_skew(a,10), rolling_skew(a**2,10)])
```

**rolling\_min**

pyg.timeseries.\_min.rolling\_min(a, n, axis=0, data=None, state=None)  
equivalent to pandas a.rolling(n).min().

- works with np.arrays
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a :** array, pd.Series, pd.DataFrame or list/dict of these timeseries

**n: int**

size of rolling window

**axis :** int, optional

0/1/-1. The default is 0.

**data: None.**

unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state: dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas



```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.rolling(10).min(); ts = rolling_min(a,10)
>>> assert abs(ts-panda).min()<1e-10
```

**Example:** nan handling

Unlike pandas, timeseries does not include the nans in the rolling calculation: it skips them. Since pandas rolling engine does not skip nans, they propagate. In fact, having removed half the data points, rolling(10) will return 99% of nans

```
>>> a[a<0.1] = np.nan
>>> panda = a.rolling(10).min(); ts = rolling_min(a,10)
>>> print('#original:', len(nona(a)), 'timeseries:', len(nona(ts)), 'panda:', len(nona(panda)))
>>> #original: 4534 timeseries: 4525 panda: 6 data points
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = rolling_min(a,10)
>>> old_ts = rolling_min_(old,10)
>>> new_ts = rolling_min(new, 10, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(rolling_min(dict(x = a, y = a**2),10), dict(x = rolling_min(a,10), y = rolling_min(a**2,10)))
>>> assert eq(rolling_min([a,a**2],10), [rolling_min(a,10), rolling_min(a**2,10)])
```

**rolling\_max**

`pyg.timeseries._max.rolling_max` (*a*, *n*, *axis*=0, *data*=None, *state*=None)  
equivalent to `pandas a.rolling(n).max()`.

- works with np.arrays
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a** : *array, pd.Series, pd.DataFrame or list/dict of these timeseries*

**n**: *int*

size of rolling window

**axis** : *int, optional*

0/1/-1. The default is 0.

**data**: **None**.

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: **dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.rolling(10).max(); ts = rolling_max(a,10)
>>> assert abs(ts-panda).max()<1e-10
```

**Example:** nan handling

Unlike pandas, timeseries does not include the nans in the rolling calculation: it skips them. Since pandas rolling engine does not skip nans, they propagate. In fact, having removed half the data points, rolling(10) will return 99% of nans

```
>>> a[a<0.1] = np.nan
>>> panda = a.rolling(10).max(); ts = rolling_max(a,10)
>>> print('#original:', len(nona(a)), 'timeseries:', len(nona(ts)), 'panda:', len(nona(panda)))
>>> #original: 4534 timeseries: 4525 panda: 6 data points
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = rolling_max(a,10)
>>> old_ts = rolling_max_(old,10)
>>> new_ts = rolling_max(new, 10, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(rolling_max(dict(x = a, y = a**2),10), dict(x = rolling_max(a,10), y = rolling_max(a**2,10)))
>>> assert eq(rolling_max([a,a**2],10), [rolling_max(a,10), rolling_max(a**2,10)])
```

**rolling\_median**

pyg.timeseries.\_median.**rolling\_median** (a, n, axis=0, data=None, state=None)  
equivalent to pandas a.rolling(n).median().

- works with np.arrays
- handles nan without forward filling.
- supports state parameters

**Parameters:**

**a** : array, pd.Series, pd.DataFrame or list/dict of these timeseries

**n** : int  
size of rolling window

**axis** : int, optional  
0/1/-1. The default is 0.

**data**: None.  
unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state**: dict, optional  
state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** agreement with pandas

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> panda = a.rolling(10).median(); ts = rolling_median(a,10)
>>> assert abs(ts-panda).max()<1e-10
```

**Example:** nan handling

Unlike pandas, timeseries does not include the nans in the rolling calculation: it skips them. Since pandas rolling engine does not skip nans, they propagate. In fact, having removed half the data points, rolling(10) will return 99% of nans

```
>>> a[a<0.1] = np.nan
>>> panda = a.rolling(10).median(); ts = rolling_median(a,10)
>>> print('#original:', len(nona(a)), 'timeseries:', len(nona(ts)), 'panda:', len(nona(panda)))
#original: 4634 timeseries: 4625 panda: 4 data points
```

**Example:** state management

One can split the calculation and run old and new data separately.

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> ts = rolling_median(a,10)
>>> old_ts = rolling_median_(old,10)
>>> new_ts = rolling_median(new, 10, **old_ts)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** dict/list inputs

```
>>> assert eq(rolling_median(dict(x = a, y = a**2),10), dict(x = rolling_median(a,10), y = rolling_median(a**2,10)))
>>> assert eq(rolling_median([a,a**2],10), [rolling_median(a,10), rolling_median(a**2,10)])
```

**rolling\_quantile**

pyg.timeseries.\_stride.rolling\_quantile(a, n, quantile=0.5, axis=0, data=None, state=None)  
equivalent to a.rolling(n).quantile(q) except... - supports numpy arrays - supports multiple q values

**Example:**

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> res = rolling_quantile(a, 100, 0.3)
>>> assert sub_(res, a.rolling(100).quantile(0.3)).max() < 1e-13
```

**Example:** multiple quantiles

```
>>> res = rolling_quantile(a, 100, [0.3, 0.5, 0.75])
>>> assert abs(res[0.3] - a.rolling(100).quantile(0.3)).max() < 1e-13
```

**Example:** state management

```
>>> res = rolling_quantile(a, 100, 0.3)
>>> old = rolling_quantile_(a.iloc[:2000], 100, 0.3)
>>> new = rolling_quantile(a.iloc[2000:], 100, 0.3, **old)
>>> both = pd.concat([old.data, new])
>>> assert eq(both, res)
```

**Parameters:**

**a** : array/timeseries  
**n** : integer

window size.

**q** : float or list of floats in [0,1]  
quantile(s).

**data:** None.

unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state:** dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Returns:**

timeseries/array of quantile(s)

## rolling\_rank

`pyg.timeseries._rank.rolling_rank(a, n, axis=0, data=None, state=None)`

returns a rank of the current value within a given window, scaled to be -1 if it is the smallest and +1 if it is the largest - works on mumpy arrays too - skips nan, no ffill

### Parameters:

**a** : array, *pd.Series*, *pd.DataFrame* or list/dict of these timeseries

**n**: int  
window size

**axis** : int, optional  
0/1/-1. The default is 0.

**data**: None.  
unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: dict, optional  
state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

### Example:

```
>>> from pyg import *; import pandas as pd; import numpy as np
>>> a = pd.Series([1.,2., np.nan, 0., 4., 2., 3., 1., 2.], drange(-8))
>>> rank = rolling_rank(a, 3)
>>> assert eq(rank.values, np.array([np.nan, np.nan, np.nan, -1, 1, 0, 0, -1, 0]))
>>> # 0 is smallest in [1,2,0] so goes to -1
>>> # 4 is largest in [2,0,4] so goes to +1
>>> # 2 is middle of [0,4,2] so goes to 0
```

### Example: numpy equivalent

```
>>> assert eq(rolling_rank(a.values, 10), rolling_rank(a, 10).values)
```

### Example: state management

```
>>> a = np.random.normal(0,1,10000)
>>> old = rolling_rank(a[:5000], 10) # grab both data and state
>>> new = rolling_rank(a[5000:], 10, **old)
>>> assert eq(np.concatenate([old.data,new]), rolling_rank(a, 10))
```

## exponentially weighted moving functions

### ewma

`pyg.timeseries._ewm.ewma(a, n, time=None, axis=0, data=None, state=None)`

ewma is equivalent to `a.ewm(n).mean()` but with... - supports np.ndarrays as well as timeseries - handles nan by skipping them - allows state-management - ability to supply a 'clock' to the calculation

### Parameters:

**a** : array/timeseries **n** : int/fraction

The number or days (or a ratio) to scale the history

**time** : Calendar, 'b/d/y/m' or a timeseries of time (use `clock(a)` to see output)

If time parameter is provided, we allow multiple observations per unit of time. i.e., converging to the last observation in time unit.

- if we have intraday data, and set time = 'd', then
- the ewm calculation on last observations per day is what is retained.

- the ewm calculation on each intraday observation is same as an ewm(past EOD + current intraday observation)

**data:** None.

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state:** dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** matching pandas

```
>>> import pandas as pd; import numpy as np; from pyg import *
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> ts = ewma(a,10); df = a.ewm(10).mean()
>>> assert abs(ts-df).max()<1e-10
```

**Example:** numpy arrays support

```
>>> assert eq(ewma(a.values, 10), ewma(a,10).values)
```

**Example:** nan handling

```
>>> a[a.values<0.1] = np.nan
>>> ts = ewma(a,10, time = 'i'); df = a.ewm(10).mean() # note: pandas assumes, 'time' pass
>>> assert abs(ts-df).max()<1e-10
```

```
>>> pd.concat([ts,df], axis=1)
>>>
>>> 1993-09-24  0.263875  0.263875
>>> 1993-09-25      NaN  0.263875
>>> 1993-09-26      NaN  0.263875
>>> 1993-09-27      NaN  0.263875
>>> 1993-09-28      NaN  0.263875
>>>
>>> 2021-02-04      NaN  0.786506
>>> 2021-02-05  0.928817  0.928817
>>> 2021-02-06      NaN  0.928817
>>> 2021-02-07  0.839168  0.839168
>>> 2021-02-08  0.831109  0.831109
```

**Example:** state management

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> old_ts = ewma_(old, 10)
>>> new_ts = ewma(new, 10, **old_ts) # instantiation with previous ewma
>>> ts = ewma(a,10)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** Support for time & clock

```
>>> daily = a
>>> monthly = daily.resample('M').last()
>>> m_ts = ewma(monthly, 3) ## 3-month ewma run on monthly data
>>> d_ts = ewma(daily, 3, 'm') ## 3-month ewma run on daily data
>>> daily_resampled_to_month = d_ts.resample('M').last()
>>> assert abs(daily_resampled_to_month - m_ts).max() < 1e-10
```

So you can run a 3-monthly ewma on daily, where within month, most recent value is used with the EOM history.

**Example:** Support for dict/list of arrays

```
>>> x = pd.Series(np.random.normal(0,1,1000), drange(-999)); y = pd.Series(np.random.normal(0,1,1000), drange(-999))
>>> a = dict(x = x, y = y)
>>> assert eq(ewma(dict(x=x, y=y),10), dict(x=ewma(x,10), y=ewma(y,10)))
>>> assert eq(ewma([x,y],10), [ewma(x,10), ewma(y,10)])
```

**Returns:**

an array/timeseries of ewma

**ewmrms**

pyg.timeseries.\_ewm.ewmrms(a, n, time=None, axis=0, data=None, state=None)

ewmrms is equivalent to  $(a**2).ewm(n).mean()**0.5$  but with... - supports np.ndarrays as well as timeseries - handles nan by skipping them - allows state-management - ability to supply a 'clock' to the calculation

**Parameters:**

a : array/timeseries n : int/fraction

The number or days (or a ratio) to scale the history

**time** : Calendar, 'b/d/y/m' or a timeseries of time (use clock(a) to see output)

**If time parameter is provided, we allow multiple observations per unit of time. i.e., converging to the last observation in time unit.**

- if we have intraday data, and set time = 'd', then
- the ewm calculation on last observations per day is what is retained.
- the ewm calculation on each intraday observation is same as an ewm(past EOD + current intraday observation)

**data: None.**

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state: dict, optional**

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** matching pandas

```
>>> import pandas as pd; import numpy as np; from pyg import *
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> ts = ewmrms(a,10); df = (a**2).ewm(10).mean()**0.5
>>> assert abs(ts-df).max()<1e-10
```

**Example:** numpy arrays support

```
>>> assert eq(ewmrms(a.values, 10), ewmrms(a,10).values)
```

**Example:** nan handling

```
>>> a[a.values<0.1] = np.nan
>>> ts = ewmrms(a,10, time = 'i'); df = (a**2).ewm(10).mean()**0.5 # note: pandas assumes,
>>> assert abs(ts-df).max()<1e-10
```

```
>>> pd.concat([ts,df], axis=1)
>>>
>>> 1993-09-24  0.263875  0.263875
>>> 1993-09-25      NaN  0.263875
>>> 1993-09-26      NaN  0.263875
>>> 1993-09-27      NaN  0.263875
>>> 1993-09-28      NaN  0.263875
>>>
>>> 2021-02-04      NaN  0.786506
>>> 2021-02-05  0.928817  0.928817
```

```
>>> 2021-02-06      NaN    0.928817
>>> 2021-02-07    0.839168    0.839168
>>> 2021-02-08    0.831109    0.831109
```

**Example:** state management

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> old_ts = ewmrms_(old, 10)
>>> new_ts = ewmrms(new, 10, **old_ts) # instantiation with previous ewma
>>> ts = ewmrms(a,10)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** Support for time & clock

```
>>> daily = a
>>> monthly = daily.resample('M').last()
>>> m_ts = ewmrms(monthly, 3) ## 3-month ewma run on monthly data
>>> d_ts = ewmrms(daily, 3, 'm') ## 3-month ewma run on daily data
>>> daily_resampled_to_month = d_ts.resample('M').last()
>>> assert abs(daily_resampled_to_month - m_ts).max() < 1e-10
```

So you can run a 3-monthly ewma on daily, where within month, most recent value is used with the EOM history.

**Example:** Support for dict/list of arrays

```
>>> x = pd.Series(np.random.normal(0,1,1000), drange(-999)); y = pd.Series(np.random.normal(0,1,1000), drange(-999))
>>> a = dict(x=x, y=y)
>>> assert eq(ewmrms(dict(x=x, y=y),10), dict(x=ewmrms(x,10), y=ewmrms(y,10)))
>>> assert eq(ewmrms([x,y],10), [ewmrms(x,10), ewmrms(y,10)])
```

**Returns:**

an array/timeseries of ewma

**ewmstd**

pyg.timeseries.\_ewm.**ewmstd**(a, n, time=None, min\_sample=0.25, bias=False, axis=0, data=None, state=None)

ewmstd is equivalent to a.ewm(n).std() but with... - supports np.ndarrays as well as timeseries - handles nan by skipping them - allows state-management - ability to supply a 'clock' to the calculation

**Parameters:**

a : array/timeseries n : int/fraction

The number or days (or a ratio) to scale the history

**time** : Calendar, 'b/d/y/m' or a timeseries of time (use clock(a) to see output)

**If time parameter is provided, we allow multiple observations per unit of time. i.e., converging to the last observation in time unit.**

- if we have intraday data, and set time = 'd', then
- the ewm calculation on last observations per day is what is retained.
- the ewm calculation on each intraday observation is same as an ewm(past EOD + current intraday observation)

**data:** None.

unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

**state:** dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** matching pandas

```
>>> import pandas as pd; import numpy as np; from pyg import *
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> ts = ewmstd(a,10); df = a.ewm(10).std()
>>> assert abs(ts-df).max()<1e-10
>>> ts = ewmstd(a,10, bias = True); df = a.ewm(10).std(bias = True)
>>> assert abs(ts-df).max()<1e-10
```

**Example:** numpy arrays support

```
>>> assert eq(ewmstd(a.values, 10), ewmstd(a,10).values)
```

**Example:** nan handling

```
>>> a[a.values<-0.1] = np.nan
>>> ts = ewmstd(a,10, time = 'i'); df = a.ewm(10).std() # note: pandas assumes, 'time' pas
>>> assert abs(ts-df).max()<1e-10
>>> ts = ewmstd(a,10, time = 'i', bias = True); df = a.ewm(10).std(bias = True) # note: pa
>>> assert abs(ts-df).max()<1e-10
```

**Example:** state management

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> old_ts = ewmstd(old, 10)
>>> new_ts = ewmstd(new, 10, **old_ts) # instantiation with previous ewma
>>> ts = ewmstd(a,10)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** Support for time & clock

```
>>> daily = a
>>> monthly = daily.resample('M').last()
>>> m_ts = ewmstd(monthly, 3) ## 3-month ewma run on monthly data
>>> d_ts = ewmstd(daily, 3, 'm') ## 3-month ewma run on daily data
>>> daily_resampled_to_month = d_ts.resample('M').last()
>>> assert abs(daily_resampled_to_month - m_ts).max() < 1e-10
```

So you can run a 3-monthly ewma on daily, where within month, most recent value is used with the EOM history.

**Example:** Support for dict/list of arrays

```
>>> x = pd.Series(np.random.normal(0,1,1000), drange(-999)); y = pd.Series(np.random.normal(0,1,1000), drange(-999))
>>> a = dict(x = x, y = y)
>>> assert eq(ewmstd(dict(x=x, y=y),10), dict(x=ewmstd(x,10), y=ewmstd(y,10)))
>>> assert eq(ewmstd([x,y],10), [ewmstd(x,10), ewmstd(y,10)])
```

### Returns:

an array/timeseries of ewma

## ewmskew

pyg.timeseries.\_ewm.**ewmskew** (a, n, time=None, bias=False, min\_sample=0.25, axis=0, data=None, state=None)

Equivalent to a.ewm(n).skew() but with... - supports np.ndarrays as well as timeseries - handles nan by skipping them - allows state-management - ability to supply a 'clock' to the calculation

### Parameters:

a : array/timeseries n : int/fraction

The number or days (or a ratio) to scale the history



**time** : Calendar, 'b/d/y/m' or a timeseries of time (use `clock(a)` to see output)

If time parameter is provided, we allow multiple observations per unit of time. i.e., converging to the last observation in time unit.

- if we have intraday data, and set time = 'd', then
- the ewm calculation on last observations per day is what is retained.
- the ewm calculation on each intraday observation is same as an ewm(past EOD + current intraday observation)

**data**: None.

unused at the moment. Allow code such as `func(live, **func_(history))` to work

**state**: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** matching pandas

```
>>> import pandas as pd; import numpy as np; from pyg import *
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> x = a.ewm(10).skew()
```

```
>>> old = a.iloc[:10]
>>> new = a.iloc[10:]
>>> for f in [ewma_, ewmstd_, ewmrms_, ewmskew_, ]:
>>>     both = f(a, 3)
>>>     o = f(old, 3)
>>>     n = f(new, 3, **o)
>>>     assert eq(o.data, both.data.iloc[:10])
>>>     assert eq(n.data, both.data.iloc[10:])
>>>     assert both - 'data' == n - 'data'
```

```
>>> assert abs(a.ewm(10).mean() - ewma(a,10)).max() < 1e-14
>>> assert abs(a.ewm(10).std() - ewmstd(a,10)).max() < 1e-14
```

**Example:** numpy arrays support

```
>>> assert eq(ewma(a.values, 10), ewma(a,10).values)
```

**Example:** nan handling

while pandas ffill values, timeseries skips nans:

```
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> a[a.values>0.1] = np.nan
>>> ts = ewma(a,10)
>>> assert eq(ts[~np.isnan(ts)], ewma(a[~np.isnan(a)], 10))
```

**Example:** initiating the ewma with past state

```
>>> old = np.random.normal(0,1,100)
>>> new = np.random.normal(0,1,100)
>>> old_ = ewma_(old, 10)
>>> new_ = ewma(new, 10, t0 = old_.t0, t1 = old_.t1) # instantiation with previous
>>> new_2 = ewma(np.concatenate([old,new]), 10)[-100:]
>>> assert eq(new_ewma, new_ewma2)
```

**Example:** Support for time & clock

```
>>> daily = pd.Series(np.random.normal(0,1,10000), drange(-9999)).cumsum()
>>> monthly = daily.resample('M').last()
>>> m = ewma(monthly, 3) ## 3-month ewma run on monthly data
>>> d = ewma(daily, 3, 'm') ## 3-month ewma run on daily data
```

```
>>> daily_resampled_to_month = d.resample('M').last()
>>> assert abs(daily_resampled_to_month - m).max() < 1e-10
```

So you can run a 3-monthly ewma on daily, where within month, most recent value is used with the EOM history.

#### Returns:

an array/timeseries of ewma

### ewmvar

pyg.timeseries.\_ewm.ewmvar(a, n, time=None, min\_sample=0.25, bias=False, axis=0, data=None, state=None)

ewmstd is equivalent to a.ewm(n).var() but with... - supports np.ndarrays as well as timeseries - handles nan by skipping them - allows state-management - ability to supply a 'clock' to the calculation

#### Parameters:

a : array/timeseries n : int/fraction

The number or days (or a ratio) to scale the history

time : Calendar, 'b/d/y/m' or a timeseries of time (use clock(a) to see output)

If time parameter is provided, we allow multiple observations per unit of time. i.e., converging to the last observation in time unit.

- if we have intraday data, and set time = 'd', then
- the ewm calculation on last observations per day is what is retained.
- the ewm calculation on each intraday observation is same as an ewm(past EOD + current intraday observation)

data: None.

unused at the moment. Allow code such as func(live, \*\*func\_(history)) to work

state: dict, optional

state parameters used to instantiate the internal calculations, based on history prior to 'a' provided.

**Example:** matching pandas

```
>>> import pandas as pd; import numpy as np; from pyg import *
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> ts = ewmvar(a,10); df = a.ewm(10).var()
>>> assert abs(ts-df).max()<1e-10
>>> ts = ewmvar(a,10, bias = True); df = a.ewm(10).var(bias = True)
>>> assert abs(ts-df).max()<1e-10
```

**Example:** numpy arrays support

```
>>> assert eq(ewmvar(a.values, 10), ewmvar(a,10).values)
```

**Example:** nan handling

```
>>> a[a.values<-0.1] = np.nan
>>> ts = ewmvar(a,10, time = 'i'); df = a.ewm(10).var() # note: pandas assumes, 'time' pas
>>> assert abs(ts-df).max()<1e-10
>>> ts = ewmvar(a,10, time = 'i', bias = True); df = a.ewm(10).var(bias = True) # note: pa
>>> assert abs(ts-df).max()<1e-10
```

**Example:** state management

```
>>> old = a.iloc[:5000]
>>> new = a.iloc[5000:]
>>> old_ts = ewmvar_(old, 10)
>>> new_ts = ewmvar(new, 10, **old_ts) # instantiation with previous ewma
```

```
>>> ts = ewmvar(a,10)
>>> assert eq(new_ts, ts.iloc[5000:])
```

**Example:** Support for time & clock

```
>>> daily = a
>>> monthly = daily.resample('M').last()
>>> m_ts = ewmvar(monthly, 3) ## 3-month ewma run on monthly data
>>> d_ts = ewmvar(daily, 3, 'm') ## 3-month ewma run on daily data
>>> daily_resampled_to_month = d_ts.resample('M').last()
>>> assert abs(daily_resampled_to_month - m_ts).max() < 1e-10
```

So you can run a 3-monthly ewma on daily, where within month, most recent value is used with the EOM history.

**Example:** Support for dict/list of arrays

```
>>> x = pd.Series(np.random.normal(0,1,1000), drange(-999)); y = pd.Series(np.random.normal(0,1,1000), drange(-999))
>>> a = dict(x=x, y=y)
>>> assert eq(ewmvar(dict(x=x, y=y),10), dict(x=ewmvar(x,10), y=ewmvar(y,10)))
>>> assert eq(ewmvar([x,y],10), [ewmvar(x,10), ewmvar(y,10)])
```

### Returns:

an array/timeseries of ewma

## ewmcor

pyg.timeseries.\_ewm.ewmcor(a, b, n, time=None, min\_sample=0.25, bias=True, axis=0, data=None, state=None)

calculates pair-wise correlation between a and b.

### Parameters:

a : array/timeseries b : array/timeseries n : int/fraction

The number or days (or a ratio) to scale the history

**time** : Calendar, 'b/d/y/m' or a timeseries of time (use clock(a) to see output)

**If time parameter is provided, we allow multiple observations per unit of time. i.e., converging to the last observation in time unit.**

- if we have intraday data, and set time = 'd', then
- the ewm calculation on last observations per day is what is retained.
- the ewm calculation on each intraday observation is same as an ewm(past EOD + current intraday observation)

**min\_sample** : float, optional

minimum weight of observations before we return a reading. The default is 0.25. This ensures that we don't get silly numbers due to small population.

**bias** : bool, optional

vol estimation for a and b should really be unbiased. Nevertheless, we track pandas and set bias = True as a default.

**axis** : int, optional

axis of calculation. The default is 0.

**data** : place holder, ignore, optional

ignore. The default is None.

**state** : dict, optional

Output from a previous run of ewmcor. The default is None.

**Example:** matching pandas

```
>>> import pandas as pd; import numpy as np; from pyg import *
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> b = pd.Series(np.random.normal(0,1,9000), drange(-8999))
>>> ts = ewmcor(a, b, n = 10); df = a.ewm(10).corr(b)
>>> assert abs(ts-df).max()<1e-10
```

**Example:** numpy arrays support

```
>>> assert eq(ewmcor(a.values, b.values, 10), ewmcor(a, b, 10).values)
```

**Example:** nan handling

```
>>> a[a.values<-0.1] = np.nan
>>> ts = ewmcor(a, b, 10, time = 'i'); df = a.ewm(10).corr(b) # note: pandas assumes, 'time'
>>> assert abs(ts-df).max()<1e-10
```

**Example:** state management

```
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> b = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> old_a = a.iloc[:5000]; old_b = b.iloc[:5000]
>>> new_a = a.iloc[5000:]; new_b = b.iloc[5000:]
>>> old_ts = ewmcor_(old_a, old_b, 10)
>>> new_ts = ewmcor(new_a, new_b, 10, **old_ts) # instantiation with previous ewma
>>> ts = ewmcor(a,b,10)
>>> assert eq(new_ts, ts.iloc[5000:])
```

## ewmcorr

pyg.timeseries.\_ewm.ewmcorr(a, n, min\_sample=0.25, bias=False, instate=None)

This calculates a full correlation matrix as a timeseries. The calculation is returned as an xarray.Dataset object, which is a multidimensional data structure.

### Parameters:

**a** : *np.array* or *a pd.DataFrame*

timeseries to calculate correlation for

**n** : *int*

days for which rolling correlation is calculated.

**min\_sample** : *float, optional*

Minimum observations needed before we calculate correlation. The default is 0.25.

**bias** : *bool, optional*

input to stdev calculations, the default is False.

**instate** : *dict, optional*

historical calculations so far.

**correlation dataset**

an xarray.Dataset unless there are only two timeserieses

**Example:** a pair of ts

```
>>> rtn = np.random.normal(0,1,10000)
>>> x0 = ewmxo(rtn, 10, 20, 30)[50:]
>>> x1 = ewmxo(rtn, 20, 40, 30)[50:]
>>> a = pd.DataFrame(np.array([x0,x1]).T, drange(-9949))
>>> res = ewmcorr(a, n)
>>> res
```

```
>>> Out[130]:
>>> 1994-06-07      NaN
>>> 1994-06-08      NaN
>>> 1994-06-09      NaN
>>> 1994-06-10      NaN
>>> 1994-06-11      NaN
```

```
>>> 2021-08-29      0.890766
>>> 2021-08-30      0.886926
>>> 2021-08-31      0.883054
>>> 2021-09-01      0.879577
>>> 2021-09-02      0.875766
>>> Length: 9950, dtype: float64
```

**Example:** a pair of ts

```
>>> rtn = np.random.normal(0,1,10000)
>>> x0 = ewmxo(rtn, 10, 20, 30)[50:]
>>> x1 = ewmxo(rtn, 20, 40, 30)[50:]
>>> x2 = ewmxo(rtn, 40, 80, 30)[50:]
>>> a = pd.DataFrame(np.array([x0,x1,x2]).T, drange(-9949), ['a','b','c'])
>>> ds = ewmcorr(a, n)
>>> ds
```

```
>>> <xarray.Dataset>
>>> Dimensions:  (index: 9950, x: 3, y: 3)
>>> Coordinates:
>>>   * index      (index) datetime64[ns] 1994-06-07 1994-06-08 ... 2021-09-02
>>>   * x          (x) <U1 'a' 'b' 'c'
>>>   * y          (y) <U1 'a' 'b' 'c'
>>> Data variables:
>>>   cor          (index, x, y) float64 1.0 nan nan nan ... 0.9402 0.7254 0.9402 1.0
```

To access individual correlations:

```
>>> a_vs_b = pd.Series(ds.loc[dict(x = 'a', y = 'b')].cor.values, ds.index.values)
```

To access all correlations to a:

```
>>> subset = ds.loc[dict(x = 'a')]
>>> a_vs_all = pd.DataFrame(subset.cor.values, ds.index.values, subset.y.values)
>>> a_vs_all
```

```
>>>
>>>      a      b      c
>>> 1994-06-07  1.0   NaN   NaN
>>> 1994-06-08  1.0   NaN   NaN
>>> 1994-06-09  1.0   NaN   NaN
>>> 1994-06-10  1.0   NaN   NaN
>>> 1994-06-11  1.0   NaN   NaN
>>> ...
>>> 2021-08-29  1.0  0.890766  0.753859
>>> 2021-08-30  1.0  0.886926  0.746812
>>> 2021-08-31  1.0  0.883054  0.739746
>>> 2021-09-01  1.0  0.879577  0.733322
>>> 2021-09-02  1.0  0.875766  0.725399
```

## ewmLR

pyg.timeseries.\_ewm.ewmLR (a, b, n, time=None, min\_sample=0.25, bias=True, axis=0, c=None, m=None, state=None)

calculates pair-wise linear regression between a and b. We have a and b for which we want to fit:

```
>>> b_i = c + m a_i
>>> LSE(c,m) = \sum w_i (c + m a_i - b_i)^2
>>> dLSE/dc = 0 <==> \sum w_i (c + m a_i - b_i) = 0 [1]
>>> dLSE/dm = 0 <==> \sum w_i a_i (c + m a_i - b_i) = 0 [2]
```

```
>>> c + mE(a) = E(b) [1]
>>> cE(a) + mE(a^2) = E(ab) [2]
```

```
>>> cE(a) + mE(a)^2 = E(a)E(n) [1] * E(a)
>>> m(E(a^2) - E(a)^2) = E(ab) - E(a)E(b)
>>> m = covar(a,b)/var(a)
>>> c = E(b) - mE(a)
```

a : array/timeseries b : array/timeseries n : int/fraction

The number or days (or a ratio) to scale the history

**time** : *Calendar, 'b/d/y/m' or a timeseries of time (use clock(a) to see output)*

**If time parameter is provided, we allow multiple observations per unit of time. i.e., converging to the last observation in time unit.**

- if we have intraday data, and set time = 'd', then
- the ewm calculation on last observations per day is what is retained.
- the ewm calculation on each intraday observation is same as an ewm(past EOD + current intraday observation)

**min\_sample** : *float, optional*

minimum weight of observations before we return a reading. The default is 0.25. This ensures that we don't get silly numbers due to small population.

**bias** : *bool, optional*

vol estimation for a and b should really be unbiased. Nevertheless, we track pandas and set bias = True as a default.

**axis** : *int, optional*

axis of calculation. The default is 0.

**c,m** : *place holder, ignore, optional*

ignore. The default is None.

**state** : *dict, optional*

Output from a previous run of [ewmcor](#). The default is None.

**Example:** numpy arrays support

```
>>> assert eq(ewmLR(a.values, b.values, 10), ewmLR(a, b, 10).values)
```

**Example:** nan handling

```
>>> a[a.values<-0.1] = np.nan
>>> ts = ewmcor(a, b, 10, time = 'i'); df = a.ewm(10).corr(b) # note: pandas assumes, 'tim
>>> assert abs(ts-df).max()<1e-10
```

**Example:** state management

```
>>> from pyg import *
>>> a = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> b = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> old_a = a.iloc[:5000]; old_b = b.iloc[:5000]
>>> new_a = a.iloc[5000:]; new_b = b.iloc[5000:]
>>> old_ts = ewmLR(old_a, old_b, 10)
>>> new_ts = ewmLR(new_a, new_b, 10, **old_ts) # instantiation with previous ewma
```

```
>>> ts = ewmLR(a,b,10)
>>> assert eq(new_ts.c, ts.c.iloc[5000:])
>>> assert eq(new_ts.m, ts.m.iloc[5000:])
```

**Example:**

```
>>> from pyg import *
>>> a0 = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> a1 = pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> b = (a0 - a1) + pd.Series(np.random.normal(0,1,10000), drange(-9999))
>>> a = pd.concat([a0,a1], axis=1)
>>> LR = ewmLR(a,b,50)
>>> assert abs(LR.m.mean()[0]-1)<0.5
>>> assert abs(LR.m.mean()[1]+1)<0.5
```

**ewmGLM**

pyg.timeseries.\_ewm.ewmGLM(a, b, n, time=None, min\_sample=0.25, bias=True, data=None, state=None)  
Calculates a General Linear Model fitting b to a.

**Parameters:**

a : a 2-d array/pd.DataFrame of values fitting b b : a 1-d array/pd.Series n : int/fraction

The number or days (or a ratio) to scale the history

**time** : Calendar, 'b/d/y/m' or a timeseries of time (use clock(a) to see output)

**If time parameter is provided, we allow multiple observations per unit of time. i.e., converging to the last observation in time unit.**

- if we have intraday data, and set time = 'd', then
- the ewm calculation on last observations per day is what is retained.
- the ewm calculation on each intraday observation is same as an ewm(past EOD + current intraday observation)

**min\_sample** : float, optional

minimum weight of observations before we return the fitting. The default is 0.25. This ensures that we don't get silly numbers due to small population.

**data** : place holder, ignore, optional

ignore. The default is None.

**state** : dict, optional

Output from a previous run of ewmGLM. The default is None.

**Theory:**

See [https://en.wikipedia.org/wiki/Generalized\\_linear\\_model](https://en.wikipedia.org/wiki/Generalized_linear_model) for full details. Briefly, we assume b is single column while a is multicolumn. We minimize least square error (LSE) fitting:

```
>>> b[i] = \sum_j m_j a_j[i]
>>> LSE(m) = \sum_i w_i (b[i] - \sum_j m_j * a_j[i])^2
```

```
>>> dLSE/dm_k = 0
>>> <==> \sum_i w_i (b[i] - \sum_j m_j * a_j[i]) a_k[i] = 0
>>> <==> E(b*a_k) = m_k E(a_k^2) + sum_{j<>k} m_k E(a_j a_k)
```

E is expectation under weights w. And we can rewrite it as:

```
>>> a2 x m = ab ## matrix multiplication
>>> a2[i,j] = E(a_i * a_j)
>>> ab[j] = E(a_j * b)
>>> m = a2.inverse x ab ## matrix multiplication
```

**Example:** simple fit

```
>>> from pyg import *
>>> a = pd.DataFrame(np.random.normal(0,1,(10000,10)), drange(-9999))
>>> true_m = np.random.normal(1,1,10)
>>> noise = np.random.normal(0,1,10000)
>>> b = (a * true_m).sum(axis = 1) + noise
```

```
>>> fitted_m = ewmGLM(a, b, 50)
```

**ewmxo**

pyg.timeseries.\_ewmxo.**ewmxo** (rtn, fast, slow, vol=None, instate=None)

This is the normalized crossover function

```
>>> res = (ewma(rtn, fast) - ewma(rtn, slow)) / (ewmstd(rtn, vol) * ou_factor(fast, slow))
```

The OU factor normalizes the result so that rms(res) is approximately 1

**rtn: timeseries**

The returns of a financial process

**fast : int/frc**

number of days. can also be 1/(1+days) if presented as a fraction

**slow : int/frc**

number of days. can also be 1/(1+days) if presented as a fraction

**vol: int/frc**

number of days. used for calculating the volatility horizon

**Example:**

```
>>> import numpy as np; import pandas as pd; from pyg import *
>>> rtn = pd.Series(np.random.normal(0,1,10000),drange(-9999,0))
>>> fast = 64; slow = 192; vol = 32; instate = None
```

**functions exposing their state****simple functions**

pyg.timeseries.\_rolling.**diff\_** (a, n=1, axis=0, data=None, instate=None)

returns a forward filled array, up to n values forward. Equivalent to diff(a,n) but returns the full state. See diff for full details

pyg.timeseries.\_rolling.**shift\_** (a, n=1, axis=0, instate=None)

Equivalent to shift(a,n) but returns the full state. See shift for full details

pyg.timeseries.\_rolling.**ratio\_** (a, n=1, data=None, instate=None)

pyg.timeseries.\_ts.**ts\_count\_** (a, axis=0, data=None, instate=None)

ts\_count\_(a) is equivalent to ts\_count(a) except vec is also returned. See ts\_count for full documentation

pyg.timeseries.\_ts.**ts\_sum\_** (a, axis=0, data=None, instate=None)

ts\_sum\_(a) is equivalent to ts\_sum(a) except vec is also returned. See ts\_sum for full documentation

pyg.timeseries.\_ts.**ts\_mean\_** (a, axis=0, data=None, instate=None)

ts\_mean\_(a) is equivalent to ts\_mean(a) except vec is also returned. See ts\_mean for full documentation

pyg.timeseries.\_ts.**ts\_rms\_** (a, axis=0, data=None, instate=None)

ts\_rms\_(a) is equivalent to ts\_rms(a) except it also returns vec see ts\_rms for full documentation

pyg.timeseries.\_ts.**ts\_std\_** (a, axis=0, data=None, instate=None)

ts\_std\_(a) is equivalent to ts\_std(a) except vec is also returned. See ts\_std for full documentation

pyg.timeseries.\_ts.**ts\_skew\_** (a, bias=False, min\_sample=0.25, axis=0, data=None, instate=None)



`ts_skew_(a)` is equivalent to `ts_skew` except `vec` is also returned. See `ts_skew` for full details

`pyg.timeseries._ts.ts_max_` (a, axis=0, data=None, instate=None)

`ts_max(a)` is equivalent to `pandas a.min()`

`pyg.timeseries._ts.ts_max_` (a, axis=0, data=None, instate=None)

`ts_max(a)` is equivalent to `pandas a.min()`

`pyg.timeseries._rolling.ffill_` (a, n=0, axis=0, instate=None)

returns a forward filled array, up to n values forward. supports state management

## ***expanding window functions***

`pyg.timeseries._expanding.expanding_mean_` (a, axis=0, data=None, instate=None)

Equivalent to `expanding_mean(a)` but returns also the state variables. For full documentation, look at `expanding_mean.__doc__`

`pyg.timeseries._expanding.expanding_rms_` (a, axis=0, data=None, instate=None)

Equivalent to `expanding_rms(a)` but returns also the state variables. For full documentation, look at `expanding_rms.__doc__`

`pyg.timeseries._expanding.expanding_std_` (a, axis=0, data=None, instate=None)

Equivalent to `expanding_mean(a)` but returns also the state variables. For full documentation, look at `expanding_std.__doc__`

`pyg.timeseries._expanding.expanding_sum_` (a, axis=0, data=None, instate=None)

Equivalent to `expanding_sum(a)` but returns also the state variables. For full documentation, look at `expanding_sum.__doc__`

`pyg.timeseries._expanding.expanding_skew_` (a, bias=False, axis=0, data=None, instate=None)

Equivalent to `expanding_mean(a)` but returns also the state variables. For full documentation, look at `expanding_skew.__doc__`

`pyg.timeseries._min.expanding_min_` (a, axis=0, data=None, instate=None)

Equivalent to `a.expanding().min()` but returns the full state: i.e. both data: the `expanding().min()` m: the current minimum

`pyg.timeseries._max.expanding_max_` (a, axis=0, data=None, instate=None)

Equivalent to `a.expanding().max()` but returns the full state: i.e. both data: the `expanding().max()` m: the current maximum

`pyg.timeseries._expanding.cumsum_` (a, axis=0, data=None, instate=None)

Equivalent to `expanding_sum(a)` but returns also the state variables. For full documentation, look at `expanding_sum.__doc__`

`pyg.timeseries._expanding.cumprod_` (a, axis=0, data=None, instate=None)

Equivalent to `cumprod(a)` but returns also the state variable. For full documentation, look at `cumprod.__doc__`

## ***rolling window functions***

`pyg.timeseries._rolling.rolling_mean_` (a, n, time=None, axis=0, data=None, instate=None)

Equivalent to `rolling_mean(a)` but returns also the state variables `t0,t1` etc. For full documentation, look at `rolling_mean.__doc__`

`pyg.timeseries._rolling.rolling_rms_` (a, n, time=None, axis=0, data=None, instate=None)

Equivalent to `rolling_rms(a)` but returns also the state variables `t0,t1` etc. For full documentation, look at `rolling_rms.__doc__`

`pyg.timeseries._rolling.rolling_std_` (a, n, time=None, axis=0, data=None, instate=None)

Equivalent to `rolling_std(a)` but returns also the state variables `t0,t1` etc. For full documentation, look at `rolling_std.__doc__`

`pyg.timeseries._rolling.rolling_sum_` (a, n, time=None, axis=0, data=None, instate=None)

Equivalent to `rolling_sum(a)` but returns also the state variables `t0,t1` etc. For full documentation, look at `rolling_sum.__doc__`

`pyg.timeseries._rolling.rolling_skew_` (a, n, time=None, bias=False, axis=0, data=None, instate=None)

Equivalent to `rolling_skew(a)` but returns also the state variables `t0,t1` etc. For full documentation, look at `rolling_skew.__doc__`

`pyg.timeseries._min.rolling_min_` (a, n, vec=None, axis=0, data=None, instate=None)

Equivalent to `rolling_min(a)` but returns also the state. For full documentation, look at `rolling_min.__doc__`

`pyg.timeseries._max.rolling_max_` (a, n, axis=0, data=None, instate=None)

Equivalent to `rolling_max(a)` but returns also the state. For full documentation, look at `rolling_max.__doc__`

`pyg.timeseries._median.rolling_median_` (a, n, axis=0, data=None, instate=None)

Equivalent to `rolling_median(a)` but returns also the state. For full documentation, look at `rolling_median.__doc__`

`pyg.timeseries._rank.rolling_rank_` (a, n, axis=0, data=None, instate=None)

Equivalent to `rolling_rank(a)` but returns also the state variables. For full documentation, look at `rolling_rank.__doc__`

`pyg.timeseries._stride.rolling_quantile_` (a, n, quantile=0.5, axis=0, data=None, instate=None)

Equivalent to `rolling_quantile(a)` but returns also the state. For full documentation, look at `rolling_quantile.__doc__`

## ***exponentially weighted moving functions***

`pyg.timeseries._ewm.ewma_` (a, n, time=None, data=None, instate=None)

Equivalent to `ewma` but returns a state parameter for instantiation of later calculations. See `ewma` documentation for more details

`pyg.timeseries._ewm.ewmrms_` (a, n, time=None, axis=0, data=None, instate=None)

Equivalent to `ewmrms` but returns a state parameter for instantiation of later calculations. See `ewmrms` documentation for more details

`pyg.timeseries._ewm.ewmstd_` (a, n, time=None, min\_sample=0.25, bias=False, axis=0, data=None, instate=None)

Equivalent to `ewmstd` but returns a state parameter for instantiation of later calculations. See `ewmstd` documentation for more details

`pyg.timeseries._ewm.ewmvar_` (a, n, time=None, min\_sample=0.25, bias=False, axis=0, data=None, instate=None)

Equivalent to `ewmvar` but returns a state parameter for instantiation of later calculations. See `ewmvar` documentation for more details

`pyg.timeseries._ewm.ewmcor_` (a, b, n, time=None, min\_sample=0.25, bias=True, axis=0, data=None, instate=None)

Equivalent to `ewmcor` but returns a state parameter for instantiation of later calculations. See `ewmcor` documentation for more details

`pyg.timeseries._ewm.ewmcorr_` (a, n, min\_sample=0.25, bias=False, instate=None)

This calculates a full correlation matrix as a timeseries. Also returns the recent state of the calculations. See `ewmcorr` for full details.

`pyg.timeseries._ewm.ewmLR_` (a, b, n, time=None, min\_sample=0.25, bias=True, axis=0, c=None, m=None, instate=None)

Equivalent to `ewmcor` but returns a state parameter for instantiation of later calculations. See `ewmcor` documentation for more details

`pyg.timeseries._ewm.ewmGLM_` (a, b, n, time=None, min\_sample=0.25, bias=True, data=None, instate=None)

Equivalent to `ewmGLM` but returns a state parameter for instantiation of later calculations. See `ewmGLM` documentation for more details

`pyg.timeseries._ewm.ewmskew_` (a, n, time=None, bias=False, min\_sample=0.25, axis=0, data=None, instate=None)

Equivalent to `ewmskew` but returns a state parameter for instantiation of later calculations. See `ewmskew` documentation for more details

## ***Index handling***

### ***df\_fillna***

`pyg.timeseries._index.df_fillna` (df, method=None, axis=0, limit=None)

Equivalent to `df.fillna()` except:

- support `np.ndarray` as well as dataframes
- support multiple methods of filling/interpolation
- supports removal of nan from the start/all of the timeseries
- supports action on multiple timeseries

**Parameters:**

**df** : dataframe/numpy array

**method** : string, list of strings or None, optional

Either a fill method (bfill, ffill, pad) Or an interpolation method: 'linear', 'time', 'index', 'values', 'nearest', 'zero', 'slinear', 'quadratic', 'cubic', 'barycentric', 'krogh', 'spline', 'polynomial', 'from\_derivatives', 'piecewise\_polynomial', 'pchip', 'akima', 'cubicspline' Or 'fnna': removes all to the first non nan Or 'nona': removes all nans

**axis** : int, optional

axis. The default is 0.

**limit** : TYPE, optional

when filling, how many nan get filled. The default is None (indefinite)

**Example:** method ffill or bfill

```
>>> from pyg import *; import numpy as np
>>> df = np.array([np.nan, 1., np.nan, 9, np.nan, 25])
>>> assert eq(df_fillna(df, 'ffill'), np.array([ np.nan, 1., 1., 9., 9., 25.]))
>>> assert eq(df_fillna(df, ['ffill','bfill']), np.array([ 1., 1., 1., 9., 9., 25.]))
>>> assert eq(df_fillna(df, ['ffill','bfill']), np.array([ 1., 1., 1., 9., 9., 25.]))
```

```
>>> df = np.array([np.nan, 1., np.nan, np.nan, np.nan, np.nan, np.nan, np.nan, 9, np.nan,
>>> assert eq(df_fillna(df, 'ffill', limit = 2), np.array([np.nan, 1., 1., 1., np.nan, :
```

`df_fillna` does not maintain state of latest 'prev' value: use `ffill_` for that.

**Example:** interpolation methods

```
>>> from pyg import *; import numpy as np
>>> df = np.array([np.nan, 1., np.nan, 9, np.nan, 25])
>>> assert eq(df_fillna(df, 'linear'), np.array([ np.nan, 1., 5., 9., 17., 25.]))
>>> assert eq(df_fillna(df, 'quadratic'), np.array([ np.nan, 1., 4., 9., 16., 25.]))
```

**Example:** method = fnna and nona

```
>>> from pyg import *; import numpy as np
>>> ts = np.array([np.nan] * 10 + [1.] * 10 + [np.nan])
>>> assert eq(df_fillna(ts, 'fnna'), np.array([1.]*10 + [np.nan]))
>>> assert eq(df_fillna(ts, 'nona'), np.array([1.]*10))
```

```
>>> assert len(df_fillna(np.array([np.nan]), 'nona')) == 0
>>> assert len(df_fillna(np.array([np.nan]), 'fnna')) == 0
```

**Returns:**

array/dataframe with nans removed/filled

## df\_index

`pyg.timeseries._index.df_index` (seq, index='inner')

Determines a joint index of multiple timeseries objects.

**Parameters:**

**seq** : *sequence whose index needs to be determined*

a (possible nested) sequence of timeseries/non-timeseries object within lists/dicts

**index** : *str, optional*

method to determine the index. The default is 'inner'.

**Returns:**

**pd.Index**

The joint index.

**Example:**

```
>>> tss = [pd.Series(np.random.normal(0,1,10), drange(-i, 9-i)) for i in range(5)]
>>> more_tss_as_dict = dict(zip('abcde', [pd.Series(np.random.normal(0,1,10), drange(-i, 9-
>>> res = df_index(tss + [more_tss_as_dict], 'inner')
>>> assert len(res) == 6
>>> res = df_index(more_tss_as_dict, 'outer')
>>> assert len(res) == 14
```

## df\_reindex

pyg.timeseries.\_index.**df\_reindex** (ts, index=None, method=None, limit=None)

A slightly more general version of df.reindex(index)

**Parameters:**

**ts** : *dataframe or numpy array (or list/dict of theses)*

timeseries to be reindexed

**index** : *str, timeseries, pd.Index.*

The new index

**method** : *str, list of str, float, optional*

various methods of handling nans are available. The default is None. See df\_fillna for a full list.

**Returns:**

**timeseries/np.ndarray (or list/dict of theses)**

timeseries reindex.

**Example:** index = inner/outer

```
>>> tss = [pd.Series(np.random.normal(0,1,10), drange(-i, 9-i)) for i in range(5)]
>>> res = df_reindex(tss, 'inner')
>>> assert len(res[0]) == 6
>>> res = df_reindex(tss, 'outer')
>>> assert len(res[0]) == 14
```

**Example:** index provided

```
>>> tss = [pd.Series(np.random.normal(0,1,10), drange(-i, 9-i)) for i in range(5)]
>>> res = df_reindex(tss, tss[0])
>>> assert eq(res[0], tss[0])
>>> res = df_reindex(tss, tss[0].index)
>>> assert eq(res[0], tss[0])
```

## presync

pyg.timeseries.\_index.**presync** ()

Much of timeseries analysis in Pandas is spent aligning multiple timeseries before feeding them into a function. presync allows easy presyncing of all paramters of a function.

**Parameters:**

**function** : *callable, optional*

function to be presynched. The default is None.

**index** : *str, optional*

index join policy. The default is 'inner'.

**method** : *str/int/list of these, optional*

method of nan handling. The default is None.

**columns** : *str, optional*

columns join policy. The default is 'inner'.

**default** : *float, optional*

value when no data is available. The default is np.nan.

**Returns:**

presynch-decorated function

**Example:**

```
>>> from pyg import *
>>> x = pd.Series([1,2,3,4], drange(-3))
>>> y = pd.Series([1,2,3,4], drange(-4,-1))
>>> z = pd.DataFrame([[1,2],[3,4]], drange(-3,-2), ['a','b'])
>>> addition = lambda a, b: a+b
```

#We get some nonsensical results:

```
>>> assert list(addition(x,z).columns) == list(x.index) + ['a', 'b']
```

#But:

```
>>> assert list(presynch(addition)(x,z).columns) == ['a', 'b']
>>> res = presynch(addition, index='outer', method = 'ffill')(x,z)
>>> assert eq(res.a.values, np.array([2,5,6,7]))
```

**Example 2:** alignment works for parameters 'buried' within...

```
>>> function = lambda a, b: a['x'] + a['y'] + b
>>> f = presynch(function, 'outer', method = 'ffill')
>>> res = f(dict(x = x, y = y), b = z)
>>> assert eq(res, pd.DataFrame(dict(a = [np.nan, 4, 8, 10, 11], b = [np.nan, 5, 9, 11, 12]
```

**Example 3:** alignment of numpy arrays

```
>>> addition = lambda a, b: a+b
>>> a = presynch(addition)
>>> assert eq(a(pd.Series([1,2,3,4], drange(-3)), np.array([[1,2,3,4]]).T), pd.Series([2,
>>> assert eq(a(pd.Series([1,2,3,4], drange(-3)), np.array([1,2,3,4])), pd.Series([2,4,6,
>>> assert eq(a(pd.Series([1,2,3,4], drange(-3)), np.array([[1,2,3,4],[5,6,7,8]]).T), pd.
>>> assert eq(a(np.array([1,2,3,4]), np.array([[1,2,3,4]]).T), np.array([2,4,6,8]))
```

**Example 4:** inner join alignment of columns in dataframes by default

```
>>> x = pd.DataFrame({'a':[2,4,6,8], 'b':[6,8,10,12.]}, drange(-3))
>>> y = pd.DataFrame({'wrong':[2,4,6,8], 'columns':[6,8,10,12]}, drange(-3))
>>> assert len(a(x,y)) == 0
>>> y = pd.DataFrame({'a':[2,4,6,8], 'other':[6,8,10,12.]}, drange(-3))
>>> assert eq(a(x,y),x[['a']]*2)
>>> y = pd.DataFrame({'a':[2,4,6,8], 'b':[6,8,10,12.]}, drange(-3))
>>> assert eq(a(x,y),x*2)
>>> y = pd.DataFrame({'column name for a single column dataframe is ignored':[1,1,1,1]}, d
>>> assert eq(a(x,y),x+1)
```

```
>>> a = presync(addition, columns = 'outer')
>>> y = pd.DataFrame({'other':[2,4,6,8], 'a':[6,8,10,12]}, drange(-3))
>>> assert sorted(a(x,y).columns) == ['a','b','other']
```

**Example 4:** ffilling, bfilling

```
>>> x = pd.Series([1.,np.nan,3.,4.], drange(-3))
>>> y = pd.Series([1.,np.nan,3.,4.], drange(-4,-1))
>>> assert eq(a(x,y), pd.Series([np.nan, np.nan,7], drange(-3,-1)))
```

but, we provide easy conversion of internal parameters of presync:

```
>>> assert eq(a.ffill(x,y), pd.Series([2,4,7], drange(-3,-1)))
>>> assert eq(a.bfill(x,y), pd.Series([4,6,7], drange(-3,-1)))
>>> assert eq(a.oj(x,y), pd.Series([np.nan, np.nan, np.nan, 7, np.nan], drange(-4)))
>>> assert eq(a.oj.ffill(x,y), pd.Series([np.nan, 2, 4, 7, 8], drange(-4)))
```

**Example 5:** indexing to a specific index

```
>>> index = pd.Index([dt(-3), dt(-1)])
>>> a = presync(addition, index = index)
>>> x = pd.Series([1.,np.nan,3.,4.], drange(-3))
>>> y = pd.Series([1.,np.nan,3.,4.], drange(-4,-1))
>>> assert eq(a(x,y), pd.Series([np.nan, 7], index))
```

**Example 6:** returning complicated stuff

```
>>> from pyg import *
>>> a = pd.DataFrame(np.random.normal(0,1,(100,10)), drange(-99))
>>> b = pd.DataFrame(np.random.normal(0,1,(100,10)), drange(-99))
```

```
>>> def f(a, b):
>>>     return (a*b, ts_sum(a), ts_sum(b))
```

```
>>> old = f(a,b)
>>> self = presync(f)
>>> args = (); kwargs = dict(a = a, b = b)
>>> new = self(*args, **kwargs)
>>> assert eq(new, old)
```

***add/sub/mul/div/pow operators***

```
pyg.timeseries._index.add_(a, b)
    addition of a and b supporting presynching (inner join) of timeseries

pyg.timeseries._index.mul_(a, b)
    multiplication of a and b supporting presynching (inner join) of timeseries

pyg.timeseries._index.div_(a, b)
    division of a by b supporting presynching (inner join) of timeseries

pyg.timeseries._index.sub_(a, b)
    subtraction of b from a supporting presynching (inner join) of timeseries

pyg.timeseries._index.pow_(a, b)
    equivalent to a**b supporting presynching (inner join) of timeseries
```

**Tutorials**

Below are some tutorials covering some aspects of pyg that may not be obvious. All the tutorials are active python notebooks available in the `../docs/lab/` directory.

## pyg.base.Dict

There are a few existing dict-extensions similar to Dict (a nice example is <https://github.com/mewwts/addict>) but Dict has a little more up its sleeve.

### initialization

```
from pyg import *
Dict(a = 1, b = 2, c = 3)
```

```
Dict(a = 1)(b = 2, c = 3)
```

```
Dict(a = 1)(b = 2)(c = lambda a, b: a+b)
```

```
Dict(a = 1) + dict(b = 2, c = 3)
```

### members access

```
d = Dict(a = 1, b = 2, c = 3)
```

```
d.a
```

```
d['a', 'b']
```

```
d[['a', 'b']]
```

But the fun starts when Dict allows you to access **functions** of its keys:

```
d[lambda a, b: a + b]
```

It is important to note that by making `d['a', 'b']` access both 'a' and 'b' keys, we abandon the right to have tuples as keys.

```
d = Dict({'a', 'b'} : 1)
import pytest
with pytest.raises(KeyError): # Dict will be trying to grab 'a' and 'b' separately
    d[('a', 'b')]
```

### adding

```
Dict(a = 1, b = 2) + dict(b = 3, c = 4) # like .update() but not in-place
```

But addition is subtly different from update in the case of tree structure:

```
tree = Dict(a = 1, b = Dict(c = 2, d = 3))
update = dict(x = 1, b = dict(c = 'new value for b.c but keep b.d', e = 4))
tree+update
```

Tree updating is actually important enough to have its own function that can operate on dict-trees

```
tree = dict(a = 1, b = dict(c = 2, d = 3)) # I only use dicts
tree_update(tree, update) # but I can still update it like a tree
```

### subtracting

You can subtract keys or list of keys

```
Dict(a = 'remove me', b = 2, c = 3) - 'a' # subtracting a key
```

```
Dict(a = 'I am gone', b = 'and so am I', c = 3) - ['a', 'b'] # subtracting a collection of keys
```

```
tree = Dict(a = 1, b = Dict(c = 'delete me', d = 'but keep me'), c = 3)
tree - ('b', 'c') ## subtracting a branch in a tree using a tuple, possible because we know
```

## modifying the keys: rename

```
Dict(a = 1, b = 2).rename('prefix_') # need to be done sufficient
```

```
Dict(a = 1, b = 2).rename('_suffix')
```

```
Dict(a = 1, b = 2).rename(upper)
```

```
Dict(a = 1, b = 2, c = 3).rename(a = 'Abraham', b = 'Barbara')
```

## modifying the values: do

```
Dict(a = 1, b = 2, c = 3).do(lambda x: x**2) # modify all values using function
```

```
Dict(a = 1, b = 2, c = 3).do([lambda x: x**2, lambda x: x-1]) # modify all values using list
```

```
Dict(a = 1, b = 2, c = 3).do([lambda x: x**2, lambda x: x-1], 'a', 'b') # modify selected keys
```

## Dict can store a calculation flow

Being able to access function of members means we can think of a Dict as a container of variables. Consider this code:

```
def func(a, b):
    c = a + b
    d = b + c
    e = a/b + d/c
    f = (d+e)/c
    return f
func(1,2)
```

How can we keep track of our calculations and debug it easily? Consider rewriting this:

```
x = Dict(a = 1, b = 2)
x = x(c = lambda a, b: a + b)
x = x(d = lambda b, c: b + c)
x = x(e = lambda a,b,c,d : a/b + d/c)
x = x(f = lambda c,d,e: (d+e)/c)
x
```

We have all the internals of the function exposed and we are able to separate calculation flow and data easily:

```
calculation_pipeline = dict(c = lambda a, b: a + b,
                             d = lambda b, c: b + c,
                             e = lambda a,b,c,d : a/b + d/c,
                             f = lambda c,d,e: (d+e)/c)

initial_values = Dict(a = 1, b = 2)

initial_values(**calculation_pipeline)
```

You can see in `pyg.base.dictable` tutorial how this is extended



## pyg.base.dictable

dictable is a table, a collection of iterable records. It is also a dict with each key's value being a column. Why not use a pandas.DataFrame? pd.DataFrame leads a dual life:

- by day an index-based optimized numpy array supporting e.g. timeseries analytics etc.
- by night, a table with keys supporting filtering, aggregating, pivoting on keys as well as inner/outer joining on keys.

As a result, the pandas interface is somewhat cumbersome. Further, the DataFrame isn't really designed for containing more complicated objects within it. Conversely, dictable only tries to do the latter and is designed precisely for holding entire research process in one place. You can think of dictable as 'one level up' on a DataFrame: a dictable will handle thousands of data frames within it with ease. Indeed, dictable should be thought of as an 'organiser of research flow' rather than as an array of primitives. In general, each row will contain some keys indexing the experiment, while some keys will contain complicated objects: a pd.DataFrame, a timeseries, yield\_curves, machine-learning experiments etc. The interface is succinct and extremely intuitive, allowing the user to concentrate on logic of the calculations rather than boilerplate.

### *Motivation: dictable as an organiser of research flow*

We start with a simple motivating example. Here is a typical workflow:

```
from pyg import *; import pandas as pd; import numpy as np
import yfinance as yf
```

```
symbols = ['MSFT', 'WMT', 'TSLA', 'AAPL', 'BAD_SYMBOL', 'C']
history = [yf.download(symbol) for symbol in symbols]
prices = [h['Adj Close'] for h in history]
rtns = [p.diff() for p in prices]
vols = [r.ewm(30).std() for r in rtns]
zscores = [r/v for r,v in zip(rtns, vols)]
zavgs = [z.mean() for z in zscores]
```

```
zavgs
```

At this point we ask ourselves: Why do we have a **nan**? Which ticker was it, and when did it go wrong?

```
bad_symbols = [s for s, z in zip(symbols, zavgs) if np.isnan(z)]; bad_symbols
```

Great, how do we remove bad symbols from all our other variables?

```
vols = [v for s, v in zip(symbols, vols) if s not in bad_symbols]
```

Now we can calculate some stuff with rtns and vols perhaps?

```
ewmas = [r.ewm(n).mean()/v for r,v in zip(rtns, vols) for n in [10, 20, 30]]
```

Things went wrong and went wrong silently too:

- We forgot to remove bad data from rtns as well as from vols so our zip function is zipping the wrong stocks together
- It is nearly impossible to discover what item in the list belong to what n and what stock

If you ever dealt with real data, the mess described above must be familiar.

## Same code, in dictable

```
from pyg import *
import yfinance as yf
s = dictable(symbol = ['MSFT', 'WMT', 'TSLA', 'AAPL', 'BAD_SYMBOL', 'C'])
s = s(history = lambda symbol: yf.download(symbol))
s = s(price = lambda history: history['Adj Close'])
s = s(rtn = lambda price: price.diff())
s = s(vol = lambda rtn: rtn.ewm(30).std())
s = s(zscore = lambda rtn, vol: rtn/vol)
s = s(zavg = lambda zscore: zscore.mean())
```

dictable `s` contains all our data.

- each row contains all the variables associated with a specific symbol
- each column corresponds to a variable
- adding a new variable is declarative and free of boiler-plate loop and zip

```
s[['symbol', 'history', 'vol', 'zavg']]
```

```
s.zavg
```

## Oh, no, we have a bad symbol, how do we remove it?

```
s = s.exc(zavg = np.nan); s.zavg
```

## Now if we want to calculate something per symbol and window...

We want to create a new table, now keyed on two values: symbol and window `n`, so we create a bigger table using cross product:

```
sn = s * dict(n = [10, 20, 30]) ## each row is now unique per symbol and window n
```

```
sn = sn(ewma = lambda rtn, n, vol: rtn.ewm(n).mean()/vol)
```

And here is Citibank's three ewma...

```
sn.inc(symbol = 'C')[['n', 'ewma']]
```

Here is a pivot table of the average of each ewma per symbol and window... Note that again, we can access functions of variables and not just the existing keys in the dictable

```
sn.pivot('symbol', 'n', lambda ewma: ewma.mean())
```

## dictable functionality

### construction

dictable is quite flexible on constructions.

```
d = dictable(a = [1, 2, 3, 4], b = ['a', 'b', 'c', 'd']); d
```

```
d = dictable(dict(a = [1, 2, 3, 4], b = ['a', 'b', 'c', 'd']), symbol = ['MSFT', 'AAPL', 'APA', 'C'])
```

```
df = pd.DataFrame(d) # can instantiate a DataFrame from a dictable with no code and vice versa
```

```
d = dictable(df); d
```

```
d = dictable([(1, 3), (2, 4), (3, 5)], ['a', 'b']); d # construction from records as tuples
```

```
d = dictable([dict(a = 1, b = 3), dict(a = 2, b = 4, d = 'new column'), dict(a = 3, b = 5, c = 6)])
d = dictable(read_csv('d:/dropbox/yoav/python/pyg/docs/constituents_csv.csv')); d = d[:6]; d
```

### row access

```
d[0] #returns a record
```

```
d[:2] ## subset rows using slice
```

```
for row in d: # iteration is by row
    print(row)
```

### column access

```
d.Name
```

```
d['Name']
```

```
d['Name', 'Sector']
```

```
d[['Name', 'Sector']]
```

### *d is a dict so supports the usual keys(), values() and items():*

```
for key, column in d.items():
    print(key, ': ', column)
```

access via **function** of variables is also supported

```
d[lambda Symbol, Sector: '%s, %s'%(Symbol, Sector)]
```

### column and row access are commutative

```
assert d[0].Name == d.Name[0] == '3M Company'
assert d[0][lambda Symbol, Sector: '%s, %s'%(Symbol, Sector)] == d[lambda Symbol, Sector: '%s, %s'%(Symbol, Sector)][0]
assert d[0]['Name'] == d['Name'][0]
assert d[:2]['Name', 'Sector'] == d['Name', 'Sector'][:2]
assert d[:2][['Name', 'Sector']] == d[['Name', 'Sector']][:2]
```

### adding records

```
d = dictable(name = ['alan', 'barbara', 'chris'], surname = ['abramson', 'brown', 'cohen'], age = [30, 40, 50])
```

```
d + dict(name = 'david', surname = 'donaldson', age = 4) ## adding a single record
```

```
d + [dict(name = 'david', surname = 'donaldson', age = 4), dict(name = 'evan', surname = 'emerson', age = 5)]
```

```
d + dict(name = ['david', 'evan'], surname = ['donaldson', 'emerson'], age = [4, 5])
```

```
d + pd.DataFrame(dict(name = ['david', 'evan'], surname = ['donaldson', 'emerson'], age = [4, 5]))
```

### adding/modifying columns

You can add a column or a constant by simply calling the dictable with the values:

```
d(gender = ['m', 'f', 'm'])(school = 'St Paul')
```

More interestingly, it can be a callable function using the other variables...

```
d = d(initials = lambda name, surname: name[0] + surname[0]); d
```

Given d is a dict, a more traditional way of setting a new key is by simple assignment:

```
d['initials'] = d(lambda name, surname: name[0] + surname[0]); d
```

Or you can use the dict.update method:

```
d.update(dict(gender = ['m', 'f', 'm'])); d
```

## do

Sometime we want to apply the same function(s) to a collection of columns. For this, 'do' will do nicely:

```
d = d.do(upper, 'initials', 'gender').do(proper, 'name', 'surname'); d
```

## removing columns

```
d = d - 'initials'; d
```

## removing rows

```
d.exc(name = 'Alan')
```

```
d.inc(name = ['Alan', 'Chris'])
```

```
d.inc(lambda age: age>1)
```

```
d.exc(lambda gender: gender == 'M')
```

```
d.exc(lambda name, surname: len(name)>len(surname))
```

## sort

```
d.sort('name', 'surname')
```

```
d.sort(lambda name: name[::-1]) # can sort on functions of variables too
```

## listby(keys)

listby is like groupby except it returns a dictable with unique keys and the other columns are returned as a list. We find that MUCH more useful usually than groupby

```
grades = dictable(name = ['alan', 'barbara', 'chris'], grades = [30,90,80], subject = 'english')
+ dictable(name = ['alan', 'david', 'esther'], grades = [40,50,70], subject = 'math',
+ dictable(name = ['barbara', 'chris', 'esther'], grades = [90,60,80], subject = 'fre
```

```
grades.listby('teacher')
```

```
grades.listby('teacher')(avg_grade = lambda grades: np.mean(grades))
```

## unlist

unlist undoes listby() assuming it is possible...

```
grades.listby('teacher').unlist()
```

## groupby(keys) and ungroup

This is similar to DataFrame groupby except that instead of a new object, a dictable is returned: The name of the grouped column is given by 'grp'. ungroup allows us to get back to original.

```
classes = grades.groupby(['teacher', 'subject'], grp = 'class')
```

```
classes[0]
```

```
classes.ungroup('class')
```

## inner join

The multiplication operation is overloaded for the join method. By default, if two dictables share keys, the join is an inner join on the keys

```
students = dictable(name = ['alan', 'barbara', 'chris', 'david', 'esthar', 'fabian'], surname = ['alan', 'barbara', 'chris', 'david', 'esthar', 'fabian'])
```

```
print('shared keys:', grades.keys() & students.keys())
grades * students
```

Are there students with no surname? We can do a xor or use division which is overloaded for xor:

```
grades / students
```

Are there students with no grades?

```
students / grades
```

```
students = dictable(name = ['Alan', 'Barbara', 'Chris', 'David', 'Esther', 'Fabian'], surname = ['alan', 'barbara', 'chris', 'david', 'esthar', 'fabian'])
```

We fixed Esther's spelling but introduced capitalization, that is OK, we are allowed to inner join on functions of keys too.

```
grades.join(students, 'name', lambda name: name.lower())
```

```
students = dictable(first_name = ['alan', 'barbara', 'chris', 'david', 'esther', 'fabian'], surname = ['alan', 'barbara', 'chris', 'david', 'esthar', 'fabian'])
```

You can inner join on different column names and both columns will be populated:

```
grades.join(pd.DataFrame(students), 'name', 'first_name')
```

## inner join (with other columns that match names)

By default, if columns are shared but are not in the join, they will be returned with a tuple containing both values

```
x = dictable(key = ['a', 'b', 'c', 'c'], x = [1,2,3,4], y = [4,5,6,7])
y = dictable(key = ['b', 'b', 'c', 'a'], x = [1,2,3,4], z = [8,9,1,2])
x.join(y, 'key', 'key') ## ignore x column for joining
```

```
x.join(y, 'key', 'key', mode = 'left') ## grab left value
```

## cross join

If no columns are shared, then a cross join is returned.

```
x = dictable(x = [1,2,3,4])
y = dict(y = [1,2,3])
x * y
```

```
x.join(y, [], []) ## you can force a full outer join
```

```
x / y == x
```

## xor (versus left and right join)

We find left/right join actually not very useful. There is usually a genuine reason for records for which there is a match and for records for which there isn't. And the treatment of these is distinct, which means a left-join operation that joins the two outcomes together is positively harmful.

The xor operator is much more useful and you can use it to recreate left/right join if we really must. Here is an example

```
students = dictable(name = ['alan', 'barbara', 'chris'], surname = ['abramsom', 'brown', 'co
new_students = dictable(name = ['david', 'esther', 'fabian'], surname = ['drummond', 'eckles

inner_join = grades * students ## grades with students
left_xor = grades / students ## grades without students

# you can...
left_join = grades * students + grades / students ## grades for which no surname is availabl
left_join

# but really you want to do:
student_grades = grades * students
unmapped_grades = grades / students ## we treat this one separately...
new_student_grades = unmapped_grades * new_students ## and grab surnames from the new stude

assert len(unmapped_grades / new_student_grades) == 0, 'students must exist either in the st

all_grades = student_grades + new_student_grades; all_grades
```

## pivot

```
x = dictable(x = [1,2,3,4])
y = dictable(y = [1,2,3,4])
xy = (x * y)
xy

xy.pivot('x', 'y', lambda x, y: x*y)
```

## a few observations:

- as per usual, can provide a function for values in table (indeed columns y) and not just keys
- the output in the cells come back as a list. This is because sometimes there are more than one row with given x and y, and sometimes there are none:

```
(xy + xy).exc(lambda x,y: x+y == 5).pivot('x', 'y', lambda x, y: x*y)
```

You can apply a sequence of aggregate functions:

```
(xy + xy).exc(lambda x,y: x+y == 5).pivot('x', 'y', lambda x, y: x*y, lambda v: len(v))
```

## pyg.mongo

MongoDB has replaced our SQL databases as it is just too much fun to use. MongoDB does have its little quirks:

- The MongoDB 'query document' that replaces the SQL WHERE statements is very powerful but you need a PhD for even the simplest of queries.
- too many objects we use (specifically, numpy and pandas objects) cannot be pushed directly easily into Mongo.
- Mongo lacks the concept of a table with primary keys. Unstructured data is great but much of how we think of data is structured.

pyg.mongo addresses all three issues:

- **q** is a much easier way to generate Mongo queries. We are happy to acknowledge TinyDB <https://tinydb.readthedocs.io/en/latest/usage.html#queries> for the idea.
- **mongo\_cursor** is a super-charged cursor and in particular, it handles encoding and decoding of objects seamlessly in a way that allows us to store all that we want in Mongo.
- **mongo\_pk\_cursor** manages a table with primary keys and full history audit. We are happy to acknowledge Arctic by the AHL Man team for the initial inspiration

## q

The MongoDB interface for query of a collection (table) is via a creation of a query document <https://docs.mongodb.com/manual/tutorial/query-documents/>. This is rather complicated for the average use. For example, if you wanted to locate James Bond in the collection, you would need to compose a query document that looks like this:

```
{ "$and": [ { "name": { "$eq": "James" } }, { "surname": { "$eq": "Bond" } } ] }
```

It's doable, but not much fun writing. Luckily... within the continuum you can write this instead:

```
from pyg import *; import re
q(name = 'James', surname = 'Bond')
```

```
(q.name == 'James') & (q.surname == 'Bond')
```

How do we create in MongoDB a query document to find all the James who are not Bond?

```
(q.surname != 'Bond') & (q.name == 'James')
```

```
~(q.surname == 'Bond') & (q.name == 'James')
```

What about records with no surname?

```
(q.name == 'James') - q.surname
```

```
q(q.surname.not_exists, name = 'James')
```

And what about records with james rather than James?

```
q(name = ['james', 'James'], surname = ['bond', 'Bond']) ## the result is long so it is repr
```

```
q(name = re.compile('^[J]ames'), surname = re.compile('^[B]ond'))
```

As you can see, q is callable and you can put expressions inside it, or you can use the q.key method.

If you have funny characters or spaces in your dict...

```
q['funny$text with # weird & characters'].exists
```

If your document is nested and there are subkeys, that is ok, you can use either:

```
(q['key.subkey'] >= 100) | ((q.key.other.exists) & (q.some.other.stuff == [1, 2]))
```

q does not have the full power of the Mongo query document but it will get you to 95% of what you want. We end with a fun James Bond query. If we want to find the bond films with all actors who played James Bond after 1980...

```
bonds = dictable(name = ['Daniel', 'Sean', 'Roger', 'Timothy'], surname = ['Craig', 'Connery'])
bonds
```

```
q(list(bonds), q.release_date > dt(1980))
```

## mongo\_cursor

The mongo cursor:

- enables saving seamlessly objects and data in MongoDB

- simplifies filtering
- simplifies projecting onto certain keys in document

## general objects insertion into documents

pymongo.Collection supports insertion of documents into it:

```
from pyg import *; import pymongo as pym; import pytest
c = pym.MongoClient()['test']['test']
c.drop() # drop all documents
c.insert_one(dict(a = 1, b = 2)) # insert a document
```

```
assert c.count_documents({}) == 1 # in order to count documents, must apply the empty query
```

We can do similar stuff with a mongo\_cursor:

```
t = mongo_table(table = 'test', db = 'test')
t.drop()
t.insert_one(dict(a = 1, b = 2))
```

```
assert len(t) == 1 #no need to specify the filter, mongo_cursor keeps track of the current f
```

Annoyingly, raw pymongo.Collection cannot encode for lots of existing objects.

```
ts = pd.Series([1.,2.], drange(2000,1))
a = np.arange(3)
f = np.float32(32.0)
with pytest.raises(Exception):
    c.insert_one(dict(a = a)) # cannot insert an array
with pytest.raises(Exception):
    c.insert_one(dict(f = f)) # cannot insert a numpy float, string or bool
with pytest.raises(Exception):
    c.insert_one(dict(ts = ts)) # cannot insert a pd.Series or DataFrame
```

Further, unless we define the encoding, new classes do not work either

```
class NewClass():
    def __init__(self, n):
        self.n = n
    def __eq__(self, other):
        return type(other) == type(self) and self.n == other.n
n = NewClass(1)
with pytest.raises(Exception):
    c.insert_one(dict(n = n))
```

Luckily, the mongo\_cursor t can insert all these happily:

```
t.drop()
t.insert_one(Dict(a = a, f = f, ts = ts, n = n))
assert len(t) == 1
t[0] ## reading it back
```

## document reading

What is nice is that when you read the document using the mongo\_cursor, you get back the **object** you saved, not just the data. Is this magic? Not really... We read the doc directly from the Collection:

```
raw_doc = c.find_one({})
assert raw_doc['n'] == '{"py/object": "__main__.NewClass", "n": 1}'
assert encode(n) == '{"py/object": "__main__.NewClass", "n": 1}'
assert decode('{"py/object": "__main__.NewClass", "n": 1}') == n
assert t.writer == encode
assert t.reader == decode
```



- When writing, the `mongo_cursor` encodes the objects pre-saving it into Mongo, in this case as a simple dict
- When reading, it uses `decode` to convert what it reads back into the object
- This is done transparently though you can have full control via specifying writer/reader functions

This all works with the assumption that the person loading and the person saving share the library so objects can be instantiated on load. If construction method has changed and the object is not back-compatible, then user will receive the undecoded object and a warning message is logged.

## document writing to files

MongoDB is great for manipulating/searching dict keys/values. The actual dataframes in each doc, we may want to save in a file system because:

- DataFrames are stored as bytes in MongoDB anyway, so they are not searchable
- MongoDB free version has limitations on size of document
- For data licensing issues, data must not sit on servers but needs to be stored on local computer
- Storing in files allows other non-python/non-MongoDB users easier access, allowing data to be detached from app. In particular, if you want to stream messages into the array/dataframe, doing it through Mongo is probably the wrong way about it. <https://github.com/man-group/arctic> attempts to do it but Mongo should probably just contain a reference to a file. You then have a listener such as OMQ appending new messages into the file (perhaps via <https://github.com/xor2k/np-append-array/> or `aws wrangler`). This will be (a) more performant, (b) require no code, and (c) new data will then magically show up in Mongo every time you read the document.

```
t2 = mongo_table('test', 'test', writer = 'parquet')
t2.drop()
doc = dict(root = 'c:/temp', a = [a,a,a], ts = dict(one = ts, two = ts), f = f, n = n) ## c
t2.insert_one(doc)
encoded = c.find_one({})
print(tree_repr(encoded))
```

You can see that starting at the root location, the document's numpy arrays and pandas have been saved to .npz and .parquet files

```
print(tree_repr(decode(encoded)))
```

```
np.load('c:/temp/a/2.npy') ## can load data directly
```

```
pd.read_parquet('c:/temp/ts/one.parquet')
```

## document access

We start by pushing a 10x10 times table into t

```
t.drop()
times_table = (dictable(a = range(10)) * dictable(b = range(10)))(c = lambda a, b: a*b)
t.insert_many(times_table)
```

## filters

We now examine how we drill down to the document(s) we want:

```
assert len(t.inc(a = 1)) == 10
assert len(t.exc(a = 1)) == 90
assert isinstance(t.inc(a = 1), mongo_cursor) ## it is chain-able
assert len(t.find(q.a == 1).find(q.b == [1,2,3,4])) == 4
```

We can use the original collection too but not in a chain-like fashion:

```
spec = q(a = 1, b = [1,2,3,4])
assert c.count_documents(spec) == 4
```

```
c.find(spec) # That is OK
with pytest.raises(AttributeError): # not OK, cannot chain queries
    c.find(q(a=1)).find(q(b = [1,2,3,4]))
```

## iteration

Just like a mongo.Cursor, c.find(spec), t is also iterable over the documents:

```
sum([doc for doc in t.find(a = 1).find(b = [1,2,3,4])], dictable())
```

```
dictable(t.find(a = 1).find(b = [1,2,3,4])) ## or just put a cursor straight into a table
```

```
t.find(a = 1).find(b = [1,2,3,4])[::] ## or simple slicing
```

## sorting

```
t.sort('c', 'b')[::]
```

## getitem of a specific document

```
t[dict(a = 7, b = 8)]
```

## column access

```
t.b
```

```
assert t.b == t.distinct('b') == c.distinct('b')
```

In MongoDB the cursor can have a 'projection' onto specific columns. In mongo\_cursor this is simplified:

```
t[['a', 'b']].find(c = 12)[::]
```

## add/remove columns

```
del t['c']
t[::]
```

```
t = t.set(c = 'not very useful but...')
t[::]
```

```
t = t.set(c = lambda a, b: a * b) ### more useful
t[::]
```

## add/remove records

```
t.inc(c = 12).drop()
t
```

```
t = t + dict(a = 2, b = 6, c = 12)
t
```

```
t = t.inc(c = 12).drop() + times_table.inc(c = 12) ## adding four records at once
t
```

```
t = t.inc(c = 12).drop().insert_many(times_table.inc(c = 12))
t[::]
```

```
t = t.raw ## remove the filter c = 12
t
```

## mongo\_pk\_table

mongo\_pk\_table is a mongo\_cursor implementing a table with primary keys. Suppose we want to have a table of people:

```
from pyg import *; import pymongo as pym; import pytest

t = mongo_table(table = 'test', db = 'test')
c = pym.MongoClient()['test']['test']
pk = mongo_table(table = 'test', db = 'test', pk = ['name', 'surname'])

t.drop()
d = dictable(name = ['alan', 'alan', 'barbara', 'chris'], surname = ['adams', 'jones', 'brown'])
pk.insert_many(d)
pk[::]
```

Now let us suppose a year has passed...

```
pk.set(age = lambda age: age + 1)
pk[::]
```

The pk-table actually maintains complete audit trail. Older records are not deleted, they just get '\_deleted' parameter set for them.

```
print(dictable(c))
```

You can see pk only looks at records where \_deleted does not exist and \_pk are set.

```
pk
```

There are obviously some small differences on how pk works but broadly, it is just like a normal mongo\_cursor with an added filter to zoom onto the records that maintain the primary-key table:

- you cannot insert docs without primary keys all present:
- the drop() command does not actually delete the documents, they are simply 'marked' as deleted.
- to get from a mongo\_pk\_cursor to mongo\_cursor simply do pk.raw

```
with pytest.raises(KeyError):
    pk.insert_one(dict(no_name_or_surname = 'James')) # cannot insert with no PK
```

```
pk.drop()
len(pk)
```

```
t[::] ## the data is there, it is just marked as _deleted
```

## mongo\_reader and mongo\_pk\_reader

Because it is so easy to do stuff in MongoDB, we could easily cause damage to the data underlying. We therefore also introduced read-only versions for the mongo\_cursor and pk\_cursor:

```
pkr = mongo_table(table = 'test', db = 'test', pk = ['name', 'surname'], mode = 'r')
pkr
```

```
with pytest.raises(AttributeError):
    pkr.drop()
```

```
r = mongo_table(table = 'test', db = 'test', mode = 'r')
with pytest.raises(AttributeError):
    r.drop()
```

```
r[::]
```

## pyg.base.cell

cell is a dict that forms part of a calculation graph. Most usefully, db\_cell is implemented to maintain persistency of the function output in MongoDB. Before we start, we will show a few examples of how a cell works. Then, we will build a toy example of trading stocks based on an exponentially weighted crossover.

- We will start by creating the system using pyg.base.dictable and pyg.timeseries.
- We then repeat the same code, this time modifying it slightly to save the data and calculation graph in MongoDB while running the calculation.
- We conclude by discussing the two approaches

### Cell 101

```
from pyg import *
a = cell(lambda x, y: x + y, x = 1, y = 2)
b = cell(lambda x, y: x * y, x = 2, y = a)
b
```

```
b.keys() ## b is a dict
```

```
b._args ## inputs
```

```
b._output ## where the output will go once we calculate it
```

```
assert b.run() ## b has not calculated yet... please run it
```

```
b() # calculated object note b().data
```

```
assert not b().run() ## b has calculated now... no need to run it
```

```
cell(lambda x, y: x ** y)(x = a, y = 2) # you can define the cell and then call it with the
```

### Workflow without saving to the database

```
from pyg import *;
import yfinance as yf # see https://github.com/ranaroussi/yfinance
constituents = dictable(read_csv('d:/dropbox/yoav/python/pyg/docs/constituents_csv.csv')).re
constituents
```

```
stocks = constituents.inc(sector = 'Energy')
stocks
```

```
stocks = stocks(history = lambda symbol, sector, name: yf.download(tickers = symbol))
```

```
stocks = stocks.inc(lambda history: len(history)>0)
```

```
stocks = stocks(adj = lambda history: getitem(value = history, key = 'Adj Close'))
```

```
stocks = stocks(rtn = lambda adj: diff(a = adj))
```

```
stocks = stocks(vol = lambda rtn: ewmstd(a = rtn, n = 30))
```

```
_data = 'data'
def crossover_(a, fast, slow, vol, instate = None):
    state = Dict(fast = {}, slow = {}, vol = {}) if instate is None else instate
    fast_ewma_ = ewma_(a, fast, instate = state.fast)
    slow_ewma_ = ewma_(a, slow, instate = state.slow)
    raw_signal = fast_ewma_.data - slow_ewma_.data
    signal_rms = ewmrms_(raw_signal, vol, instate = state.vol)
    normalized = raw_signal/v2na(signal_rms.data)
    return Dict(data = normalized, state = Dict(fast = fast_ewma_.state, slow = slow_ewma_.state, vol = vol_ewma_.state))

crossover_.output = ['data', 'state']

def crossover(a, fast, slow, vol, state = None):
    return crossover_(a, fast, slow, vol, instate = state)
```

### *some more functions to calculate the profits & loss as well as the signal/noise ratio*

```
def signal_pnl(signal, rtn, vol):
    return shift(signal) * (rtn/vol)

def information_ratio(pnl):
    return 16 * ts_mean(pnl) / ts_std(pnl)
```

```
forecasts = stocks * dictable(fast = [2,4,8], slow = [6,12,24], forecast = ['fast', 'medium', 'slow'])
```

```
forecasts = forecasts(signal = lambda rtn, fast, slow: crossover_(rtn, fast = fast, slow = slow, vol = vol))
```

```
forecasts = forecasts(pnl = lambda signal, rtn, vol: signal_pnl(signal = signal, rtn = rtn, vol = vol))
```

```
forecasts = forecasts(ir = lambda pnl: information_ratio(pnl = pnl))
```

```
print(forecasts.pivot('symbol', 'forecast', 'ir', [last, f12]))
```

## *Workflow while saving to MongoDB*

### *Table creation*

We create three tables depending on the primary keys we will be using.

```
idb = partial(mongo_table, db = 'demo', table = 'items', pk = 'item')
sdb = partial(mongo_table, db = 'demo', table = 'stock', pk = ['item', 'symbol'])
fdb = partial(mongo_table, db = 'demo', table = 'forecast', pk = ['item', 'symbol', 'forecast'])

idb().insert_one(Dict(item = 'constituents', data = constituents))
```

### *Any code differences?*

Most of the code remains the same as above, except:

- We wrap it inside a periodic\_cell so it is calculated daily
- We add reference to where we want to store it in MongoDB by specifying the db as well as the primary keys of that table
- To run the function, we need to call the cell. This: loads the cell from the database (if found), checking if it even needs running and if so, runs it.

```
stocks = stocks(history = lambda symbol, sector, name: periodic_cell(yf.download, tickers =
                                db = sdb, item = 'history', symbol = symbol))
```

## Accessing the data in MongoDB

The data is now in the database and can be accessed:

```
get_data('stock', 'demo', symbol = 'MMM')
```

```
stocks = stocks.inc(lambda history: len(history.data)>0)
```

```
stocks = stocks(adj = lambda history, symbol: periodic_cell(getitem, value = history, key =
                                db = sdb, symbol = symbol, item =
```

```
stocks = stocks(rtn = lambda adj, symbol: periodic_cell(diff, a = adj,
                                db = sdb, symbol = symbol, item = 'r
```

```
stocks = stocks(vol = lambda rtn, symbol: periodic_cell(ewmstd, a = rtn, n = 30,
                                db = sdb, symbol = symbol, item = 'v
```

## Calculating the forecasts & saving them

```
forecasts = stocks * dictable(fast = [2,4,8], slow = [6,12,24], forecast = ['fast', 'medium'])
```

```
forecasts = forecasts(signal = lambda rtn, fast, slow, symbol, forecast: periodic_cell(cross
                                db = fdb, symbol = symbol, forecast = forecast,
```

```
forecasts = forecasts(pnl = lambda signal, rtn, vol, symbol, forecast: periodic_cell(signal
                                db = fdb, symbol = symbol, forecast =
```

```
forecasts = forecasts(ir = lambda pnl: information_ratio(pnl = pnl.data))
```

```
print(forecasts.pivot('symbol', 'forecast', 'ir', [last, f12]))
```

## Accessing & running the graph once the graph has been created

We can access the data or the cell:

```
get_cell('forecast', 'demo', symbol = 'APA', forecast = 'fast', item = 'signal')
```

And now that the graph has been created, you can actually trigger it just by loading. i.e. The code below will give you the fast signal for APA and will ensure it is up-to-date too:

```
c = get_cell('forecast', 'demo', symbol = 'APA', forecast = 'fast', item = 'signal')
c = c.go()
print(c.data)
```

## Point-in-time, cache and persistency

The pk-tables save a full history of all your data. To avoid hitting the database all the time, we also have a local GRAPH singleton that caches all the cells by their address. The cell has few basic operations we need to understand:

- **cell.run()**: Returns True/False if the cell needs to be calculated. `db_cell()` just check for values in its output, `periodic_cell` will also check if it is a new business day.
- **cell.go()**: This calculates the cell and saves the result to the database (and to GRAPH). The cell itself is not loaded but all its inputs are loaded
  - `cell.go(0)` : calculate only if there is a need

- `cell.go(1)` : calculate me but my parents only if there is a need
- `cell.go(2)` : calculate me & my parents but my grandparents only if there is a need
- `cell.go(-1)` : calculate everything
- **`cell.load()`**: This loads the data from GRAPH, if not in GRAPH, loads it from MongoDB (and updates also the GRAPH)
  - `cell.load(-1)` : Clear the data from the GRAPH
  - `cell.load(0)` : Load & update me from GRAPH, if not, from MongoDB, if not, just return good old me
  - `cell.load(1)` : If you cannot find me, throw an Exception
  - `cell.load(date)` : Load my version as valid on date. If none exists, throw.
  - `cell.load([value])`: Force GRAPH to clear, only load from DB. Same as `cell.load(-1).load(value)`
- **`cell()`**: This loads & then go

```
from pyg import *; from functools import partial
db = partial(mongo_table, db = 'demo', table = 'persistency', pk = 'key')
db().raw.drop()
```

```
def f(a, b):
    return a+b
```

```
## now we set up a fake calculation tree:
x = db_cell(f, a = 1, b = 2, db = db, key = 'x')
y = db_cell(f, a = x, b = 2, db = db, key = 'y')
z = db_cell(f, a = x, b = y, db = db, key = 'z')

## and run it by running the final value we want
z = z()
```

```
## we can access the data:
get_data('persistency', 'demo', key = 'x')
```

```
t0 = dt() ## first breakpoint
```

```
x = db_cell(f, a = 10, b = 20, db = db, key = 'x').go()
y = db_cell(f, a = x, b = 20, db = db, key = 'y').go()
z = db_cell(f, a = x, b = y, db = db, key = 'z').go()
```

```
## and here is the new data
get_data('persistency', 'demo', key = 'x')
```

```
## and here is the data valid at our first breakpoint. get_data/get_cell always go to the da
get_data('persistency', 'demo', key = 'x', _deleted = t0)
```

We can ask a cell to load itself, but remember: it will go to GRAPH first by default. The GRAPH has only one copy of the cell, while in MongoDB, every time we recalculate/save a new version of the cell, we mark the old version in the database as “deleted” but otherwise keep it. To force a cell to load itself from the database, encase the breakpoint in a list...

```
db_cell(db = db, key = 'x').load([t0])
```

We can force a full recalculation of the tree in a single line of code:

```
db_cell(db = db, key = 'z')(go = -1, mode = [t0]).data ## Should be 8, same as the old value
```

## Comparison of the two workflows

Saving to the database has negatives:

- does require some (but really not much) additional code to specify where each data item goes to
- slows down the calculation

Conversely,

- We get full persistency: We can access each part of the graph with full visibility on the inputs, the function used to calculate the result, the function output(s), the location of where the data is stored and the time it was last updated as well as the periodicity it is calculated.
- We get full audit, past calculations remain available to track (and indeed, rerun) if anything goes wrong
- Each node will manage its schedule, ensuring data is up-to-date
- We can run just the parts of the graph we are interested in (and can run in parallel)

## To save or not to save?

Luckily we don't really need to decide on one workflow or the other as both can happily coexist. We can build a calculation graph and decide that some key points in the calculation we want to save while intermediate calculations we can calculate on the fly and not save at all. We have met the crossover function. Here we implement it 'on the fly' while saving just final value to db

```
from pyg import *; import pandas as pd; import numpy as np; from functools import partial
```

```
def fake_ts(ticker):
    return pd.Series(np.random.normal(0,1,1000), drange(-999))
db = partial(mongo_table, db = 'test', table = 'test', pk = ['key'])
db().raw.drop()
```

```
appl = db_cell(fake_ts, ticker = 'appl', key = 'appl_rtn', db = db)()

#I am never saving these, In fact, I don't want to see these in calculation log.
a = cell(ewma, a = appl, n = 30)
b = cell(ewma, a = appl, n = 50)

# I may want to save these nodes but haven't made up my mind
# I do want to see the calculations in the log though...
# I replace db by the primary keys of table (here 'key').
# This allows as to see the calculation log as it happens.
# data is not saved to db though until I switch to db = db as opposed to db = 'key'

c = db_cell(sub_, a = a, b = b, key = 'calculate difference of ewma', db = 'key')
d = db_cell(ewmrms, a = c, n = 100, key = 'root mean square of difference', db = 'key')

# The final crossover I definitely want to save to db:
final_value = db_cell(div_, a = c, b = d, key = 'appl_crossover', db = db)()
```

```
db().key
```

So although we had several intermediate steps, we decided to save just the final crossover in the database. If we look at the inputs for the function, you can see that the values are not saved in the database, though the full calculation tree *is*. Therefore, one can reload the node and then recalculate all the intermediate values on the fly

```
loaded_and_recalculated = (db()[dict(key = 'appl_crossover')] - 'data').go(1) ## but once re
assert_eq(loaded_and_recalculated.data, final_value.data)
```

## Behind the scene: cell\_func

Behind the scene of cell, there is machinery designed to make it work smoothly and transparently in most cases. However, sometimes the user may need to dig deeper. Here is an example for code that fails...

```
from pyg import *
import pytest
```



```
def twox(x):
    return x*2
a = cell(a = 1)
c = cell(twox, x = a)

with pytest.raises(KeyError):
    c()
```

c tries to run the function. The function demands parameter x. When looking at the cells provided, cell 'a' does not contain anything like 'x' so the function fails.

```
a = cell(data = 1)
cell(twox, x = a)()
```

'data' key has a preferred status so although 'x' is not in the cell, we assume but default that 'data' parameter is the one the cell wants to present to the world. This is controlled by cell\_output function:

```
cell_output(a)
```

```
a = cell(data = 1, myoutput = 3, output = 'myoutput') ## you can decide your output is different
cell_output(a), cell_item(a)
```

```
cell(twox, x = a)()
```

That is good but what happens if the cell has MORE than one output or we want to direct the function to grab another key?

```
a = cell(a = 1) ## this has failed...
cell(cell_func(twox, x = 'a'), x = a)() ## when you grab x, use 'a' as key
```

What if you need the cell itself rather than the items in it?

```
def add_a_and_b(x):
    return x.a + x.b

x = cell(a = 1, b = 2)

cell(cell_func(add_a_and_b, unitemized = 'x'), x = x)()
```

We can see that the cell x itself is presented to the function and x.a + x.b is calculated and data == 3

## pyg.base.join

Only read this if you are a seasoned dictable user. In data science, we usually have data in multiple tables and we want to pull specific columns together for an analysis. We will first look at **join** function and then examine the **perdictable** decorator.

### Join

#### Example: Using join function to transfer money to a bank

We begin by setting up a mini database:

```
from pyg import *
customers = dictable(customer = ['alan', 'barbara', 'charles'], address = ['1 Abba Avenue',
products = dictable(product = ['apple', 'banana', 'cherry'], price = [1,2,3], supplier = ['
customer_products = dictable(customer = ['alan', 'alan', 'charles', 'charles'], product = ['
banks = dictable(bank = ['allied', 'barclays'], account = [5556, 2461])

print('Customers\n', customers, '\n\nProducts\n', products, '\n\nCustomer_products\n', custo
```

### Simple join: inner join between tables

Suppose we want to know how much money is to be transferred from each bank. - We only care about the fields 'bank', 'amount' and 'price' - each field is pulled from different tables, - need to specify customer & product as the keys we will join on:

```
join(dict(bank = customers, amount = customer_products, price = products), on = ['customer',
```

### Defaults for fields we want to left-join on...

The function we need to run to transfer money looks like this, so actually, we would like to have account details too.

```
def transfer_money(bank, amount, price, account = 'default'):
    ## if account == 'default' transfer money slowly, else transfer quickly
    ## return
    pass
```

We can grab the account details from the 'banks' table:

```
join(dict(bank = customers, amount = customer_products, price = products, account = banks),
```

but we just **lost** Chase transactions as we don't have its account details. However, money is transferred perfectly (albeit slowly) even without account id. So instead....

```
join(dict(bank = customers, amount = customer_products, price = products, account = banks),
      on = ['customer', 'product', 'bank'],
      defaults = dict(account = 'default'))
```

### Renaming & calculating fields

We also want to ensure we don't transfer money that we already transferred... so we need to grab an expiry column based on purchase\_date in customer\_product table

```
join(dict(bank = customers, amount = customer_products, price = products, account = banks,
          on = ['customer', 'product', 'bank'],
          renames = dict(expiry = lambda purchase_date: dt(purchase_date, '1b')), ## it takes 1
          defaults = dict(account = 'default'))
```

### Perdictable

perdictable takes the same operation one step further and actually runs the function. We also use the function signature to determine the defaults parameter. Here is another example: ### Example: Oil prices In Finance, there are contracts called Futures, each Future contract has an expiry. E.g. Futures contracts for Oil are contracts agreeing the delivery of oil to a particular place in a particular month. Once that month is gone, that contract is no longer traded and the oil needs to be delivered.

```
from pyg import *
oil = dictable(y = dt().year-1, m = range(3, 13, 3)) + dictable(y = dt().year, m = range(3,
oil = oil(ticker = lambda y, m: 'OIL_%i_%s'%(y, m if m>9 else '0%i'%m))
oil
```

y,m and ticker will form our primary keys

```
pk = ['y', 'm', 'ticker']
expiry = perdictable(lambda y, m: dt(y,m+1,1), on = pk)(y = oil, m = oil)
expiry
```

```
def fake_ts(ticker, expiry):
    return 500 + pd.Series(np.random.normal(0,1,100), drange(dt_bump(expiry,-99), expiry)).c
```

To add a price for each of the futures, we first wrap fake\_ts and then run it:

```
price = perdictable(fake_ts, on = pk)(ticker = oil, expiry = expiry)
price
```

We have wrapped a function so that we get a price for **each** of these contracts. This allows us to move from operating on single timeseries, to run it on multiple rows from multiple tables

```
rtn = perdictable(diff, on = pk)(a = price, expiry = expiry)
yesterday_price = perdictable(shift, on = pk)(a = price, expiry = expiry)
percentage_return = perdictable(div_, on = pk)(a = rtn, b = yesterday_price, expiry = expiry)
percentage_return
```

### perdictable and caching

This is nice but (a) what have we gained? and (b) why do we keep using expiry as a variable? The answer is to do with caching actually. If we rerun prices, we should get brand new data, since fake\_ts just generates random prices... perdictable identifies rows that have been run and are now 'expired' It uses provided old data and does not recalculate. If either expiry or old values are not provided then it calculates everything.

```
new_price = perdictable(fake_ts, on = pk)(ticker = oil, data = price, expiry = expiry)
(new_price.relabel(data = 'new') * price.relabel(data = 'old')).sort('y', 'm')
```

### perdictable with the cell framework

We can run the function and use a cell to store the output...

```
c = cell(perdictable(fake_ts, on = pk), ticker = oil, expiry = expiry)()
c.data
```

```
recalculated_cell = c.go(1) ## force a recalculation
recalculated_cell.data
```

We observe that the cell, when recalculates, automatically caches the history and does not recalculate fake\_ts. This is not magic. When a cell calculates its function, it provides the function with the variables it needs. Once calculated, it stores the output in **data** and will be able to provide **data** to the function next time, allowing it to avoid re-running expired calculations. Then cell will store the functions's result back in the **data** key for later use and this is repeated.

### perdictable API

Parameters **on**, **renames** and **defaults** parameters determine the way the data is joined. If defaults is missing, the defaults from the function are used:

```
function = lambda price, quantity = 1: price * quantity
price = dictable(product = ['apple', 'banana', 'cherry'], price = [1,2,3])
quantity = dictable(product = ['apple', 'banana', 'damson'], quantity = [2,3,4])
perdictable(function, on = 'product')(price = price, quantity = quantity) ## cherry should a
```

If you want to see the full calculations and inputs to the function set **include\_inputs=True**:

```
perdictable(function, on = 'product', include_inputs = True)(price = price, quantity = quant
```

If you want output column to be not data, use **col**:

```
perdictable(function, on = 'product', include_inputs = True, col = 'cost')(price = price, qu
```

The **if\_none** parameter determines how data is calculated for rows that have expired but their data is None:

```
expiry = dictable(product = ['apple', 'banana', 'cherry'], expiry = [dt(-2), dt(-1), dt(1)])
previous_data = dictable(product = ['apple', 'banana', 'cherry'], data = [None, 'some value', 'some value'])
perdictable(function, on = 'product', include_inputs = True, if_none = False)(price = price,
```

```
perdictable(function, on = 'product', include_inputs = True, if_none = True)(price = price,
```

Some function want to receive historic data and they use it themselves. Parameter **output\_is\_input** controls this. For example: If your function is pulling historic prices from yahoo finance, you can use existing data to ask yahoo for only recent ones.

```
def running_total_costs(price, quantity=1, data=0):
    return data + price * quantity
```

```
previous_data = dictable(product = ['apple', 'banana', 'cherry', 'damson'], data = [10, 20,
perdictable(running_total_costs, on = 'product', include_inputs = True)(price = price, quant
```

```
## if you don't want existing data to be presented to the function:
perdictable(running_total_costs, on = 'product', include_inputs = True, output_is_input = Fa
```

## Conclusions

pyg.base.join allows us to create joined table with the variables we need. This is leveraged by perdictable so that the 'atomic' data we work with is not a single timeseries but a whole table of timeseries data indexed by some keys. We can use various perdictable parameters to control cache policy. All this is done with very little additional code, allowing us to manage quite a lot of data items with very little effort while managing caching expired items.

## pyg.timeseries

Given pandas, why do we need this timeseries library? pandas is amazing but there are a few features in pyg.timeseries designed to enhance it.

There are three issues with pandas that pyg.timeseries tries to address:

- pandas works on pandas objects (obviously) but not on numpy arrays.
- pandas handles nan within timeseries inconsistently across its functions. This makes your results sensitive to reindexing/resampling. E.g.:
  - a.expanding() & a.ewm() **ignore** nan's for calculation and then forward-fill the result.
  - a.diff(), a.rolling() **include** any nans in the calculation, leading to nan propagation.
- pandas is great if you have the full timeseries. However, if you now want to run the same calculations in a live environment, on recent data, you have to append recent data at the end of the DataFrame and rerun.

pyg.timeseries tries to address this:

- pyg.timeseries agrees with pandas 100% (if there are no nan in the dataframe) while being of comparable speed
- pyg.timeseries works seamlessly on pandas objects and on numpy arrays, with no code change.
- pyg.timeseries handles nan consistently across all its functions, 'ignoring' all nan, making your results consistent regardless of reindexing/resampling.
- pyg.timeseries exposes the state of the internal function calculation. The exposure of internal state allows us to calculate the output of additional data **without** re-running history. This speeds up of two very common problems in finance:
  - risk calculations, Monte Carlo scenarios: We can run a trading strategy up to today and then generate multiple scenarios and see what-if, without having to rerun the full history.
  - live versus history: pandas is designed to run a full historical simulation. However, once we reach "today", speed is of the essence and running a full historical simulation every time we ingest a new price, is just too slow. That is why most fast trading is built around fast state-machines. Of course, making sure research & live versions do the same thing is tricky. pyg gives you the ability to run two systems in parallel with almost the same code base: run full history overnight and then run today's code base instantly, instantiated with the output of the historical simulation.

## Agreement between pyg.timeseries and pandas

```
from pyg import *; import pandas as pd; import numpy as np
s = pd.Series(np.random.normal(0,1,10000), _drange(-9999)); a = s.values
t = pd.Series(np.random.normal(0,1,10000), _drange(-9999))
```

```
assert abs(s.count() - ts_count(s)) < 1e-10
assert abs(s.mean() - ts_mean(s)) < 1e-10
```

```
assert abs(s.sum() - ts_sum(s)) < 1e-10
assert abs(s.std() - ts_std(s)) < 1e-10
assert abs(s.skew() - ts_skew(s)) < 1e-10
```

```
assert abs(ewma(s, 10) - s.ewm(10).mean()).max() < 1e-10
assert abs(ewmstd(s, 10) - s.ewm(10).std()).max() < 1e-10
assert abs(ewmvar(s, 10) - s.ewm(10).var()).max() < 1e-10
assert abs(ewmcor(s, t, 10) - s.ewm(10).corr(t)).max() < 1e-10
```

```
assert abs(expanding_sum(s) - s.expanding().sum()).max() < 1e-10
assert abs(expanding_mean(s) - s.expanding().mean()).max() < 1e-10
assert abs(expanding_std(s) - s.expanding().std()).max() < 1e-10
assert abs(expanding_skew(s) - s.expanding().skew()).max() < 1e-10
assert abs(expanding_min(s) - s.expanding().min()).max() < 1e-10
assert abs(expanding_max(s) - s.expanding().max()).max() < 1e-10
assert abs(expanding_median(s) - s.expanding().median()).max() < 1e-10
```

```
assert abs(rolling_sum(s,10) - s.rolling(10).sum()).max() < 1e-10
assert abs(rolling_mean(s,10) - s.rolling(10).mean()).max() < 1e-10
assert abs(rolling_std(s,10) - s.rolling(10).std()).max() < 1e-10
assert abs(rolling_skew(s,10) - s.rolling(10).skew()).max() < 1e-10
assert abs(rolling_min(s,10) - s.rolling(10).min()).max() < 1e-10
assert abs(rolling_max(s,10) - s.rolling(10).max()).max() < 1e-10
assert abs(rolling_median(s,10) - s.rolling(10).median()).max() < 1e-10
assert abs(rolling_quantile(s,10,0.3)[0.3] - s.rolling(10).quantile(0.3)).max() < 1e-10 ##
```

### Quick performance comparison

pyg, when run on pandas dataframes rather than arrays, is of comparable speed to pandas

```
compare = dictable(op = ['rolling_sum', 'rolling_mean', 'rolling_std', 'rolling_min', 'rolling_max', 'rolling_median'],
                  pyg = [rolling_sum, rolling_mean, rolling_std, rolling_min, rolling_max, rolling_median],
                  pandas = [s.rolling(10).sum, s.rolling(10).mean, s.rolling(10).std, s.rolling(10).min, s.rolling(10).max, s.rolling(10).median])

compare += dictable(op = ['expanding_sum', 'expanding_mean', 'expanding_std', 'expanding_min', 'expanding_max', 'expanding_median'],
                  pyg = [expanding_sum, expanding_mean, expanding_std, expanding_min, expanding_max, expanding_median],
                  pandas = [s.expanding().sum, s.expanding().mean, s.expanding().std, s.expanding().min, s.expanding().max, s.expanding().median])

compare += dictable(op = ['ewma', 'ewmstd', 'ewmvar'],
                  pyg = [ewma, ewmstd, ewmvar],
                  pandas = [s.ewm(10).mean, s.ewm(10).std, s.ewm(10).var]).do(lambda v: timer(v, 100000))

print(compare(winner = lambda pyg, pandas: 'pyg' if pyg < pandas * 0.8 else 'pandas' if pyg > pandas * 0.8))
```

### pyg and numpy arrays

pyg supports numpy arrays natively. Indeed, pyg is 3-5 times faster on numpy arrays.

```
a = s.values
assert abs(ts_count(a) - ts_count(s)) < 1e-10
assert abs(ts_mean(a) - ts_mean(s)) < 1e-10
assert abs(ts_sum(a) - ts_sum(s)) < 1e-10
assert abs(ts_std(a) - ts_std(s)) < 1e-10
assert abs(ts_skew(a) - ts_skew(s)) < 1e-10
```

```
assert abs(ewma(s, 10) - ewma(a,10)).max() < 1e-10
assert abs(ewmstd(s, 10) - ewmstd(a,10)).max() < 1e-10
assert abs(ewmvar(s, 10) - ewmvar(a,10)).max() < 1e-10
assert abs(ewmcor(s, t, 10) - ewmcor(a, t.values, 10)).max() < 1e-10
```

```

assert abs(expanding_sum(s) - expanding_sum(a)).max() < 1e-10
assert abs(expanding_min(s) - expanding_min(a)).max() < 1e-10
assert abs(expanding_max(s) - expanding_max(a)).max() < 1e-10
assert abs(expanding_mean(s) - expanding_mean(a)).max() < 1e-10
assert abs(expanding_std(s) - expanding_std(a)).max() < 1e-10
assert abs(expanding_skew(s) - expanding_skew(a)).max() < 1e-10
assert abs(expanding_median(s) - expanding_median(a)).max() < 1e-10

```

```

assert abs(rolling_sum(s,10) - rolling_sum(a,10)).max() < 1e-10
assert abs(rolling_min(s,10) - rolling_min(a,10)).max() < 1e-10
assert abs(rolling_max(s,10) - rolling_max(a,10)).max() < 1e-10
assert abs(rolling_mean(s,10) - rolling_mean(a,10)).max() < 1e-10
assert abs(rolling_std(s,10) - rolling_std(a,10)).max() < 1e-10
assert abs(rolling_skew(s,10) - rolling_skew(a,10)).max() < 1e-10
assert abs(rolling_median(s,10) - rolling_median(a,10)).max() < 1e-10

```

## pandas treatment of nan

Suppose we have weekly data that at some point we resample to daily... The two look the same...

```

t0 = dt_bump('20210301', '-999w')
days = drange(t0, '20210301', '1b')
weekly = pd.Series(np.random.normal(0,1,1000), drange(t0, None, '1w')); weekly.name = 'weekly'
daily = weekly.reindex(days); daily.name = 'daily'
pd.concat([weekly,daily], axis = 1)

```

... but any calculation using the daily will yield a different result from a calculation on the weekly which is then resampled to daily:

```
pd.concat([weekly.ewm(4).mean().reindex(days), daily.ewm(4).mean()], axis = 1) ## The result
```

```
pd.concat([weekly.diff().reindex(days), daily.diff()], axis = 1) ## The result depends on wh
```

Indeed, for diff, daily.diff() is all nan!

## pyg.timeseries treatment of nans

pyg treats nan as if they are not there, so the fact that we resampled the data and introduced lots of nan's does not affect the calculations. We find this to be a more logical (and less error prone) approach.

```
nona(pd.concat([ewma(weekly, 4).reindex(days), ewma(daily,4)], axis = 1)) ## The two match e
```

```
nona(pd.concat([diff(weekly).reindex(days), diff(daily)], axis = 1)) ## The result depends o
```

## Using pyg.timeseries to manage state

One of the problem in timeseries analysis is writing research code that works in analysing past data but ideally, the same code can be used in live application. One easy approach is “stick the extra data point at the end and run it again from 1980”. This leaves us with a single code base but for many live applications (e.g. live trading), this is not viable.

Further, given our positions today, we may want to run simulations of “what happens next?” to understand what the system is likely to do should various events occur. Risk calculations are expensive and re-running 10k Monte Carlo scenarios, each time running from 1980 is expensive.

Conversely, we can run research and live systems on two separate code base. This makes live systems responsive but six months down the line, we realise research code base and live code base did not do quite the same thing.

pyg approaches this problem by exposing the internal state of each of its calculation. Each function has two versions:

- function(...) returns the calculation as performed by pandas

- `function_(...)` returns a dictionary of `dict(data = , state = )`. The data agrees with `function(...)` while the state is a dict we can instantiate new calculations with.

```
from pyg import *
history = pd.Series(np.random.normal(0,1,1000), range(-1000,-1))
history_signal = ewma(history, 10)
history_signal # The output consists of 'data' and 'state' where data matches a normal ewma

live = pd.Series(np.random.normal(0,1,10), range(9))
live_signal = ewma(live, 10, state = history_signal.state) ## I only feed in live timeseries
'live: from today onwards', live_signal

joint_data = pd.concat([history, live])
joint_signal = ewma(joint_data, 10)
assert eq(live_signal, joint_signal[dt(0):]) # The live signal is the same, even though it
joint_signal[dt(0):]
```

This allows us to set up three parallel pipelines that share a virtually identical codebase:

workflow	historic data	live data	risk analysis
when run?	research/overnight	live	overnight
data source?	ts = long timeseries	a = short ts/array	1000's of sims
speed?	slow, non-critical	instantaneous	quick
apply f to data	$x_ = f_{\text{ts}}$	$x = f(a, **x_)$	same as live
apply g	$y_ = g_{\text{ts}}(x_)$	$y = g(a, x, **y_)$	same as live
final result h	$z_ = h_{\text{ts}}(x_, y_)$	$z = h(a, x, y, **z_)$	same as live

Note that for live trading or risk analysis, we tend to switch and run on numpy arrays rather than pandas object. This speeds up the calculations while introduces no code change. In the example below we explore how to create state-aware, functions within pyg. The paradigm is that for most functions, `function_` will return not just the timeseries output but also the states.

### Example: creating a function exposing its state

Suppose we try to write an ewma crossover function (the difference of two ewma). We want to normalize it by its own volatility. Traditionally we will write:

```
def pandas_crossover(a, fast, slow, vol):
    fast_ewma = a.ewm(fast).mean()
    slow_ewma = a.ewm(slow).mean()
    raw_signal = fast_ewma - slow_ewma
    signal_rms = (raw_signal**2).ewm(vol).mean()**0.5
    signal_rms[signal_rms==0] = np.nan
    normalized = raw_signal/signal_rms
    return normalized

a = pd.Series(np.random.normal(0,1,10000), range(-9999)); fast = 10; slow = 30; vol = 50
pandas_x = pandas_crossover(a, fast, slow, vol)
pandas_x
```

We can quickly rewrite it using pyg:

```
def crossover(a, fast, slow, vol):
    fast_ewma = ewma(a, fast)
    slow_ewma = ewma(a, slow)
    raw_signal = fast_ewma - slow_ewma
    signal_rms = ewmrms(raw_signal, vol)
    signal_rms = v2na(signal_rms)
    normalized = raw_signal/signal_rms
    return normalized
```



```
x = crossover(a, fast, slow, vol)
assert abs(x-pandas_x).max() < 1e-10
x
```

And with very little additional effort, we can write a new function that also exposes the internal state:

```
_data = 'data'
def crossover_(a, fast, slow, vol, instate = None):
    state = Dict(fast = {}, slow = {}, vol = {}) if instate is None else instate
    fast_ewma_ = ewma_(a, fast, instate = state.fast)
    slow_ewma_ = ewma_(a, slow, instate = state.slow)
    raw_signal = fast_ewma_.data - slow_ewma_.data
    signal_rms = ewmrms_(raw_signal, vol, instate = state.vol)
    normalized = raw_signal/v2na(signal_rms.data)
    return Dict(data = normalized, state = Dict(fast = fast_ewma_.state, slow = slow_ewma_.state))

crossover_.output = ['data', 'state'] # output declares the function to have a dict output a

def crossover(a, fast, slow, vol, state = None):
    return crossover_(a, fast, slow, vol, instate = state).data

x_ = crossover_(a, fast, slow, vol)
assert eq(x, x_.data) and eq(x, crossover(a, fast, slow, vol))
x_.data
```

The three give identical results and we can also verify that `crossover_` will allow us to split the evaluation to the long-history and the new data:

```
history = a[:9900]
live = a[9900:].values
x_history = crossover_(history, 10, 30, 50)
x_live = crossover(live, 10, 30, 50, state = x_history.state)
x_ = crossover_(a, fast, slow, vol)
assert eq(x_live, x_.data[9900:].values)
```

Have we gained anything?

```
pandas_old = timer(pandas_crossover, 100, time = True)(history, 10, 30, 50)
x_history = crossover_(history, 10, 30, 50)
x_history_time = timer(crossover_, 100, time = True)(history, 10, 30, 50)
x_live = timer(crossover, 100, time = True)(live, 10, 30, 50, state = x_history.state)
'pandas: ', pandas_old.microseconds//1000, 'pyg history: ', x_history_time.microseconds//1000
```

We see that pyg is already faster than pandas. Running just the new data using numpy arrays, is about 4-5 times faster still. Indeed, running 10k 100-day forward scenarios take about 2 seconds at most.

```
scenarios = np.random.normal(0, 1, (100, 10000))
x_scenarios = timer(crossover)(scenarios, 10, 30, 50, state = x_history.state)
```

Using cells, our code looks like this, with live and historical codebase looking pretty similar

```
x_history = cell(crossover_, a = history, fast = 10, slow = 30, vol = 50)()
x_live = cell(crossover, a = live, fast = 10, slow = 30, vol = 50, state = x_history)()
x_history

pd.concat([pd.Series(x_live.data, pandas_x.index[-100:]), pandas_x.iloc[-100:]], axis = 1)
```

## pyg.timeseries decorators

There are a few decorators that are relevant to timeseries analysis ## pd2np and compiled We write most of our underlying functions assuming the function parameters are 1-d numpy arrays. If you want them numba.jit compiled, please use the compiled operator.



```

from pyg import *
import pandas as pd; import numpy as np
@pd2np
@compiled
def sumsq(a, total = 0.0):
    res = np.empty_like(a)
    for i in range(a.shape[0]):
        if np.isnan(a[i]):
            res[i] = np.nan
        else:
            total += a[i]**2
            res[i] = total
    return res

```

It is not surprising that `sumsq` works for arrays. Notice how `np.isnan` is handled to ensure nans are skipped.

```

a = np.arange(5)
sumsq(a)

```

**pd2np** will convert a pandas Series to arrays, run the function and convert back to pandas. This will only work for a 1-dimensional objects, so no df nor 2-d np.ndarray.

```

s = pd.Series(a, drange(-4))
sumsq(s)

```

## loop

We decorate `sumsq` with the **loop** decorator. Once we introduce loop, The function will loop over columns of a DataFrame or a numpy array:

```

@loop(pd.DataFrame, dict, list, np.ndarray)
@pd2np
@compiled
def sumsq(a, total = 0):
    res = np.empty_like(a)
    for i in range(a.shape[0]):
        if np.isnan(a[i]):
            res[i] = np.nan
        else:
            total += a[i]**2
            res[i] = total
    return res

df = pd.DataFrame(dict(a = a, b = a+1), drange(-4))
df

```

```
sumsq(df)
```

Indeed, since we asked it to loop over dict, list and numpy array (2d)

```
sumsq(dict(a = a, b = a+1))
```

```
sumsq(df.values)
```

## presync: manage indexing and date stamps

Suppose the function takes two (or more) timeseries.

```

@presync(index = 'inner')
@loop(pd.DataFrame, np.ndarray)
@pd2np

```

```
def product(a, b):
    return a * b
```

```
a = np.arange(5); b = np.arange(5)
product(a,b)
```

What happens when the weights and the timeseries are unsynchronized?

```
a_ = pd.Series(a, drange(-4)) ; a_.name = 'a'
b_ = pd.Series(b, drange(-3,1)); b_.name = 'b'
pd.concat([a_, b_], axis=1)
```

```
product(a_, b_) ## just the inner values
```

```
product.oj(a_, b_) ## outer join
```

```
product.oj.ffill(a_, b_) ## outer join and forward-fill
```

## presync and numpy arrays

When we deal with thousands of equities, one way of speeding calculations is by stacking them all onto huge dataframes. This does work but one is always busy fiddling with 'the universe' one is trading. We took a slightly different approach:

- We define a global timestamp.
- We then sample each timeseries to that global timestamp, dropping the early history where the data is all nan. (df\_fillna(ts, index, method = 'fna')).
- We then do our research on these numpy arrays.
- Finally, once we are done, we resample back to the global timestamp.

While we are in numpy arrays, we can 'inner join' by recognising the 'end' of each array shares the same date. Indeed df\_index, df\_reindex and presync all work seamlessly on np.ndarray as well as DataFrames, under that assumption that **the end of all arrays are in sync**.

We find this approach saves on memory and on computation time. It also lends itself to being able to retrieve and create specific universes for specific trading ideas. It is not without its own issues but that is a separate discussion.

```
a = np.arange(5); b = np.arange(1,5)
a, b
```

```
product(a, b)
```

```
us = calendar('US')
dates = pd.Index(us.drange('-40y', 0, '1b'))
```

```
universe = dictable(stock = ['msft', 'appl', 'tsla'], n = [10000, 8000, 7000])
universe = universe(ts = lambda n: pd.Series(np.random.normal(0,1,n+1), us.drange('-%ib'%n,
universe
```

```
universe = universe(rtn = lambda ts: ts.values)
universe = universe(price = lambda rtn : cumsum(rtn))
universe = universe(vol = lambda rtn: ewmstd(rtn, 30))
universe
```

```
presync(lambda tss: np.array(tss).T)(universe.vol)
```


```
universe = universe.do(lambda value: np_reindex(value, dates), 'rtn', 'price', 'vol')
universe
```

```
vol = pd.concat(universe.vol, axis = 1); vol.columns = universe.stock
vol
```


## pyg.timeseries.ewma

The ewm functions implement the concept of time which we think is worthwhile explaining. We start with an example


```
from pyg import *; import numpy as np; import pandas as pd
rtn = pd.Series(np.random.normal(0.01,1,5218), drange(2000,2020, '1b'))
price = cumsum(rtn); price.name = 'price'
smooth = ewma(price, 50); smooth.name = 'smooth'
pd.concat([price, smooth], axis = 1).plot()
```



```
## now suppose somehow we lost 4 years of data...
bad = price.copy()
bad[dt(2008):dt(2012)] = np.nan
smooth_bad = ewma(bad, 50); smooth_bad.name = 'smooth_bad'
pd.concat([bad, smooth_bad], axis = 1).plot()
```



```
smooth_with_time = ewma(bad, 50, time = 'b')
smooth_with_time.name = 'smooth_with_business_day_clock'
pd.concat([bad, smooth_bad, smooth_with_time], axis = 1).plot()
```



What happened here? How can smooth with clock track better? The answer is that if you provide a clock, ewma can recognise that 4 years have passed. The old data is irrelevant, it forgets the old position and start with most of the weight on the more recent observations

### What happens if the clock does not move at all?

Suppose we now choose to calculate daily ewma but we are doing this with intraday data. If the number of data points per day is constant and known, then this can be done with ease. Using time parameter, we can do this even for an irregularly spaced timeseries. We first just create some fake data:


```
import datetime
bar = datetime.timedelta(minutes = 5)
all_bars = [t for t in range(2020, 2021, bar)]
ts = pd.Series(np.random.normal(0.01/24, 1, len(all_bars)), all_bars)
price = cumsum(ts); price.name = 'intraday'

bars = np.array([t for t in range(2020, 2021, bar) if t.hour>=8 and t.hour<=17]) ## trading
irregular = bars[np.random.normal(0,1,len(bars))>0]

ts = pd.Series(np.random.normal(0.01/24, 1, len(bars)), bars)
price = cumsum(ts); price.name = 'intraday'
price = price[irregular] ## remove half the bars randomly

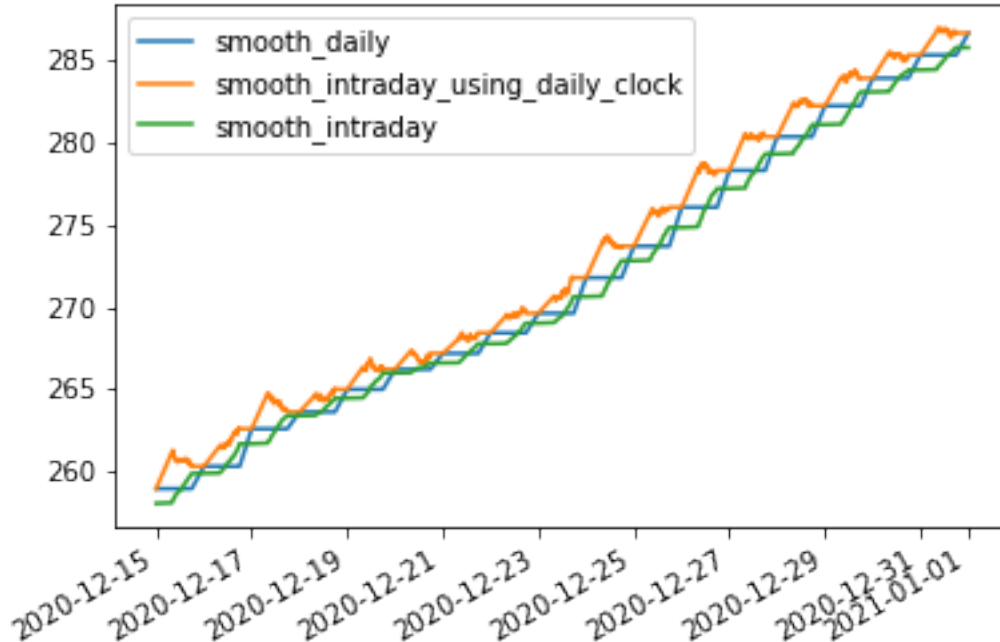
days = range(2020,2021,1)

daily = price.reindex(days, method = 'ffill'); daily.name = 'daily'
pd.concat([price, daily], axis = 1).ffill().plot()
```



- By setting clock to **daily** we tell ewma that all the hourly data in the same day are 'on same clock'. And it will use historic end-of-day prices while updating today the last point until the clock moves to tomorrow's reading
- If we set the clock to **fraction**, it will update continuously throughout the day

```
smooth_intra = ewma(price, 20, 'f'); smooth_intra.name = 'smooth_intraday' ## roughly matching
smooth_intra_using_d = ewma(price, 20, time = 'd'); smooth_intra_using_d.name = 'smooth_intra_using_d'
smooth_daily = ewma(daily, 20); smooth_daily.name = 'smooth_daily'
pd.concat([smooth_daily, smooth_intra_using_d, smooth_intra], axis = 1).ffill()[dt(2020,12,15):dt(2021,1,1)]
```



- **smooth\_daily** is calculated on daily basis and is constant within the day and experiences jumps on EOD
- **time = 'd'** option front-runs daily, but at the price of being more volatile intra-day. On end-of-day the two versions agree
- **time = 'f'** is a smoother version of daily. It is leading, but not by much

```
pd.concat([smooth_intra_using_d.reindex(days, method = 'ffill'), smooth_daily], axis = 1)
## on end-of-day we have an exact match between time = 'd' and daily smooth
```

### ***What are valid time parameters?***

- None: If None is provided, any (non-nan) observation is considered to be a clock ticking
- i: index of timeseries. The clock ticks also for nan observations. This is the default for pandas
- f: fraction of day
- b/d/w/m/q/y: business day/daily/weekly/monthly/quarterly or yearly
- Calendar: the business day as defined by the calendar provided
- For full control, you can provide a timeseries of non-decreasing times matching the original array

### **Indices and tables**

- **genindex**
- **modindex**
- **search**

# Python Module Index

## *p*

`pyg`

`pyg.base._getitem`

`pyg.base._inspect`

`pyg.base._types`